The Bonnie Situation, Pulp Fiction (1994), Quentin Tarantino

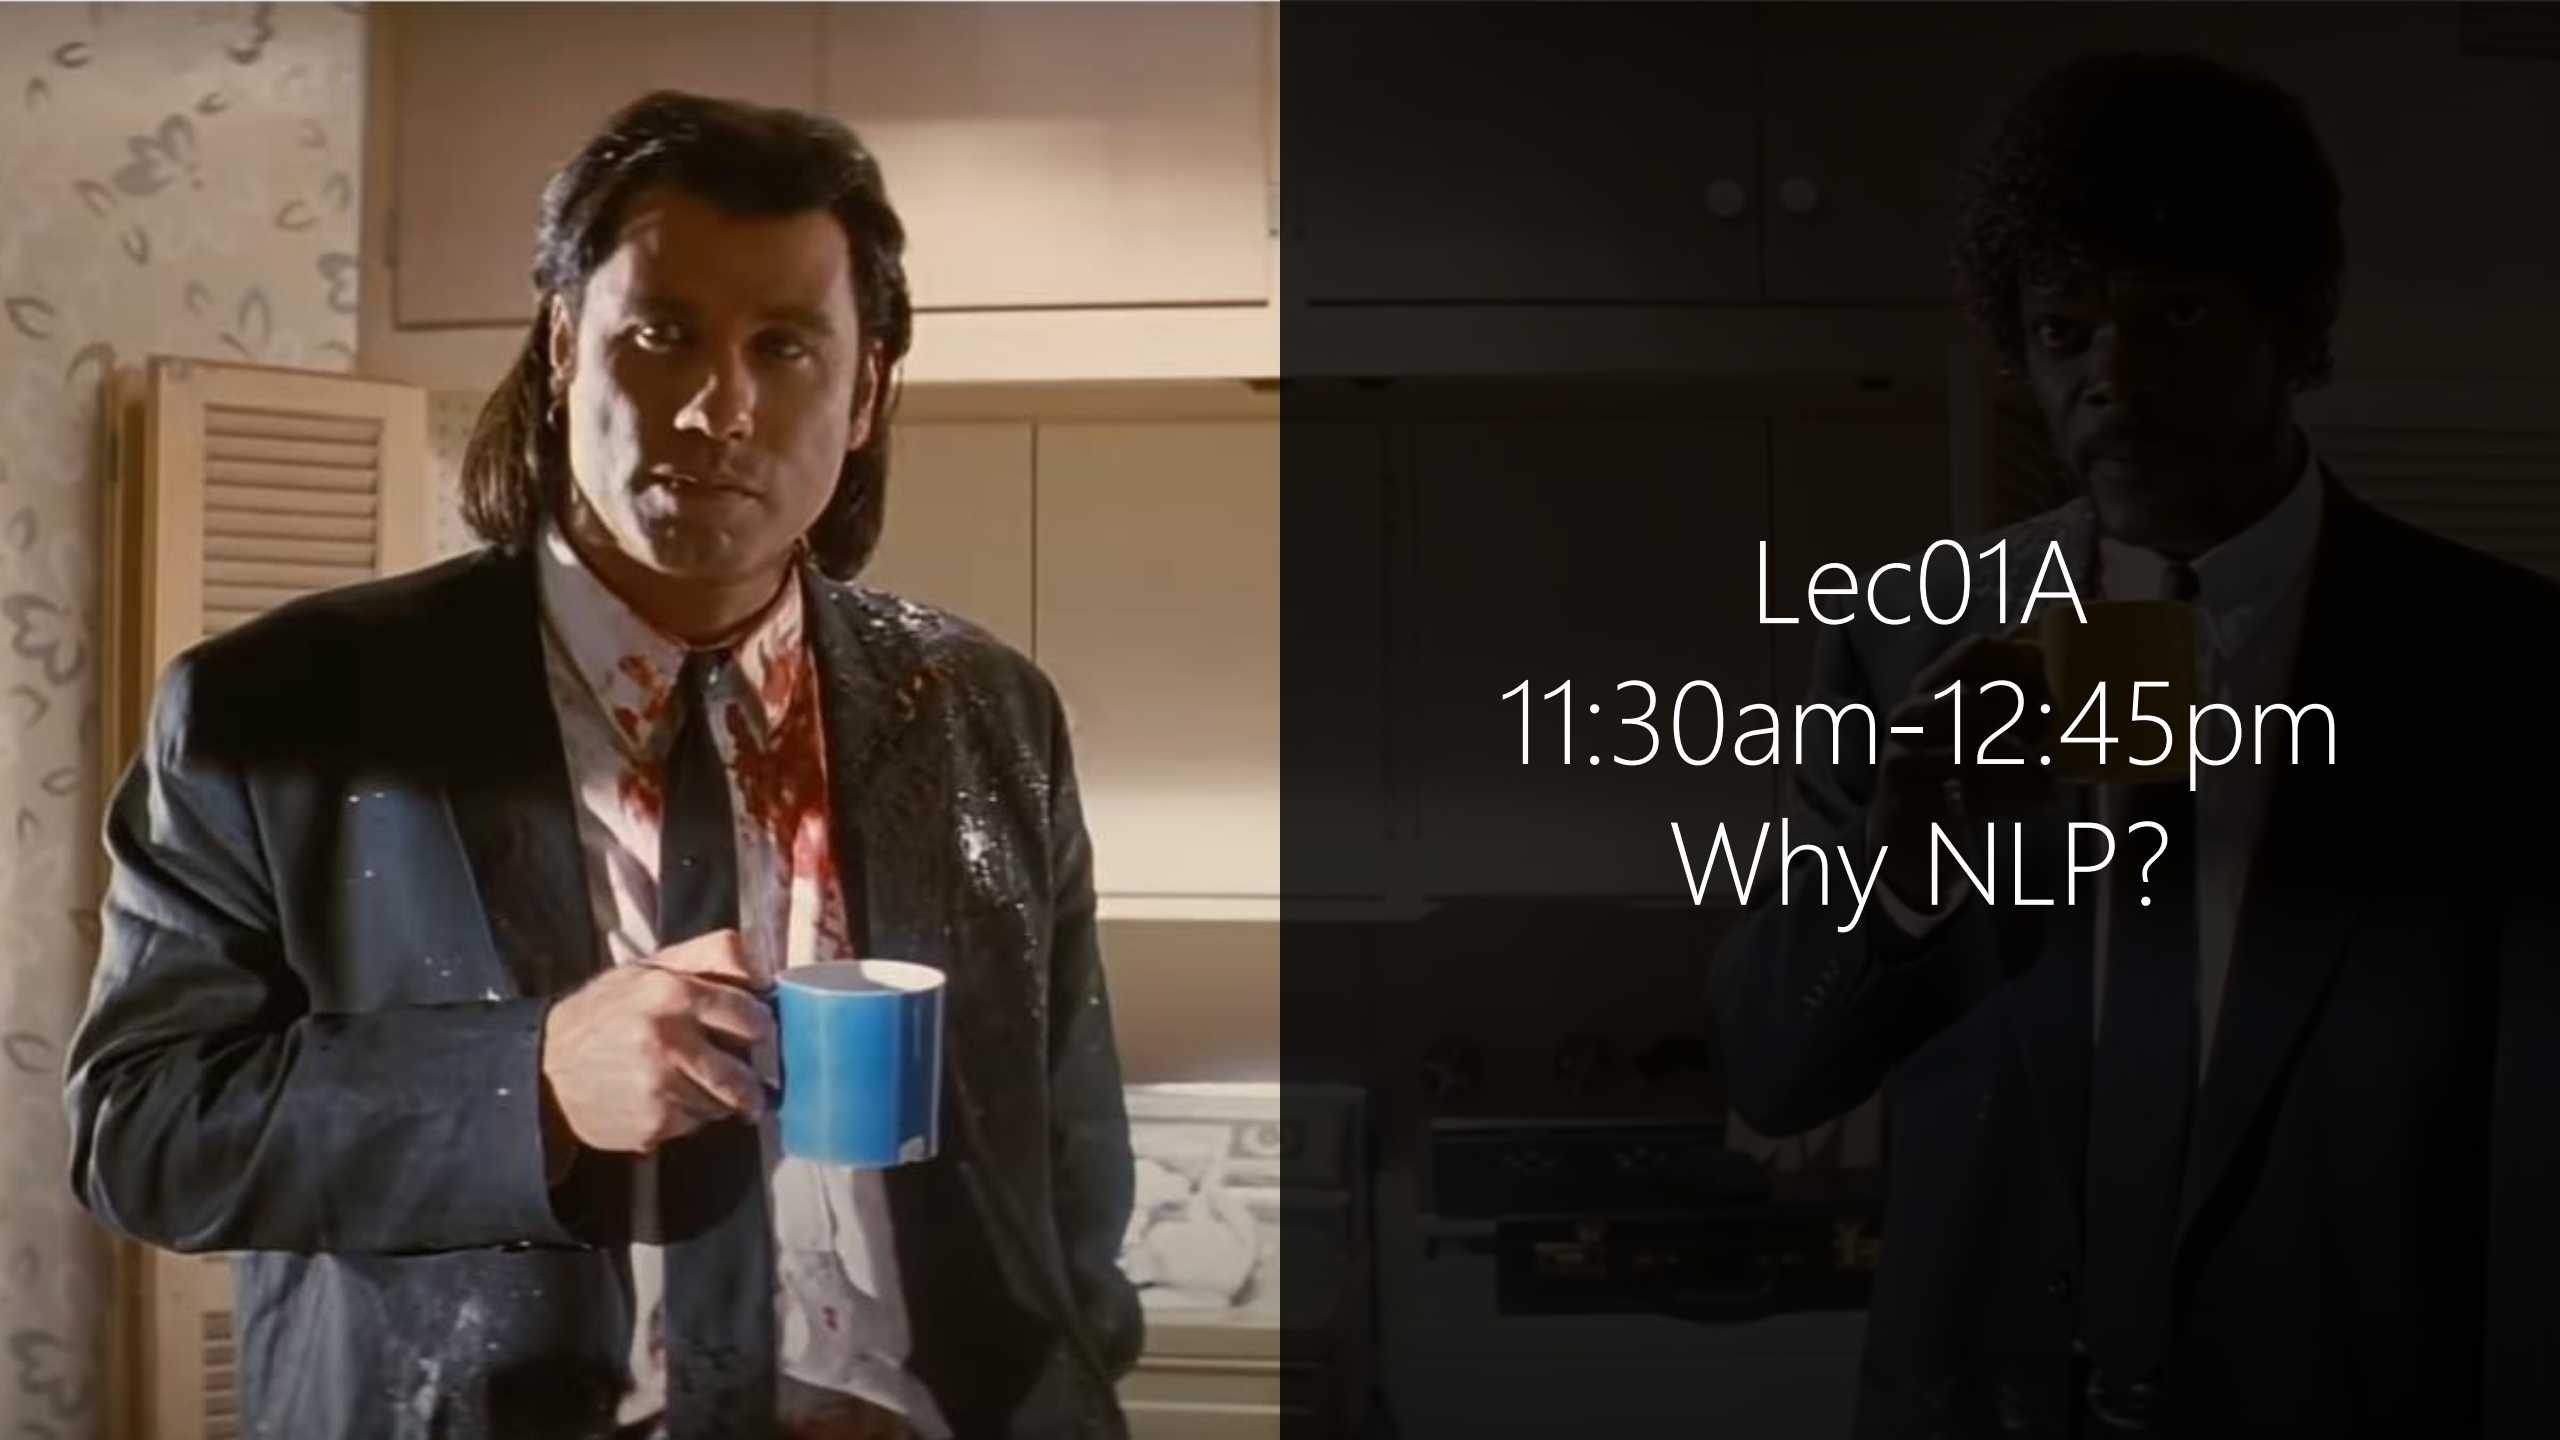

Lec01A
11:30am-12:45pm
Why NLP?

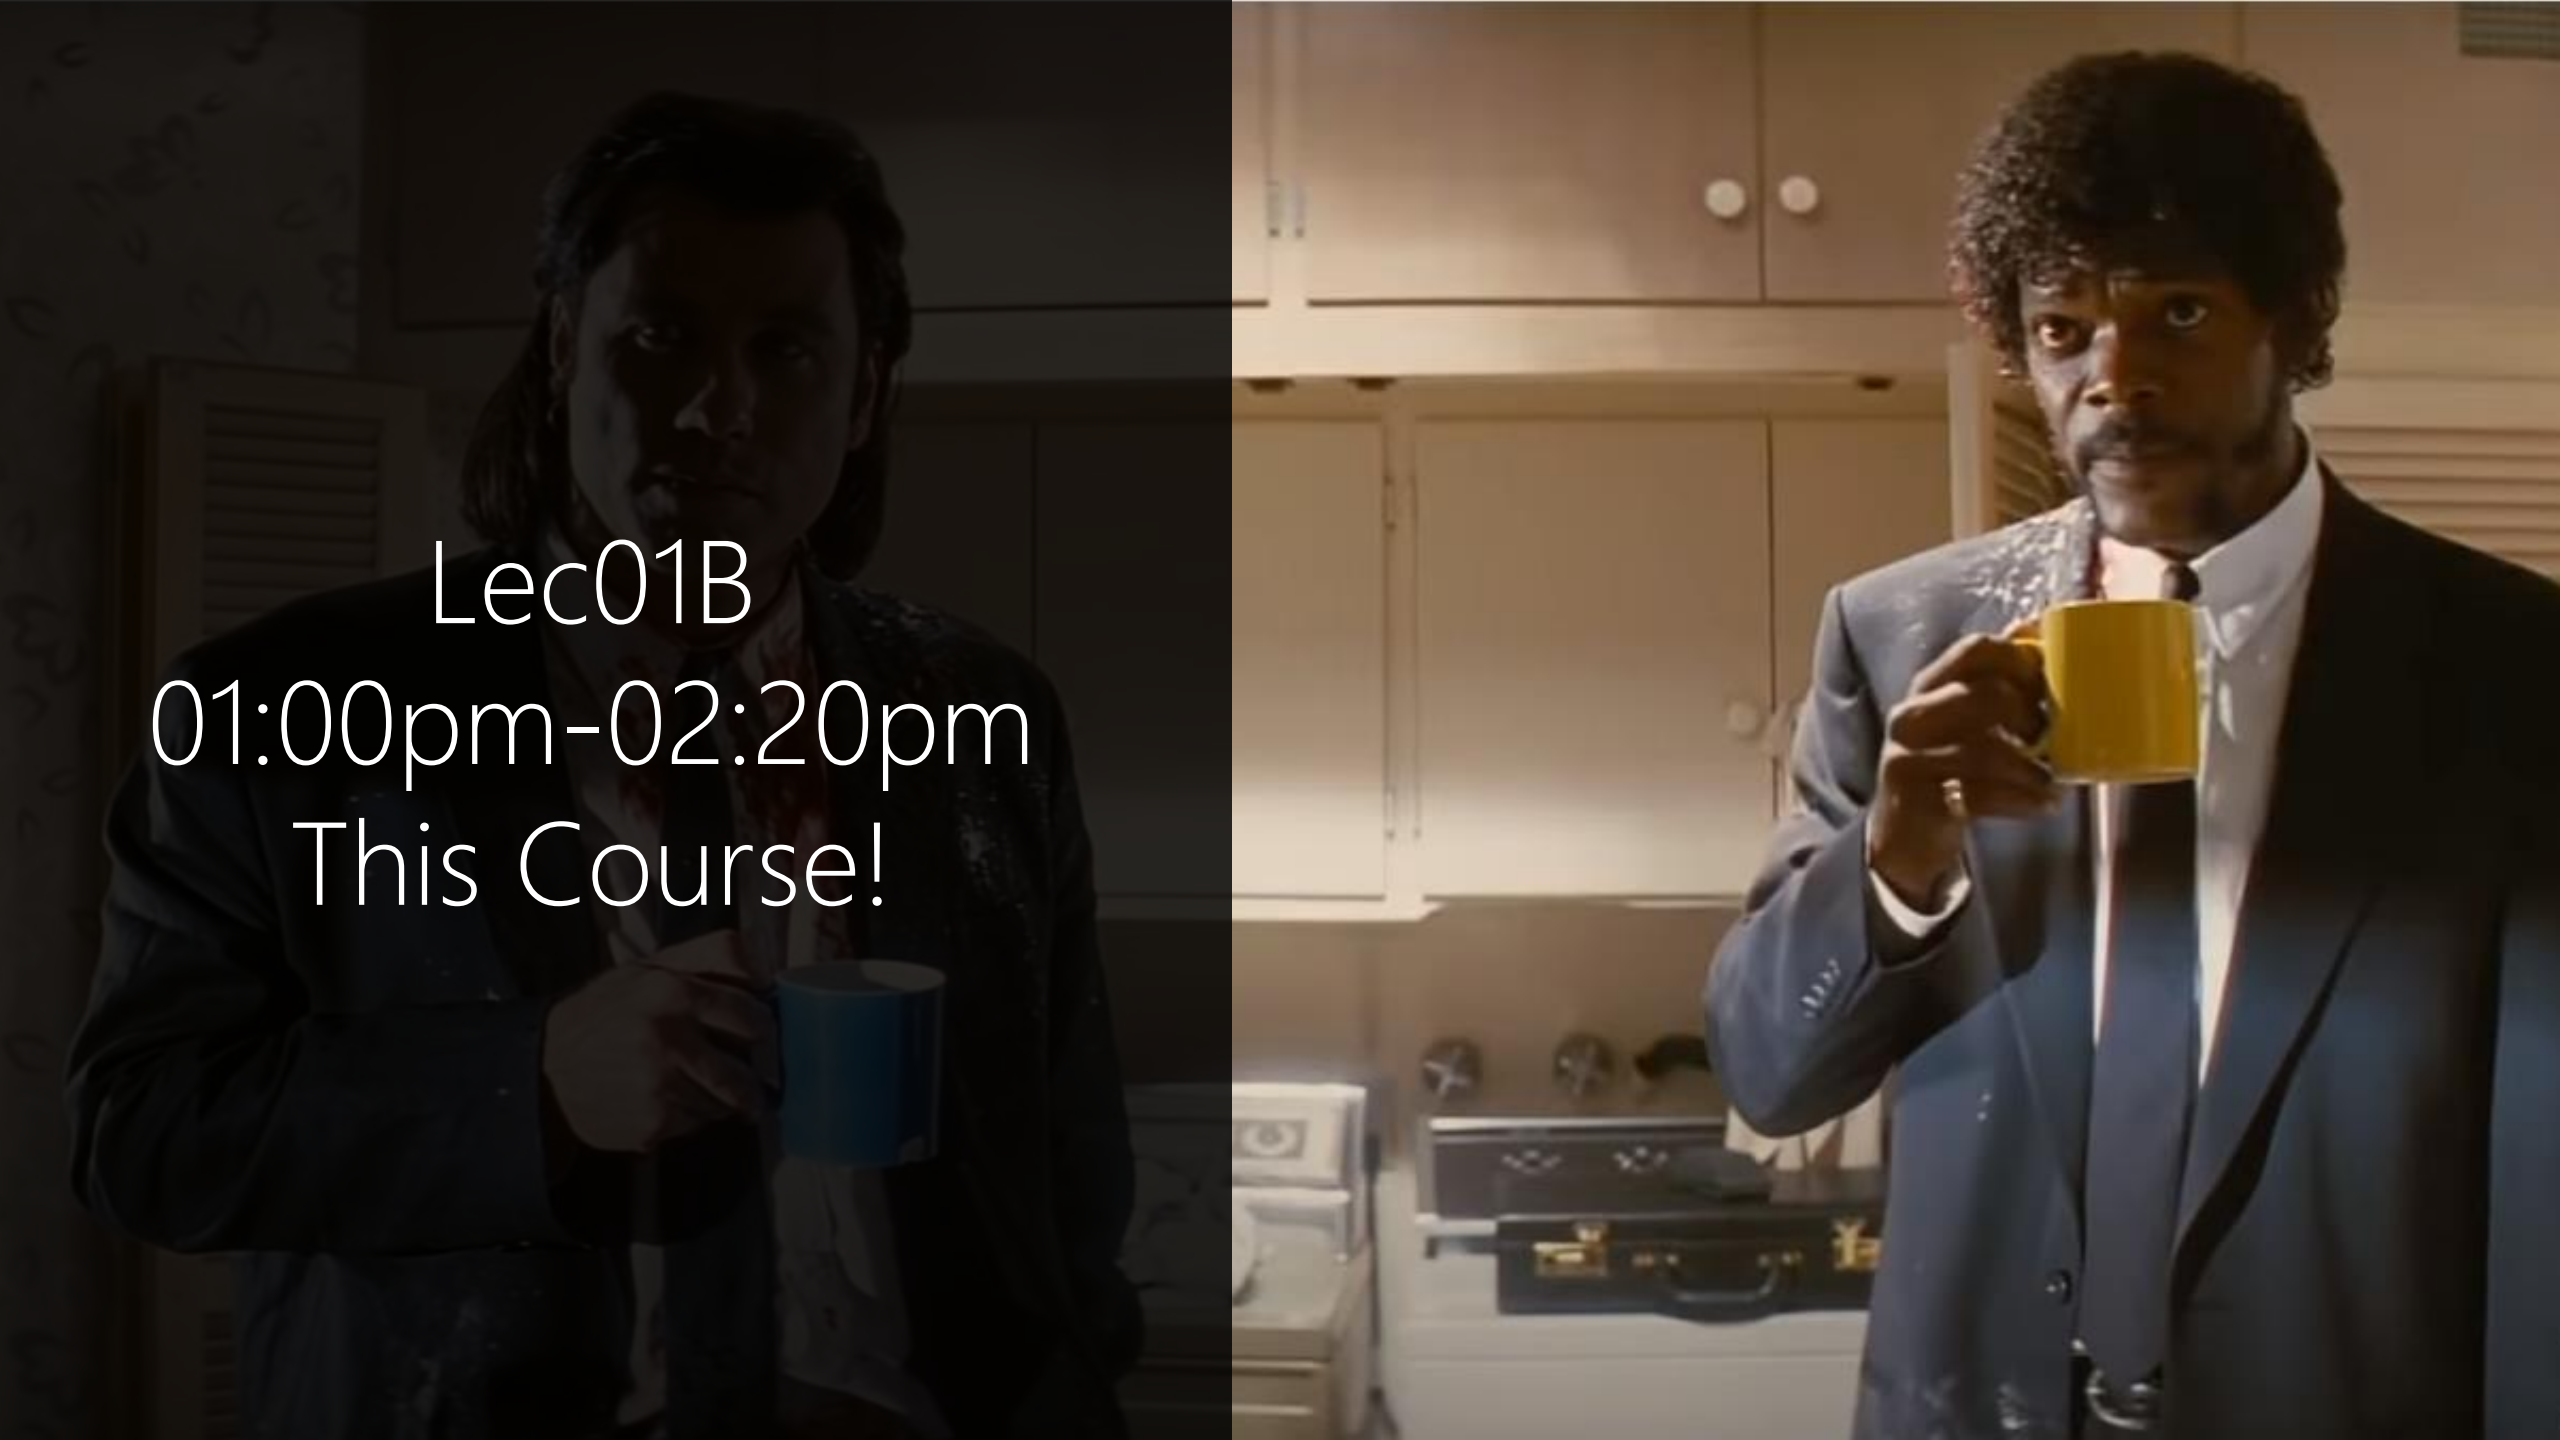
Lec01B
01:00pm-02:20pm
This Course!

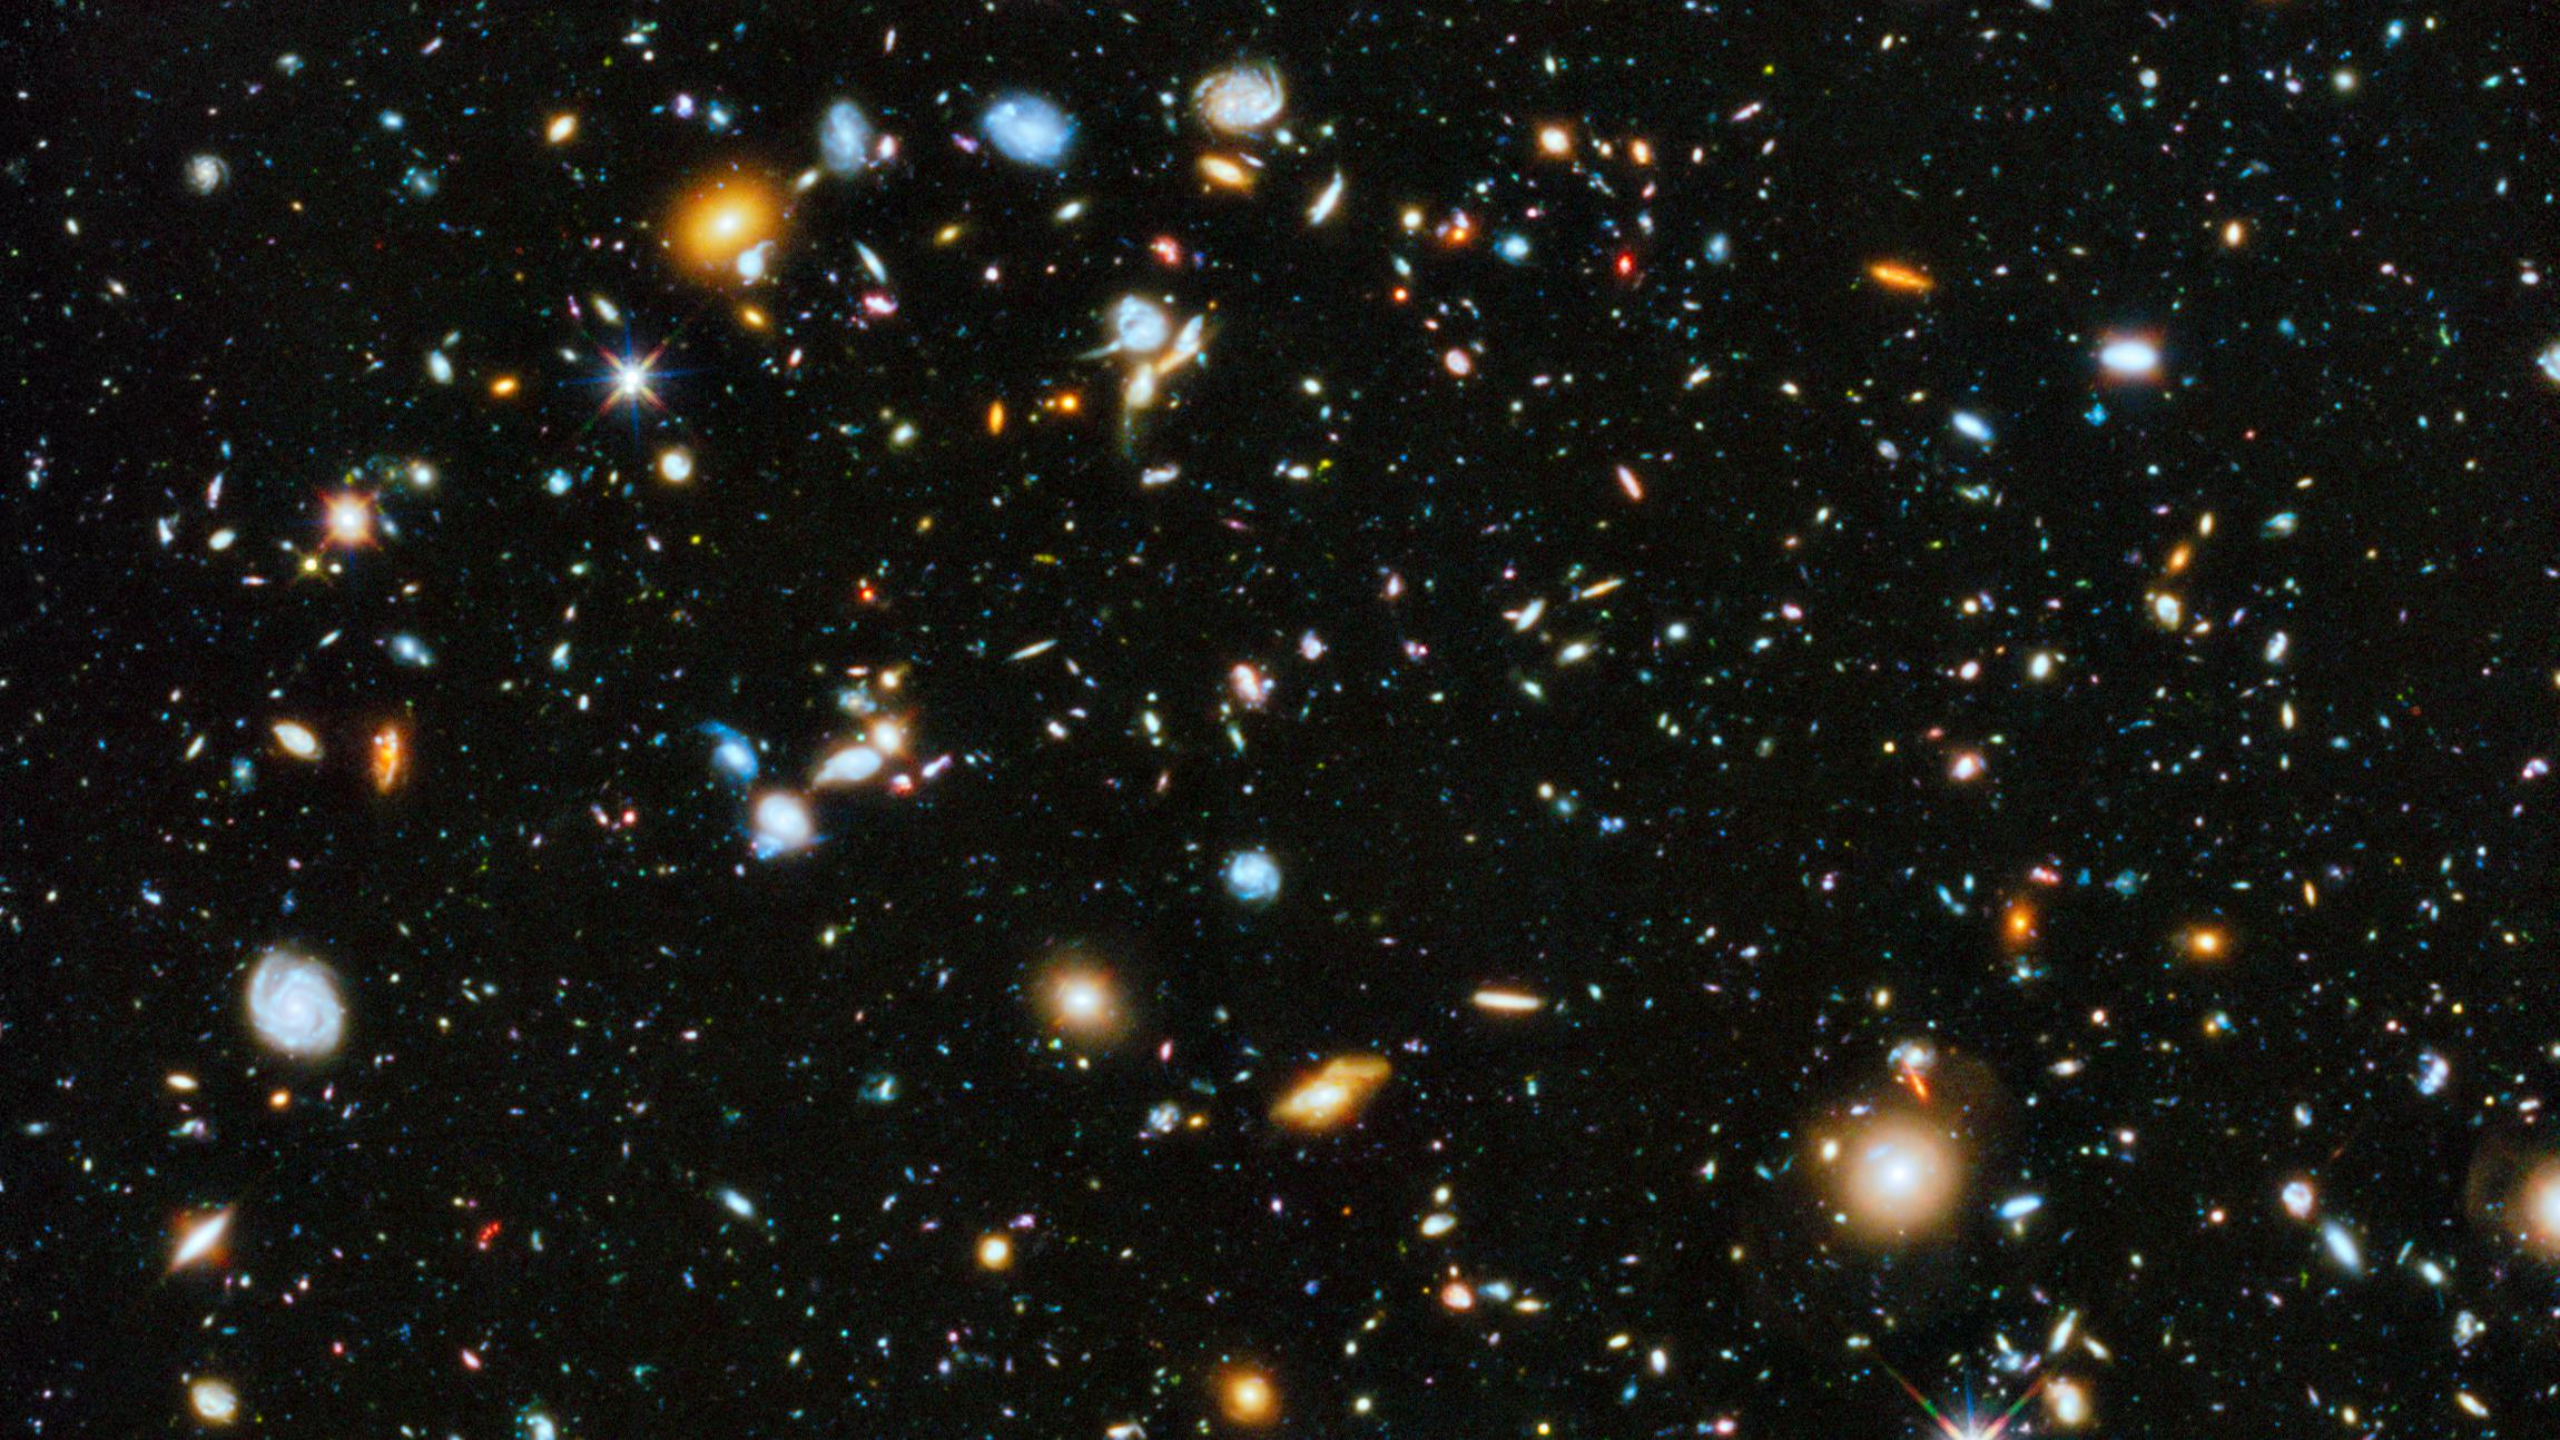


HUBBLE UNVEILS ITS MOST COLORFUL VIEW OF THE UNIVERSE
(ZOOM AND PAN)
BROWSE THE LINK BELOW AND PLAY!
https://hubblesite.org/contents/media/videos/2014/27/766-Video.html

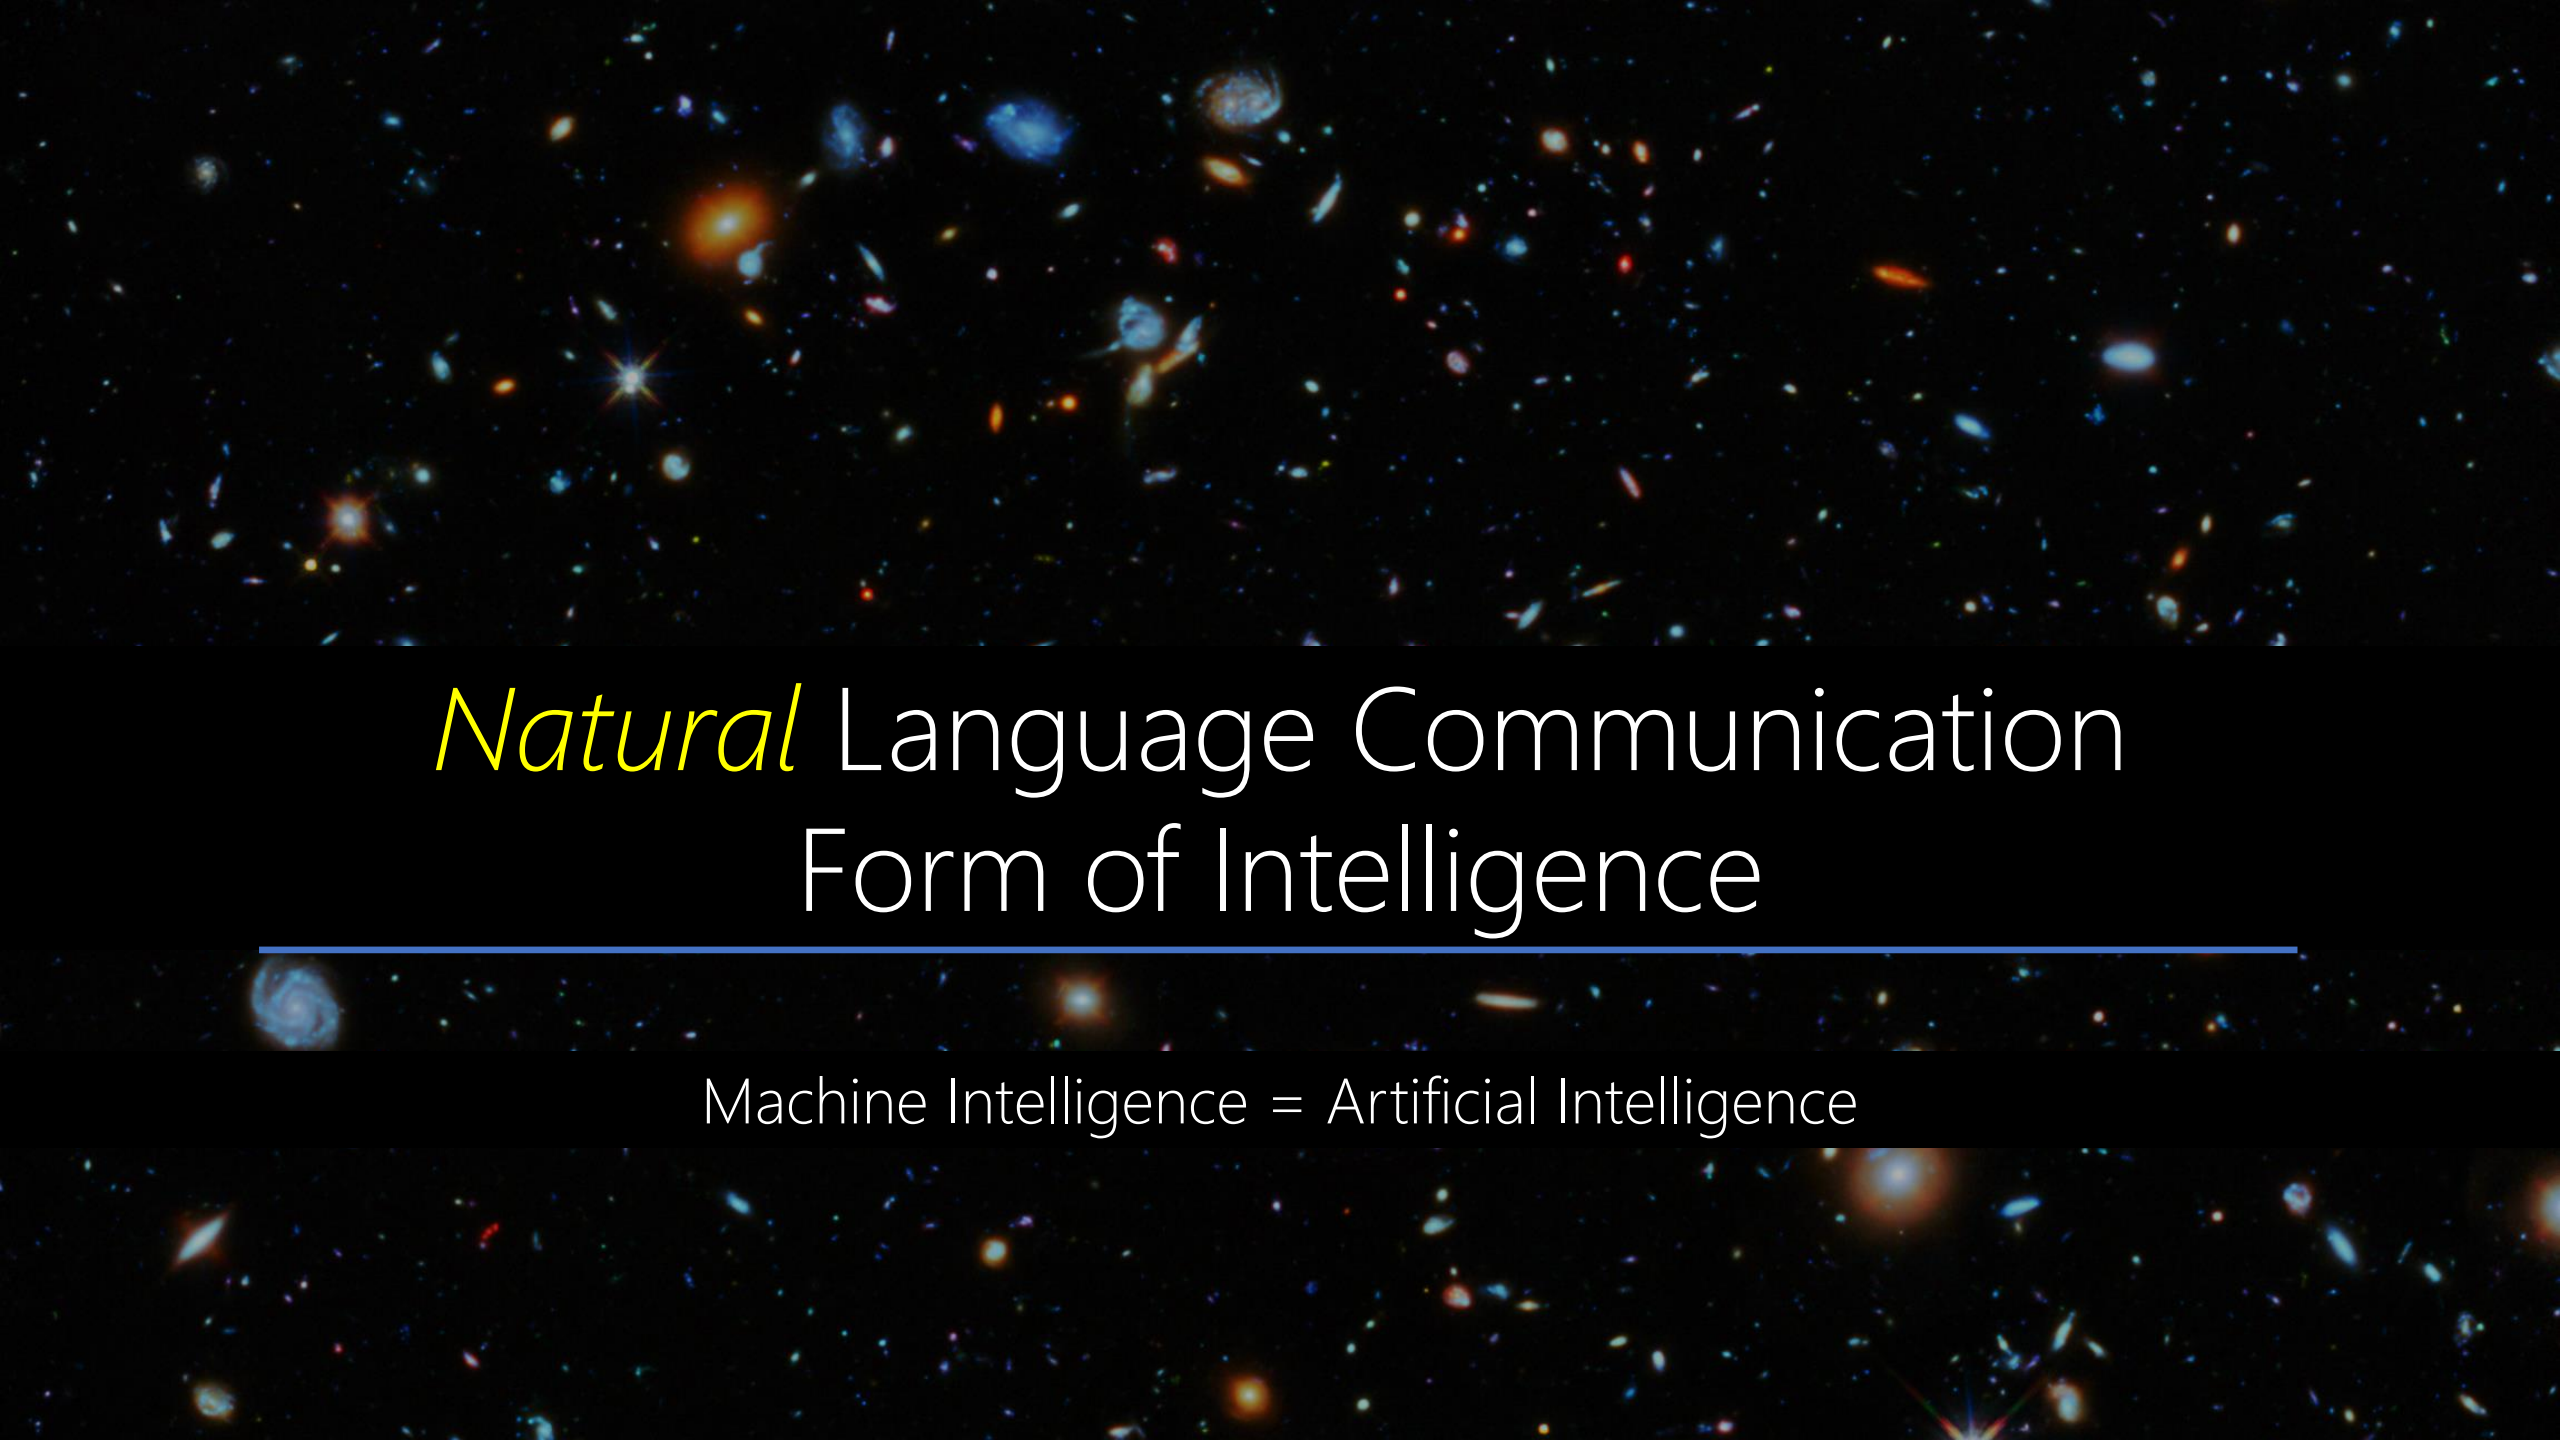

Natural Language Communication Form of Intelligence

Machine Intelligence = Artificial Intelligence

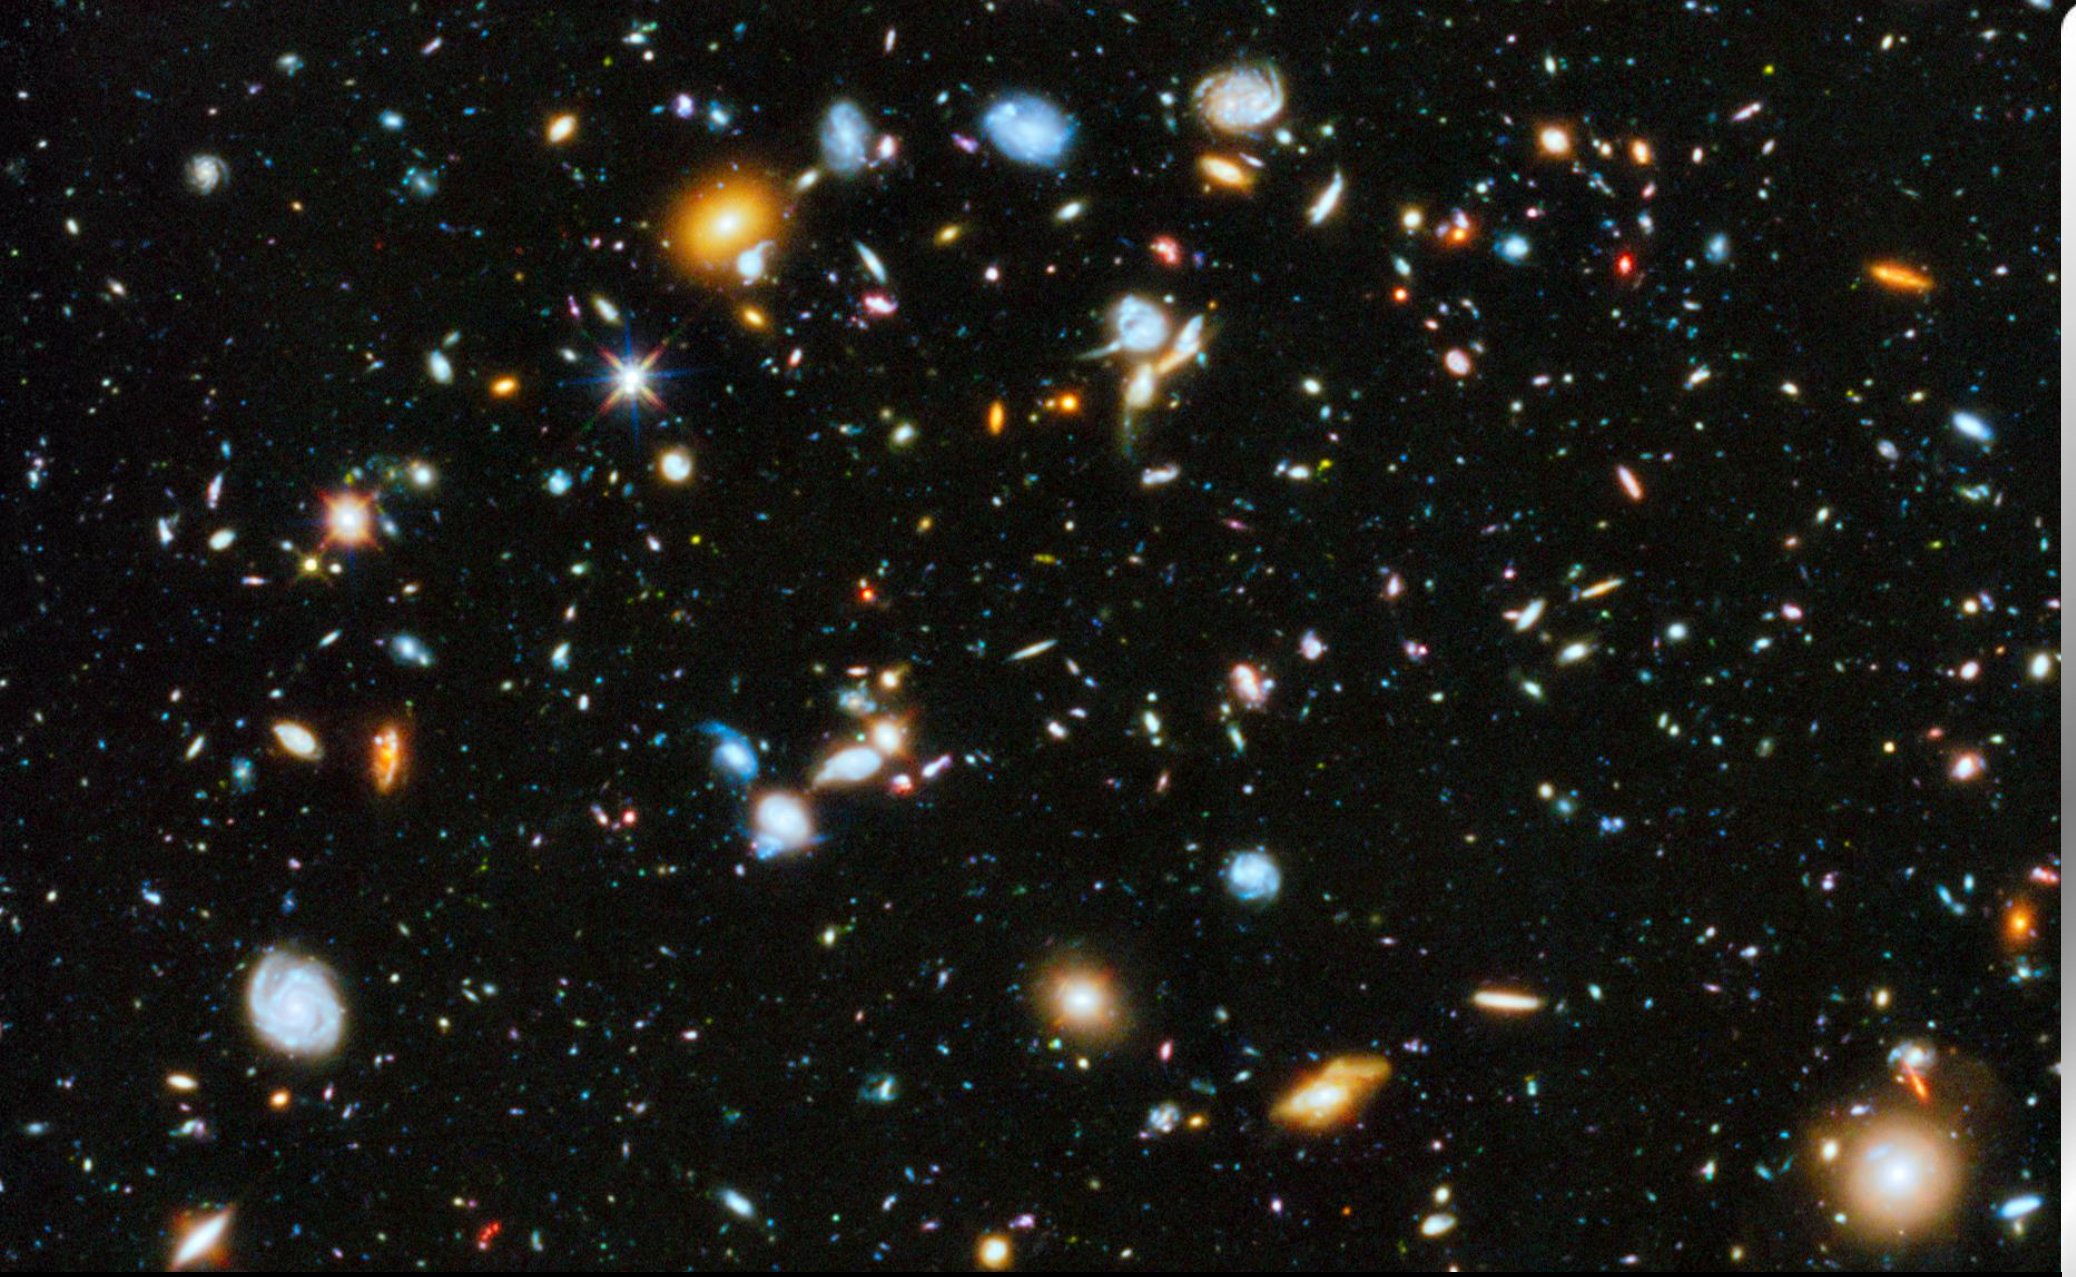

2001: A Space Odyssey (1968) - A Conversation with HAL!

https://www.youtube.com/watch?v=r13I-TuDcWI

The Turing Test

by Alan Turing in 1950

A test of a machine's ability to exhibit intelligent behavior equivalent to, or indistinguishable from, that of a human

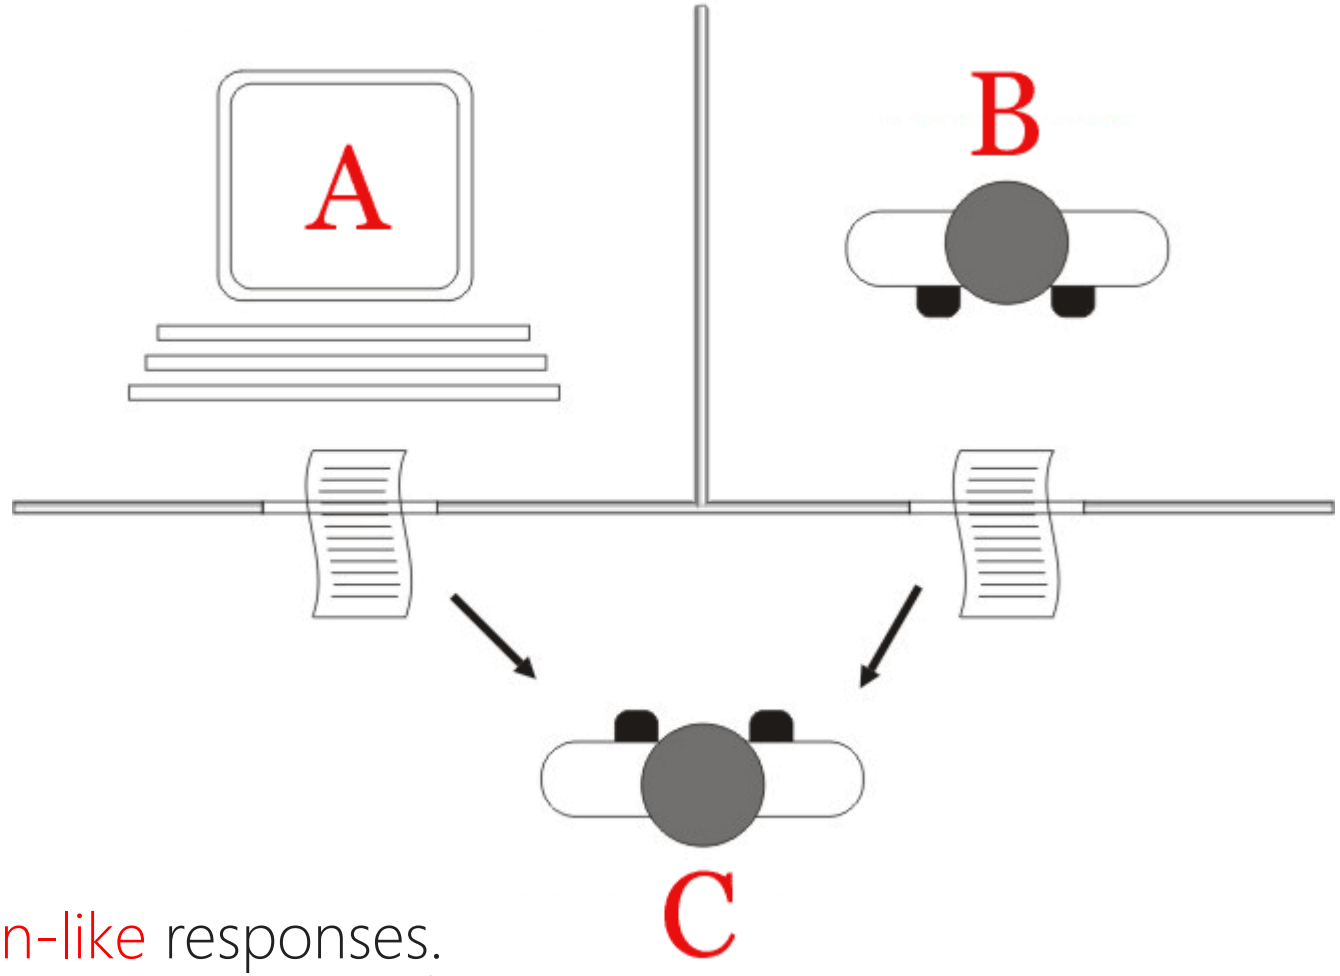

A machine designed to generate human-like responses.

Intelligence

The Turing Test

by Alan Turing in 1950

Human player C, the interrogator, is given the task of trying to determine which player – A or B – is a computer and which is a human. The interrogator is limited to using the responses to written questions to make the determination.

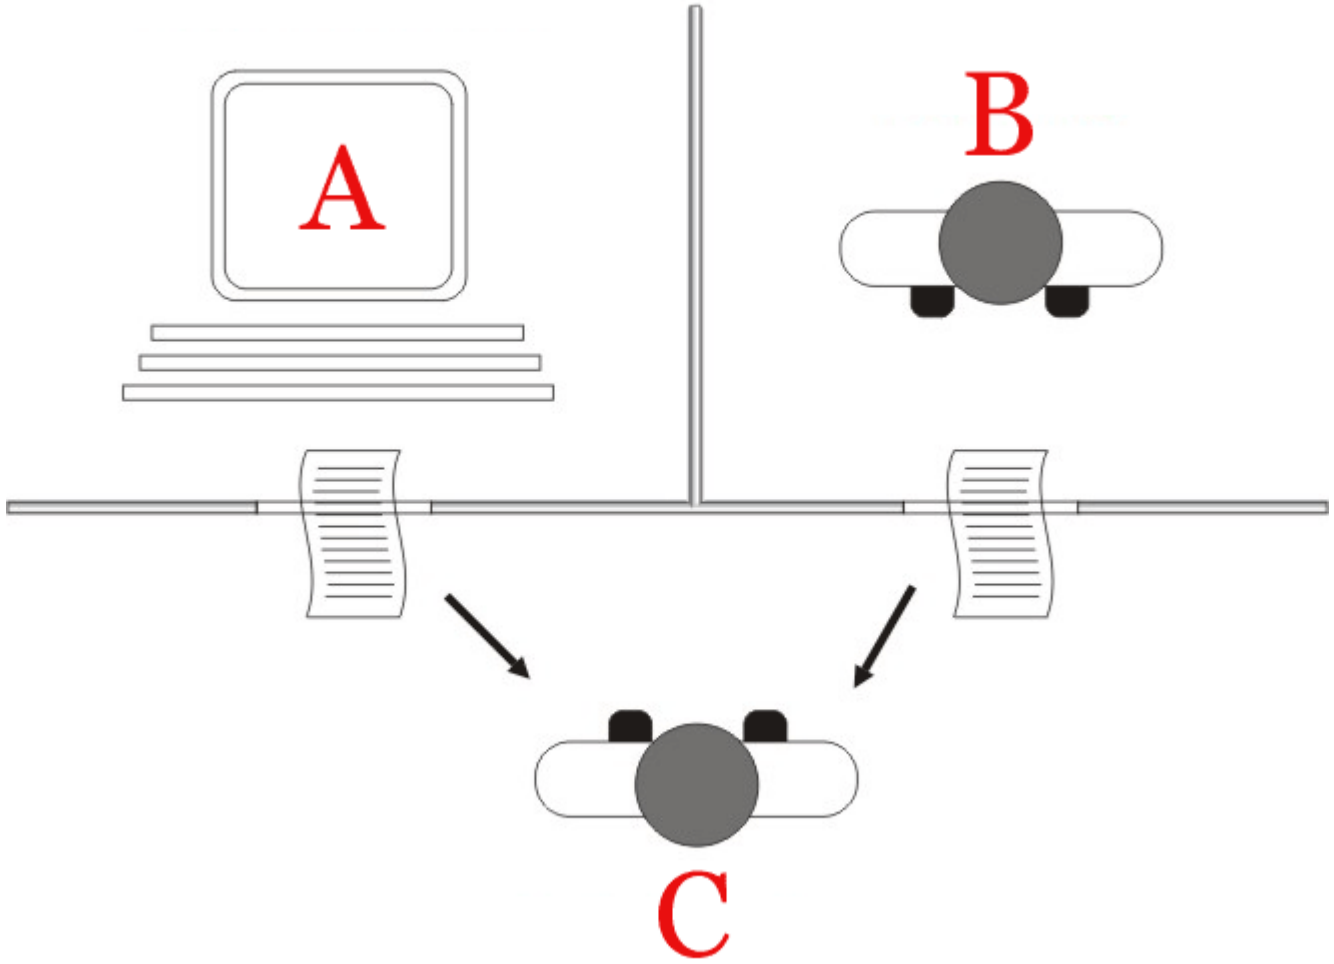

The Turing Test

by Alan Turing in 1950

Fast vs. Slow?
Speech vs. Text?
Wrong vs. Right?
Inaccurate vs. Precise?

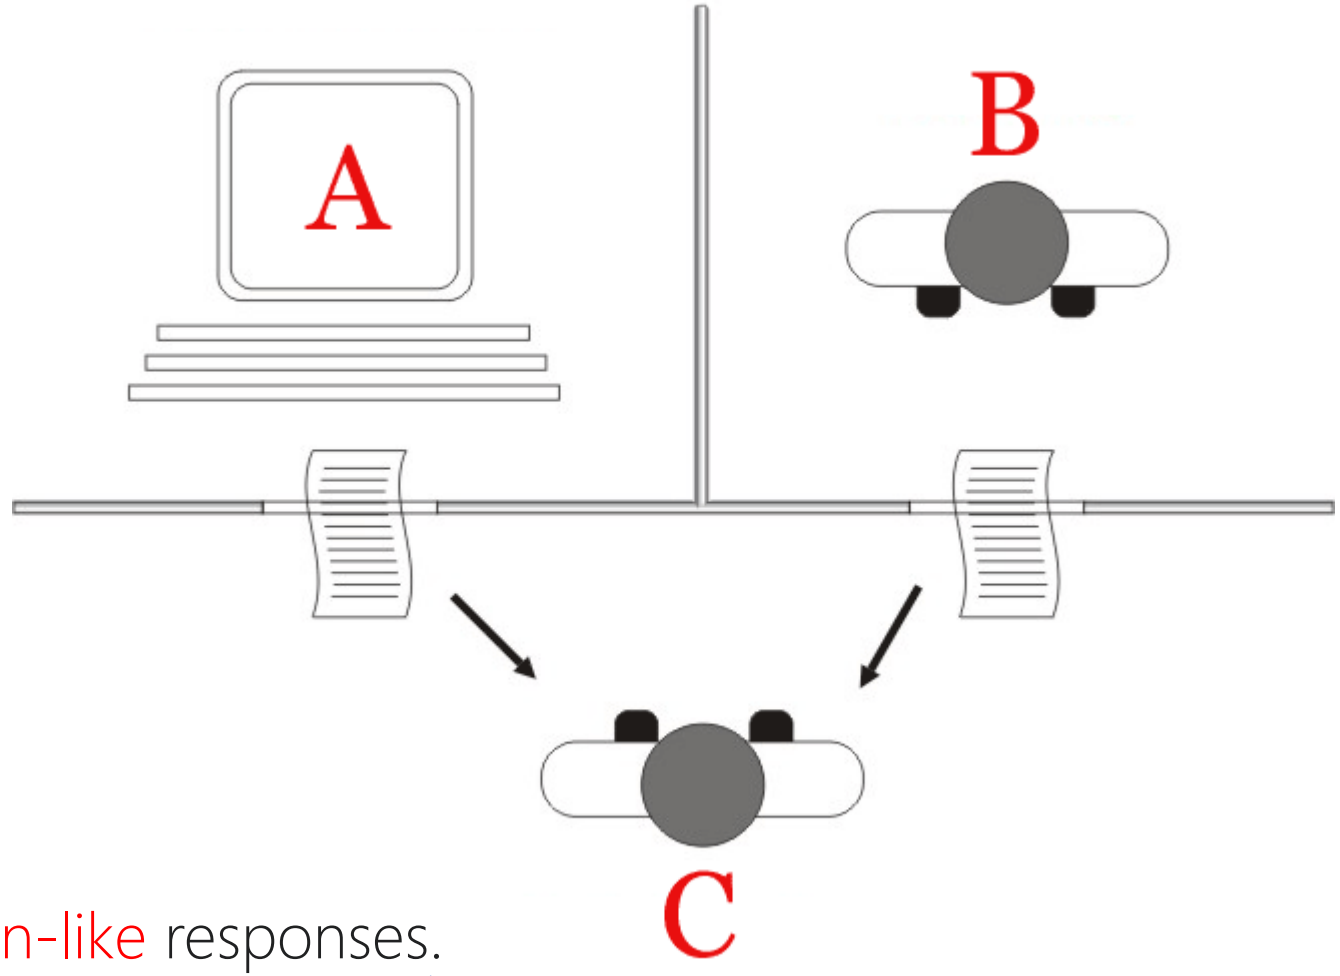

A machine designed to generate human-like responses.

Intelligence

The Turing Test

by Alan Turing in 1950

Audio Signals (Speech: Utterance)

- Talking
- Hearing

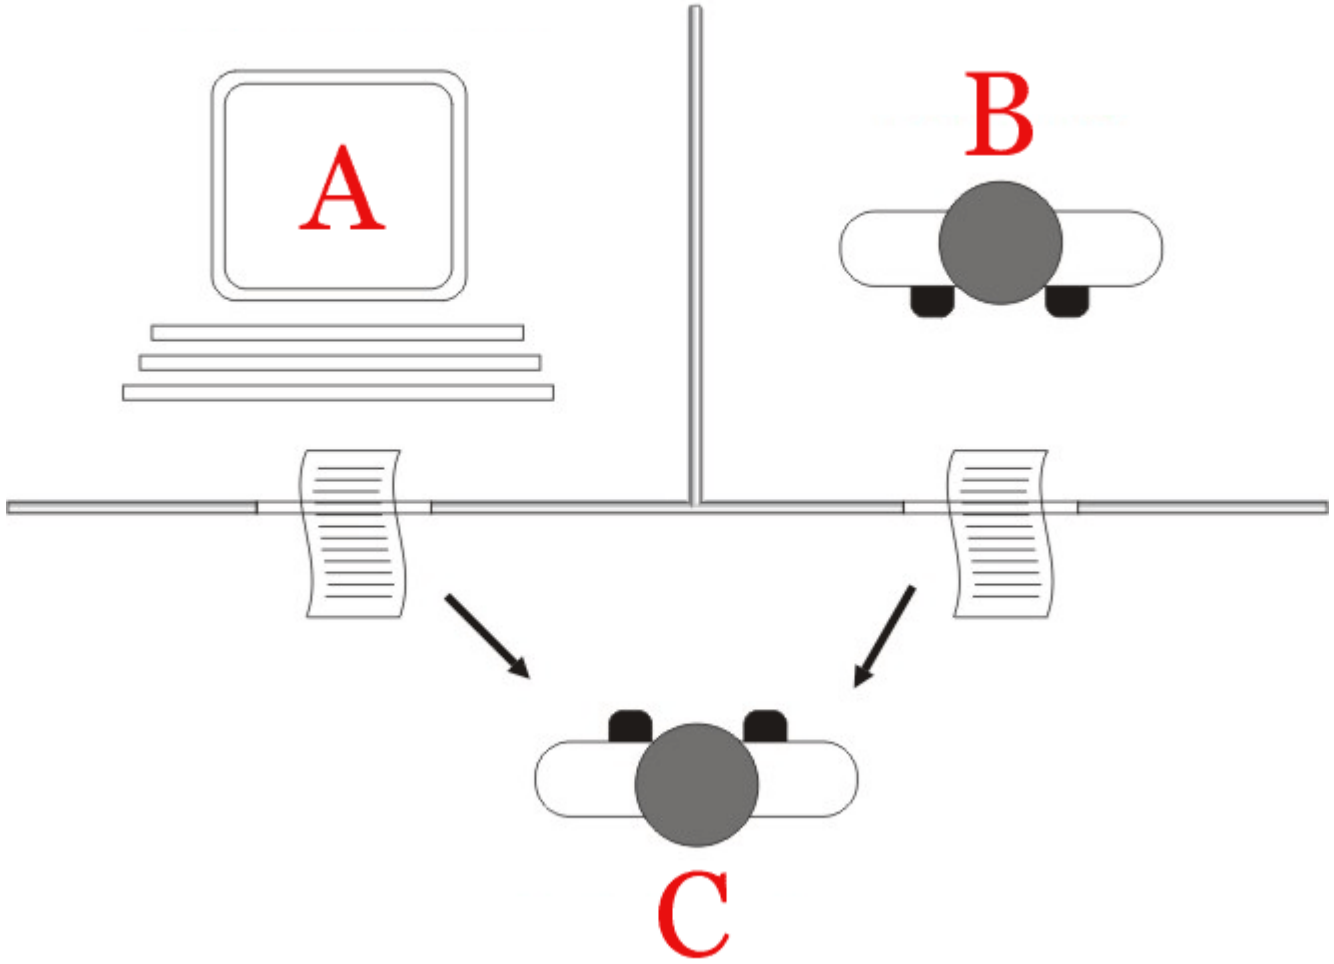

Sundar Pichai, Google's CEO

Google Duplex: AI Assistant Calls Local Businesses To Make Appointments

https://www.youtube.com/watch?v=D5VN56jQMWM

The Turing Test

by Alan Turing in 1950

Audio Signals
- Talking
- Hearing

Visual Signals
- Emotion
- Gesture (Body Language)

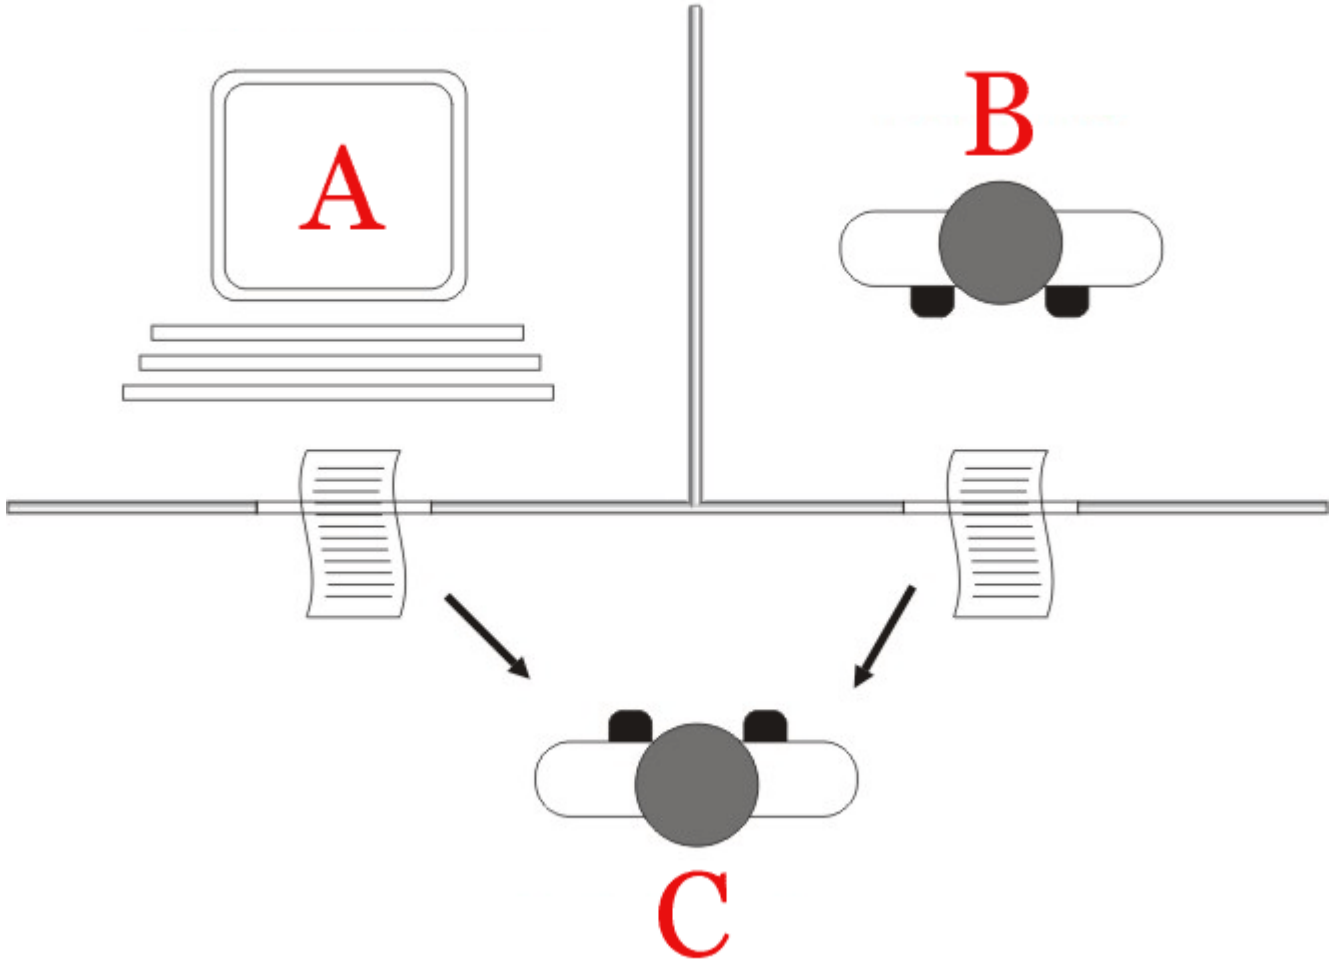

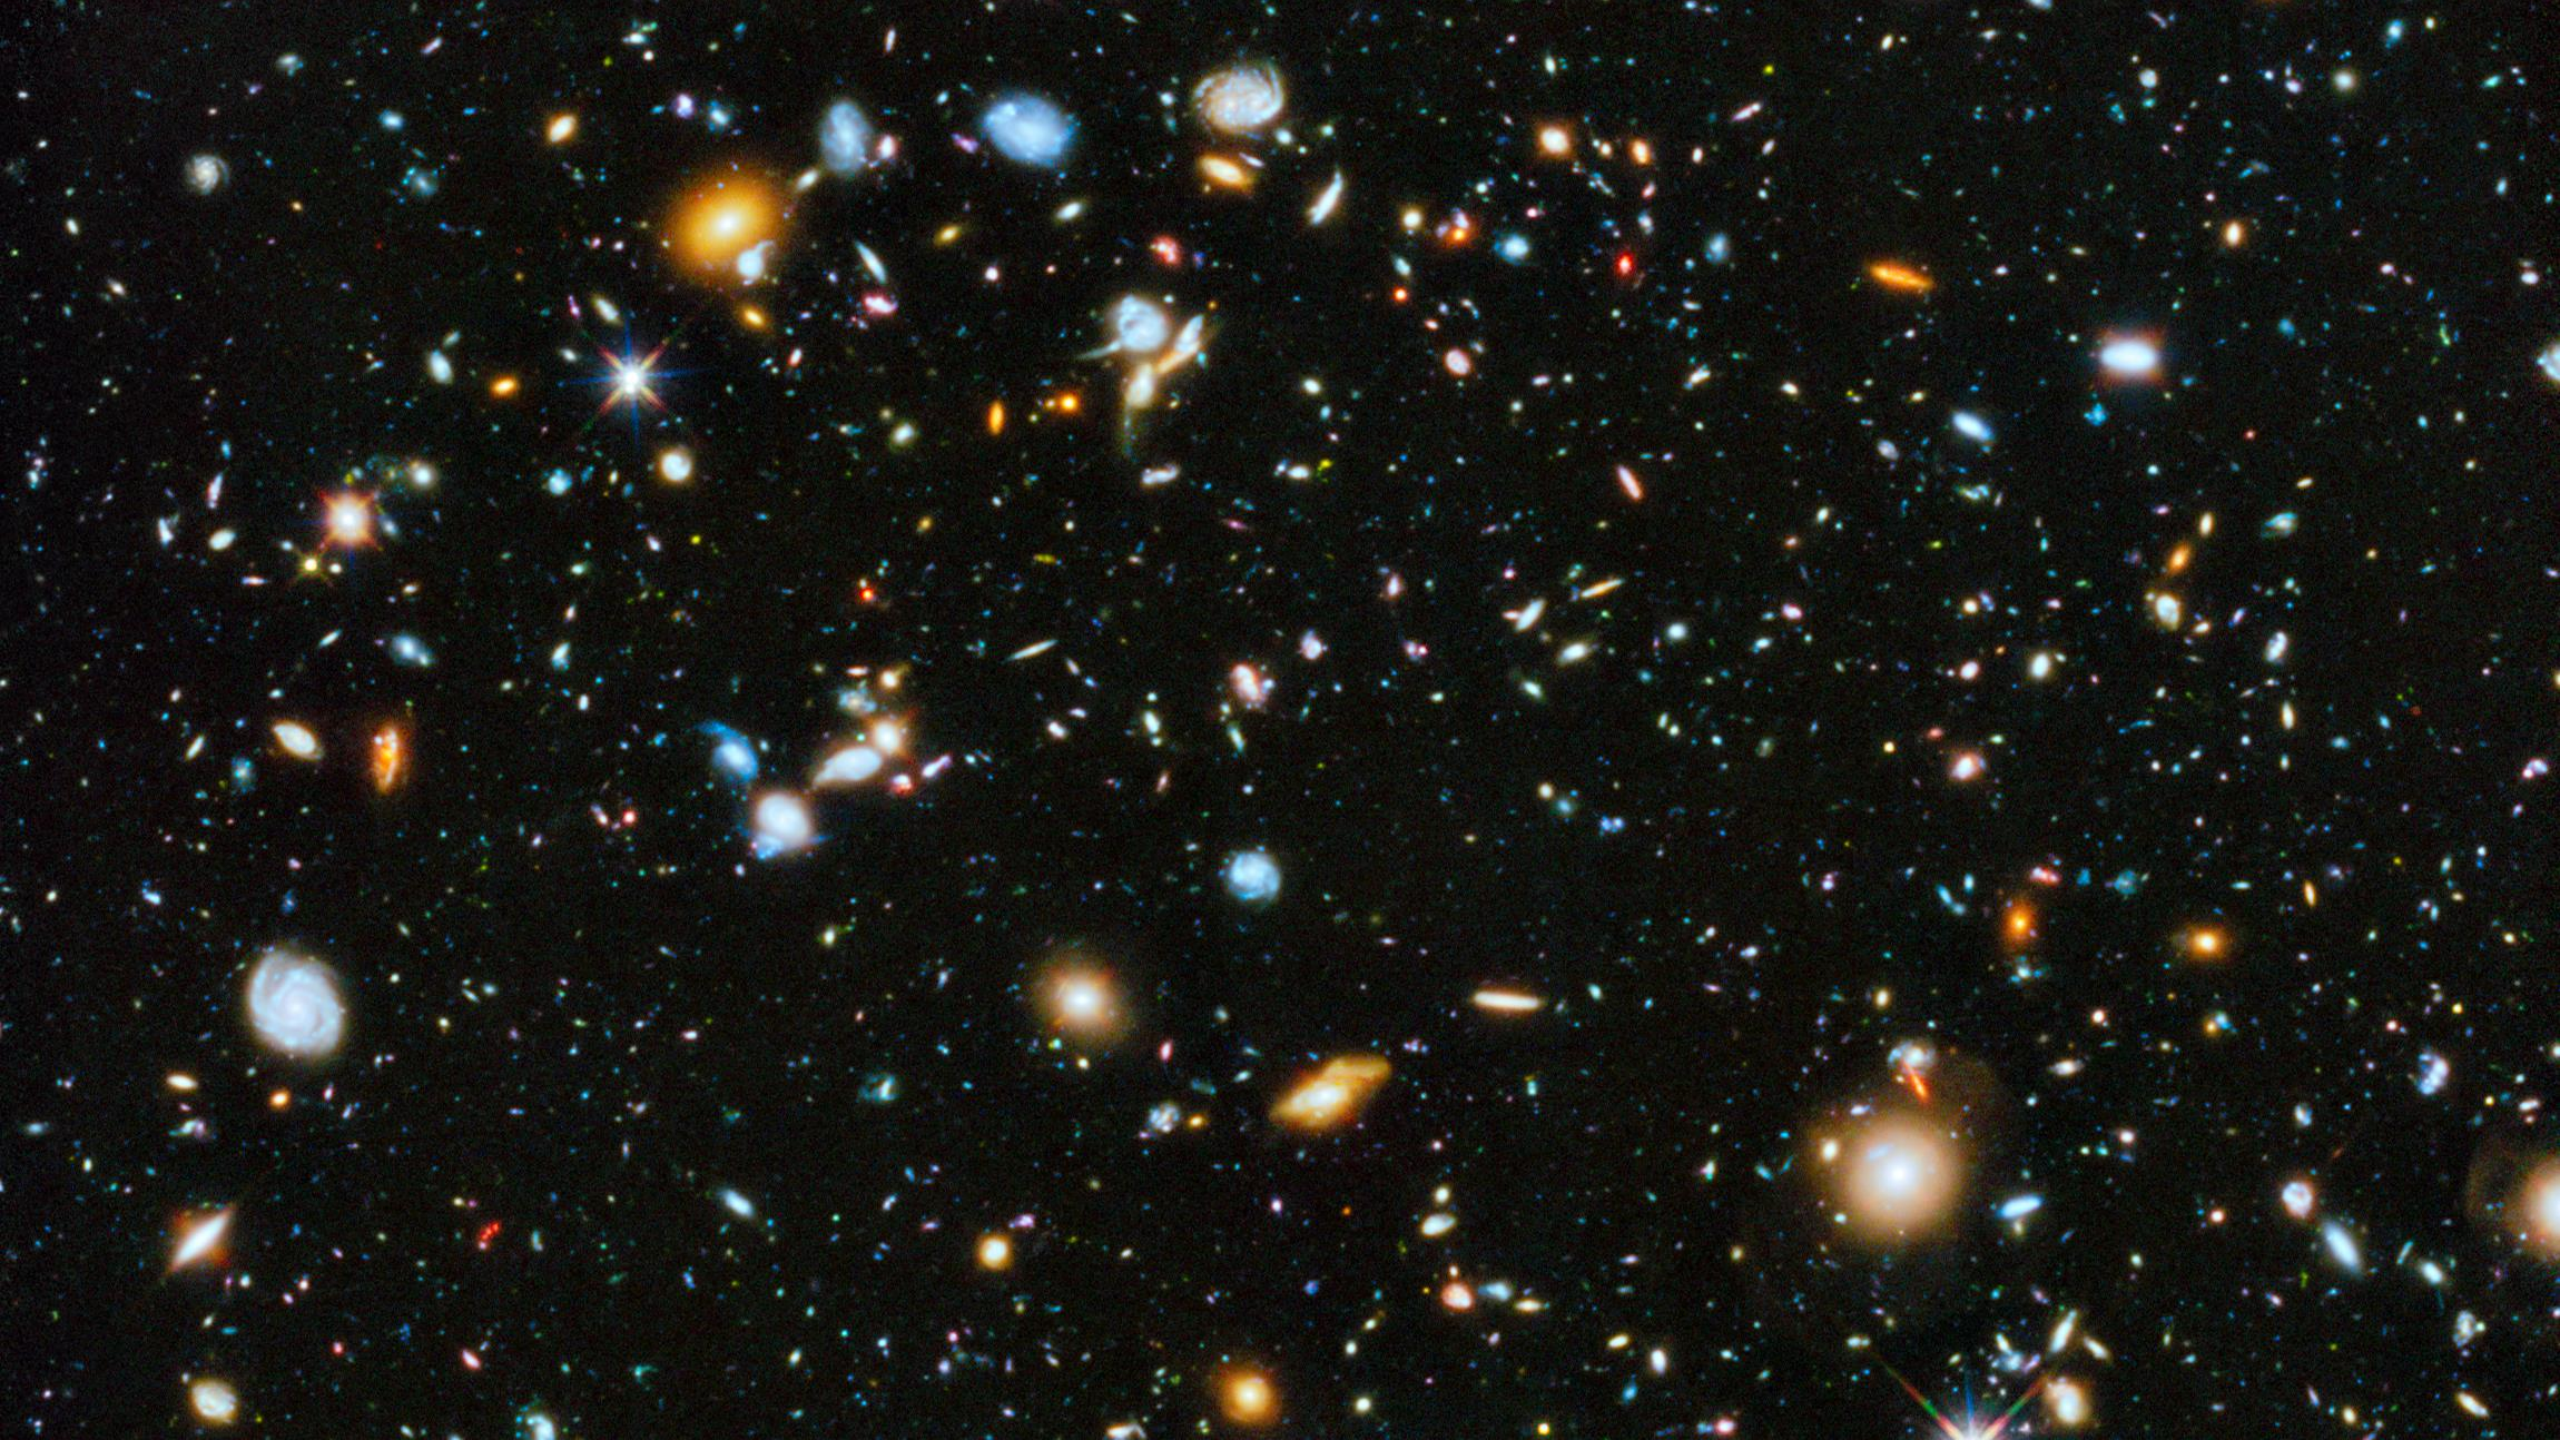

2001: A Space Odyssey (1968) - HAL Reads Lips

https://www.youtube.com/watch?v=XDO8OYnmkNY

The Turing Test

by Alan Turing in 1950

Audio Signals
- Talking
- Hearing

Visual Signals
- Emotion
- Gesture (Body Language)
- Lips Reading

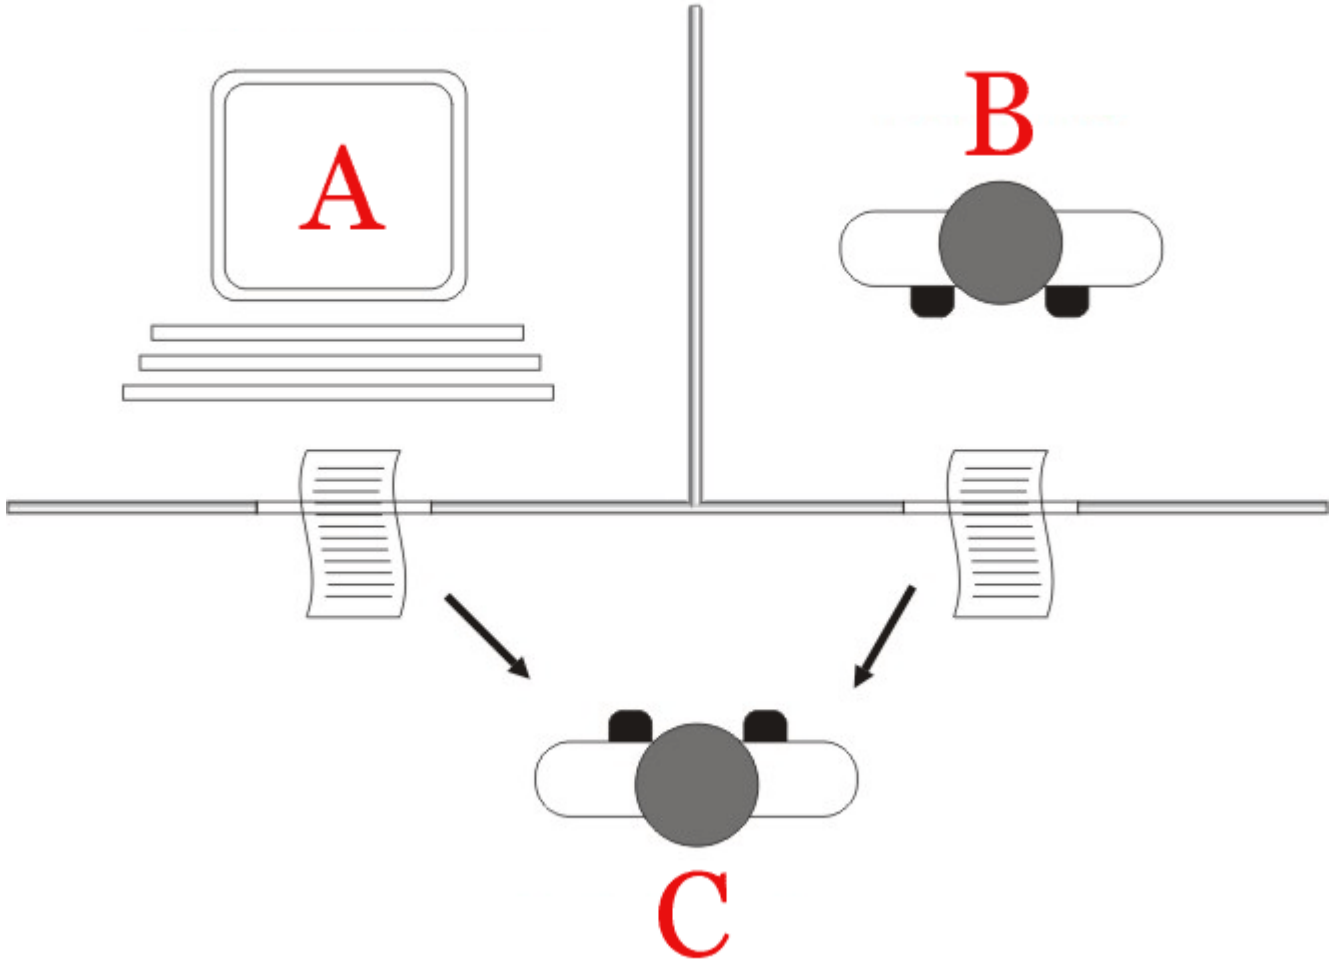

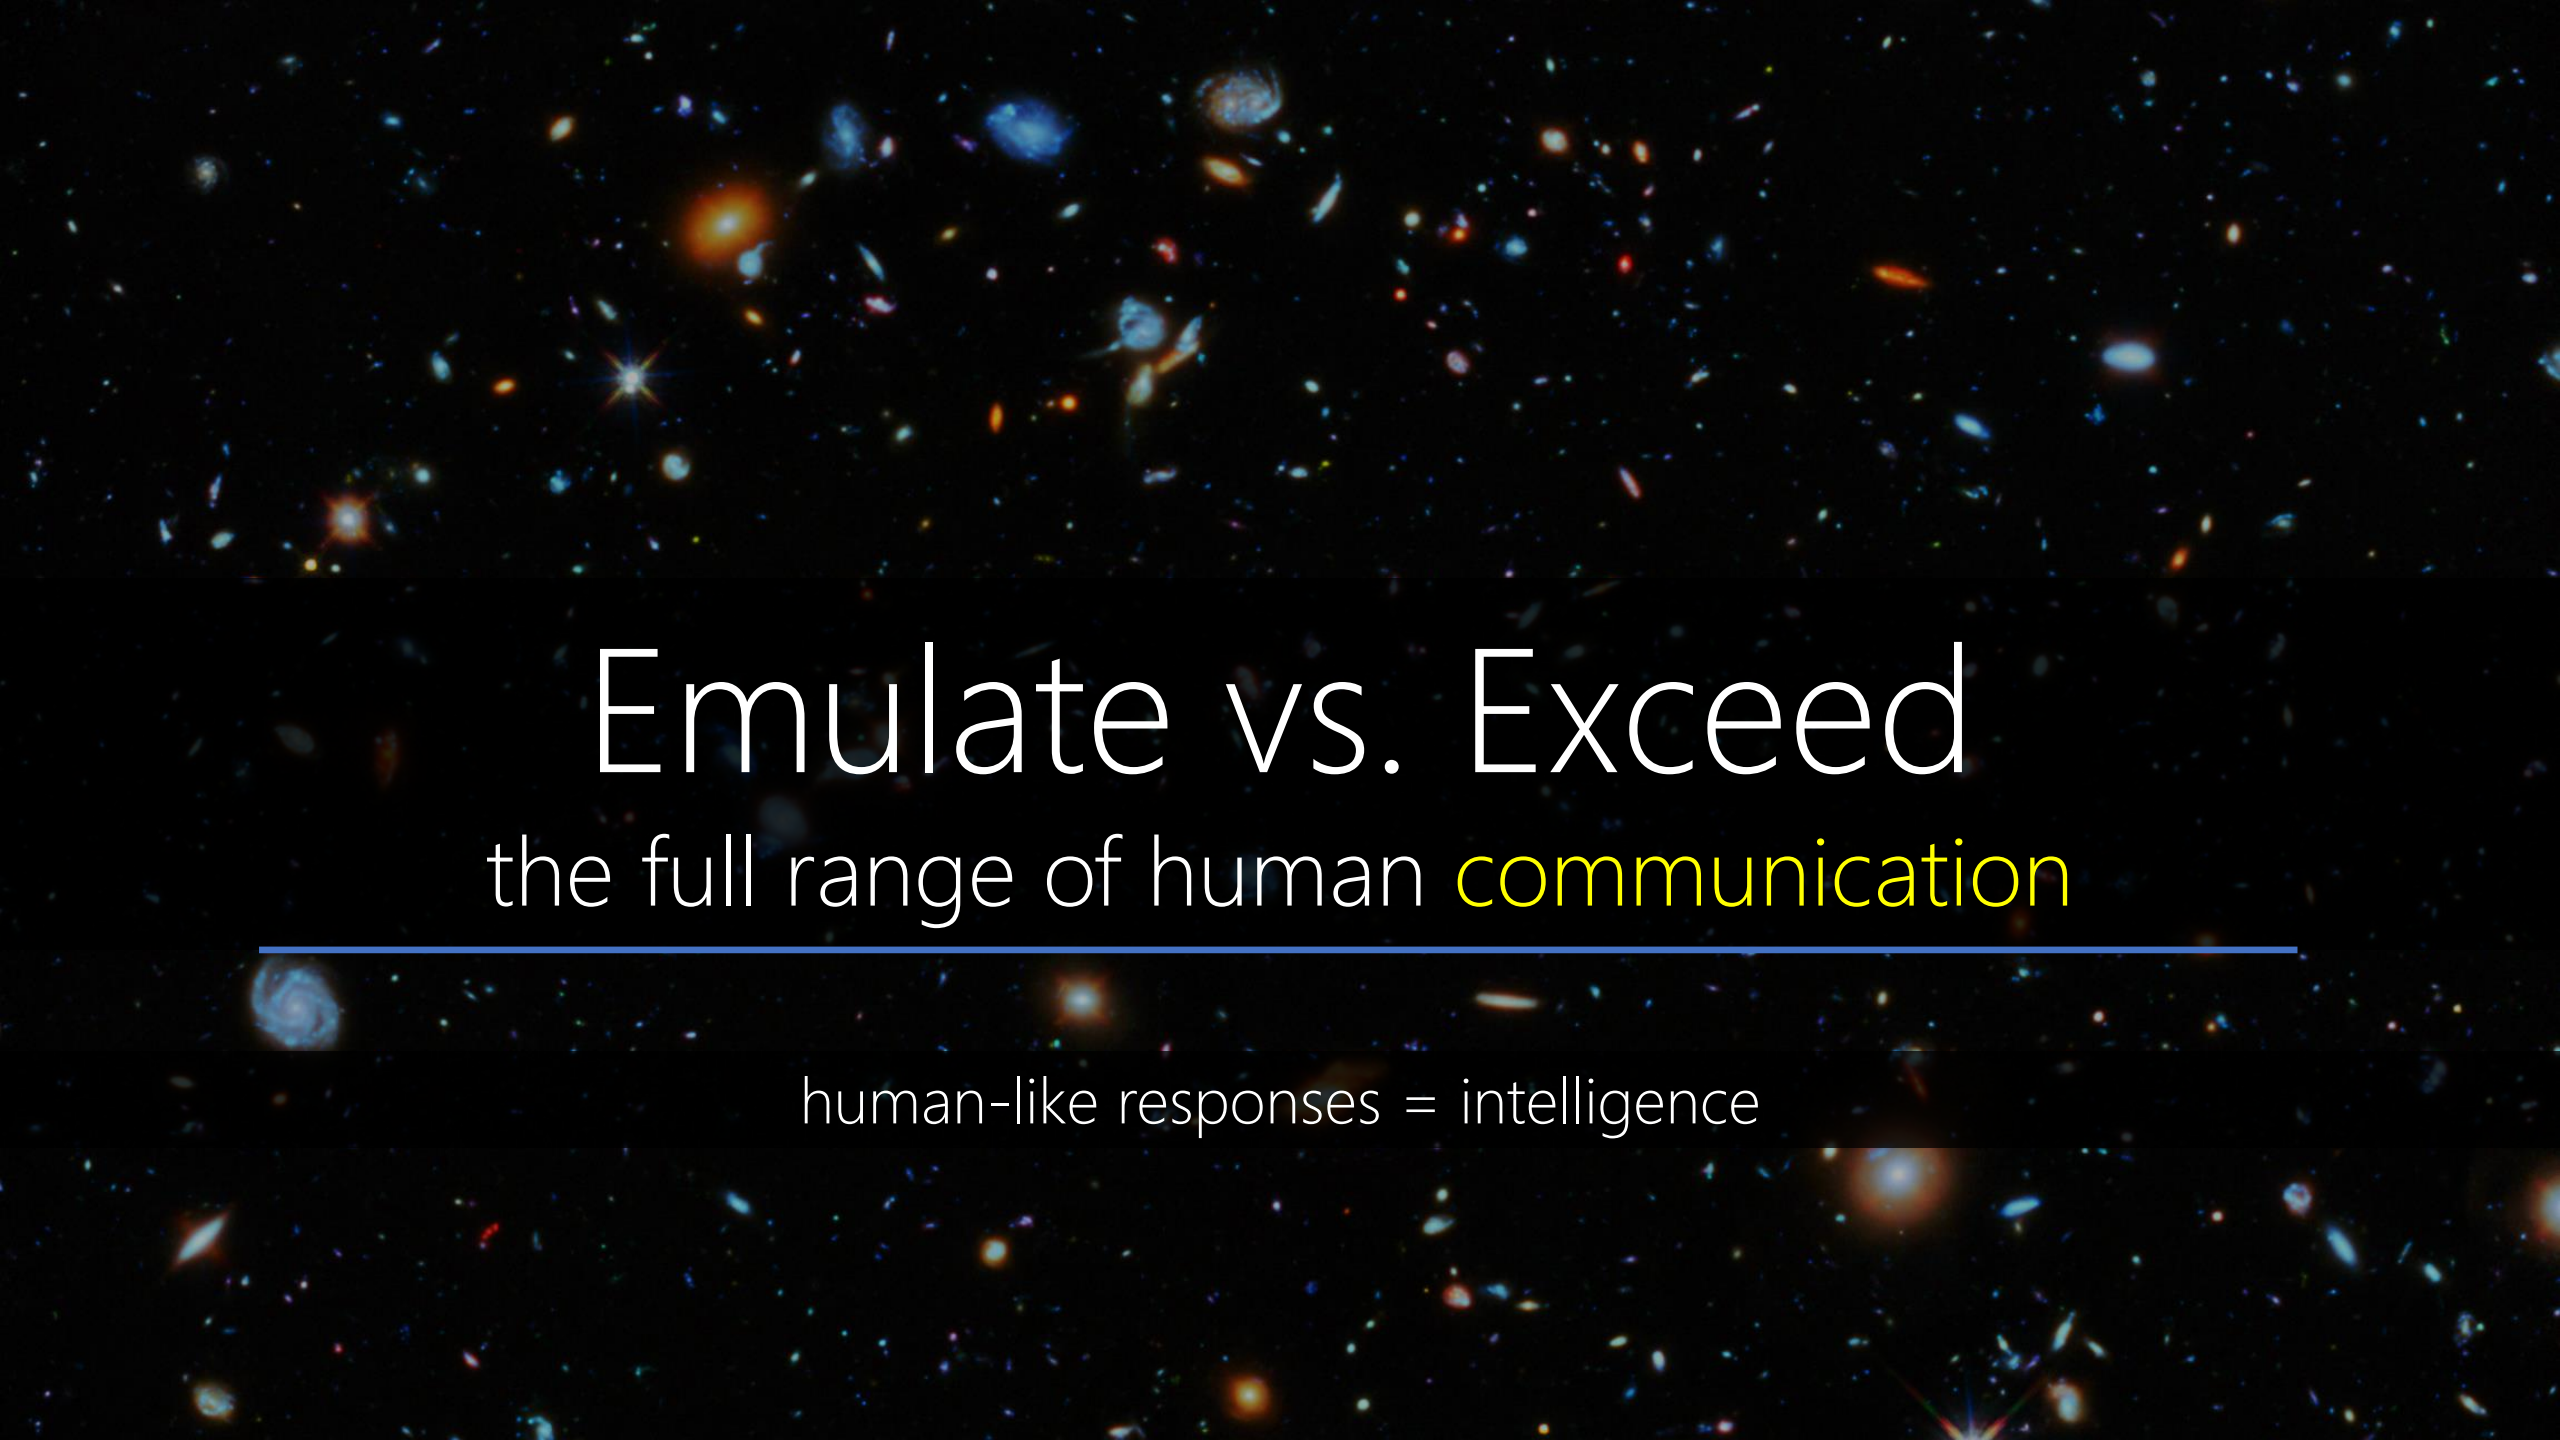

Emulate vs. Exceed

the full range of human communication

human-like responses = intelligence

Natural Language Communication
Speech, Text, Emotion, …

Dialog: Understand and Respond

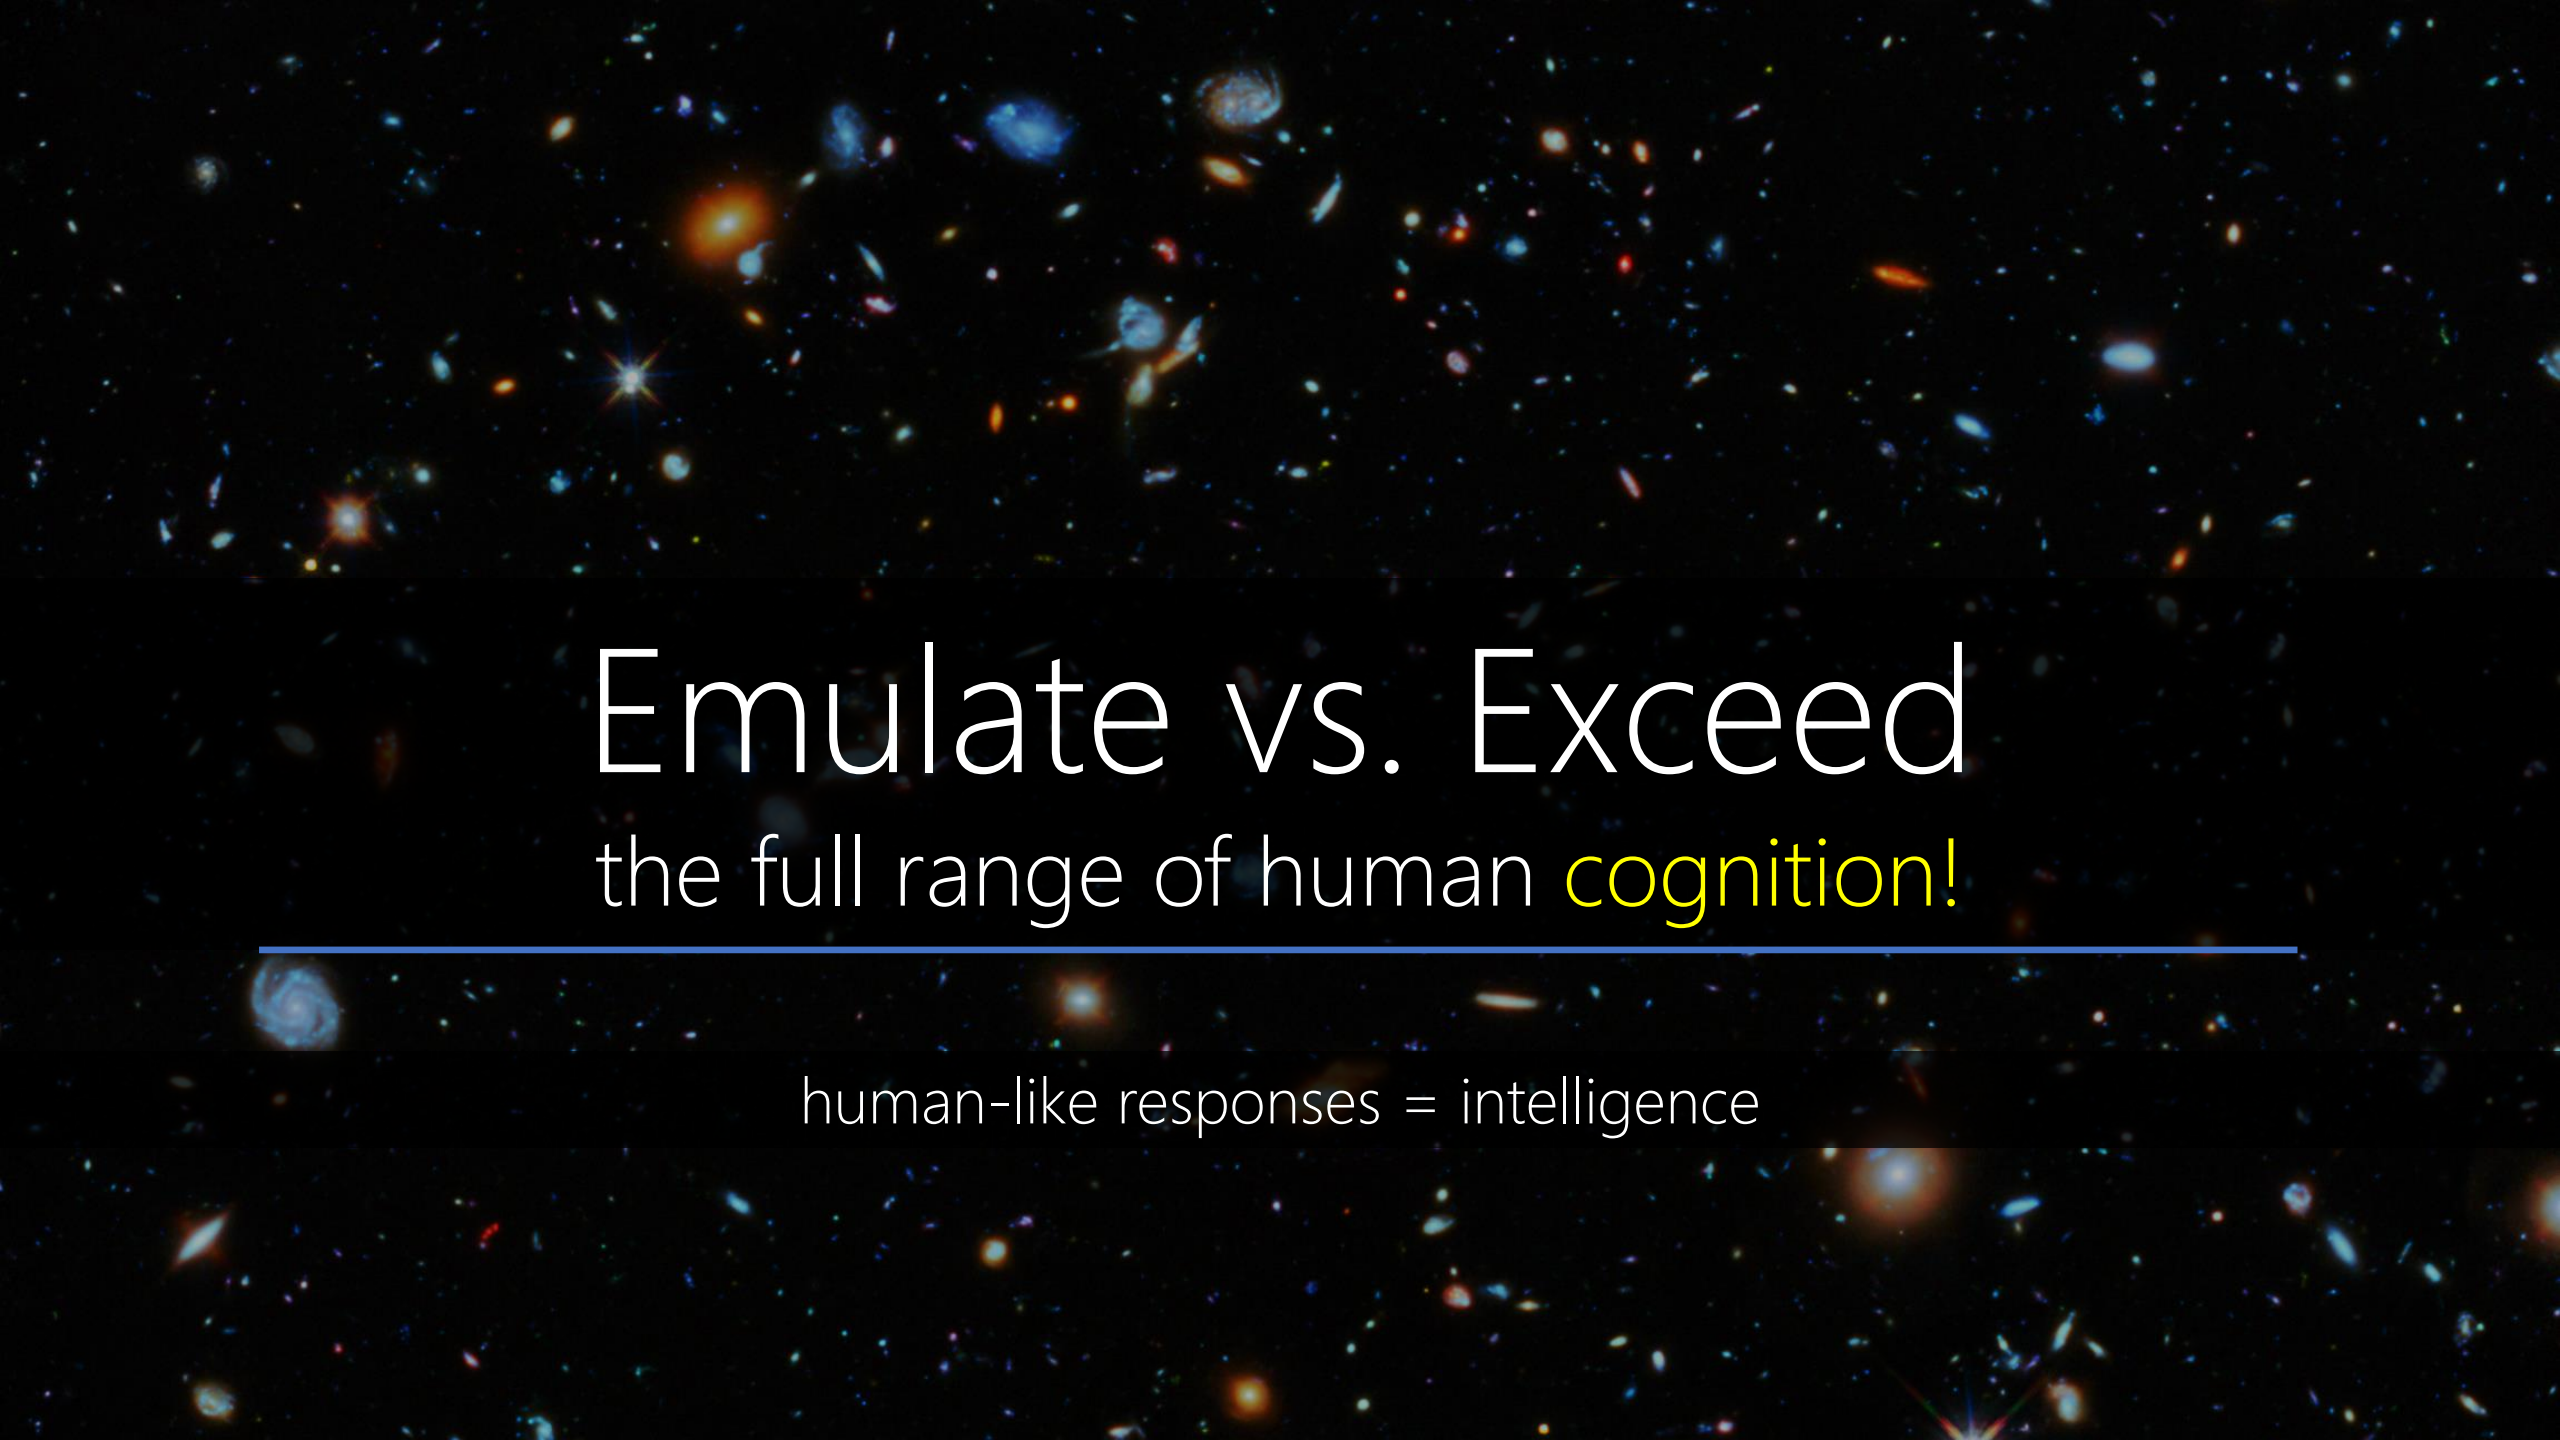
Emulate vs. Exceed
the full range of human cognition!
human-like responses = intelligence

2001: A Space Odyssey (1968) – HAL: I'm Sorry, Dave!

https://www.youtube.com/watch?v=Wy4EfdnMZ5g

Tweet

Yann LeCun
@ylecun

Language is an imperfect, incomplete, and low-bandwidth serialization protocol for the internal data structures we call thoughts.

9:36 AM · Mar 6, 2021 · Twitter for Android

Despite its many practical applications, language is perhaps number 300 in the priority list for AI research. It would be a great achievement if AI could attain the capabilities of an orangutan, which do not include language!
- Yann LeCun (computer vision researcher)

Ryan Abernathey @rabernat · Mar 6

Replying to @ylecun

Disagree. Thought as we know it would be impossible without language.

17 3 72

Yann LeCun @ylecun · Mar 6

The vast majority of human knowledge, skills, and thoughts are not verbalizable.

11 5 94

Show replies

Handle is Boston Dynamics' newest design. It can jump four feet in the air and zip around at nine miles per hour.
https://www.youtube.com/watch?v=7h8mX9ZMs7g

- The Matrix (1999), Lana & Lilly Wachowski

Research Priorities for Artificial Intelligence

State of AI Report

October 1, 2020

Nathan Benaich Ian Hogarth

About the authors

Nathan Benaich

Nathan is the General Partner of **Air Street Capital**, a venture capital firm investing in AI-first technology and life science companies. He founded RAAIS and London.AI, which connect AI practitioners from large companies, startups and academia, and the RAAIS Foundation that funds open-source AI projects. He studied biology at Williams College and earned a PhD from Cambridge in cancer research.

Ian Hogarth

Ian is an **angel investor** in 60+ startups. He is a Visiting Professor at UCL working with Professor Mariana Mazzucato. Ian was co-founder and CEO of Songkick, the concert service used by 17M music fans each month. He studied engineering at Cambridge where his Masters project was a computer vision system to classify breast cancer biopsy images. He is the Chair of Phasecraft, a quantum software company.

Language models: Welcome to the Billion Parameter club

Huge models, large companies and massive training costs dominate the hottest area of AI today, NLP.

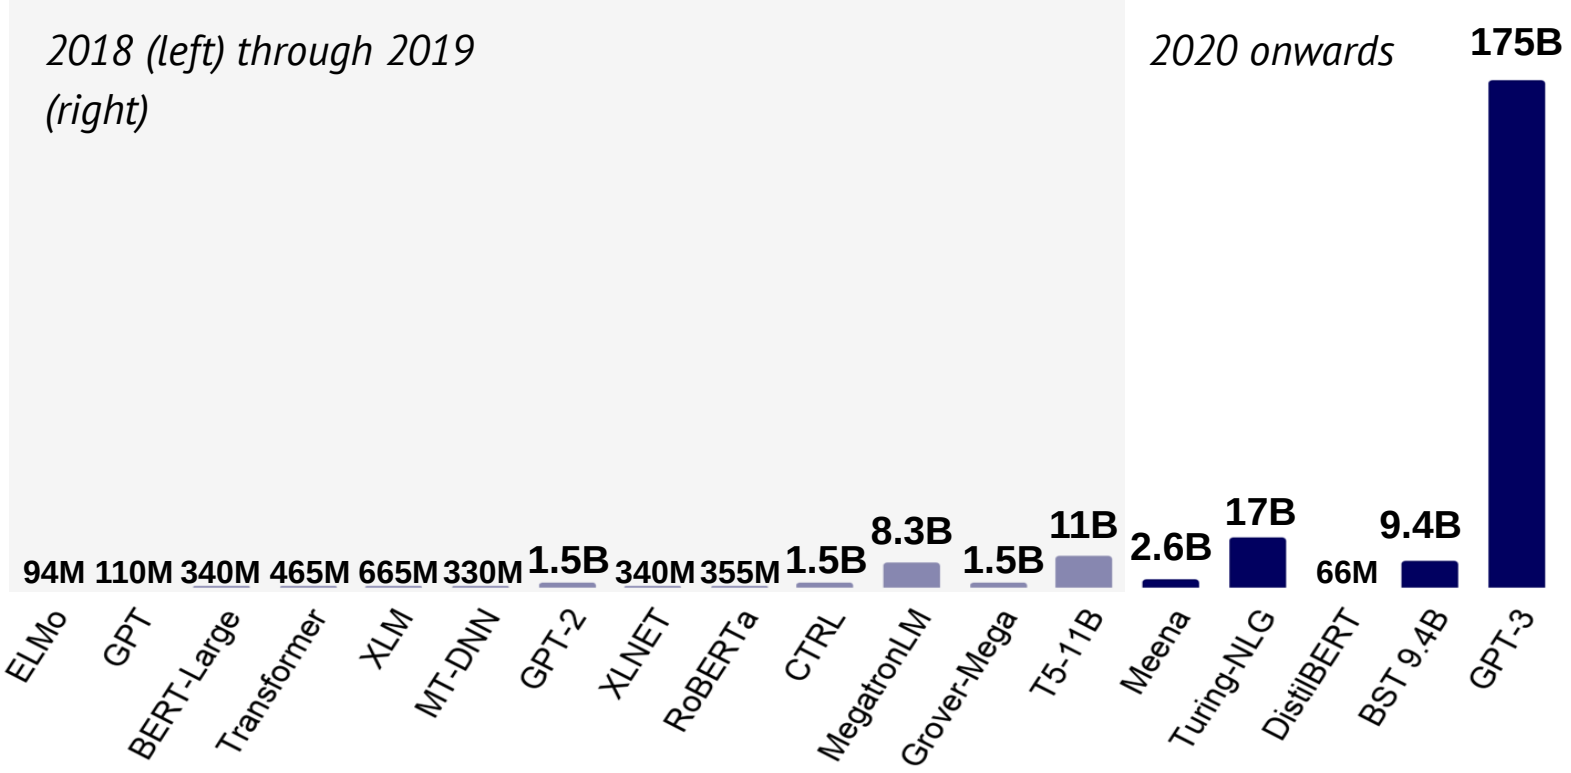

Note: The number of parameters indicates how many different coefficients the algorithm optimizes during the training process.

NLP benchmarks take a beating: Over a dozen teams outrank the human GLUE baseline

It was only 12 months ago that the human GLUE benchmark was beat by 1 point. Now SuperGLUE is in sight.

- GLUE (General Language Understanding Evaluation) and it's more challenging sibling SuperGLUE are benchmarks that evaluate NLP systems at a range of tasks spanning logic, common sense understanding, and lexical semantics. The human benchmark on GLUE is reliably beat today (right) and the SuperGLUE human benchmark is almost surpassed too!

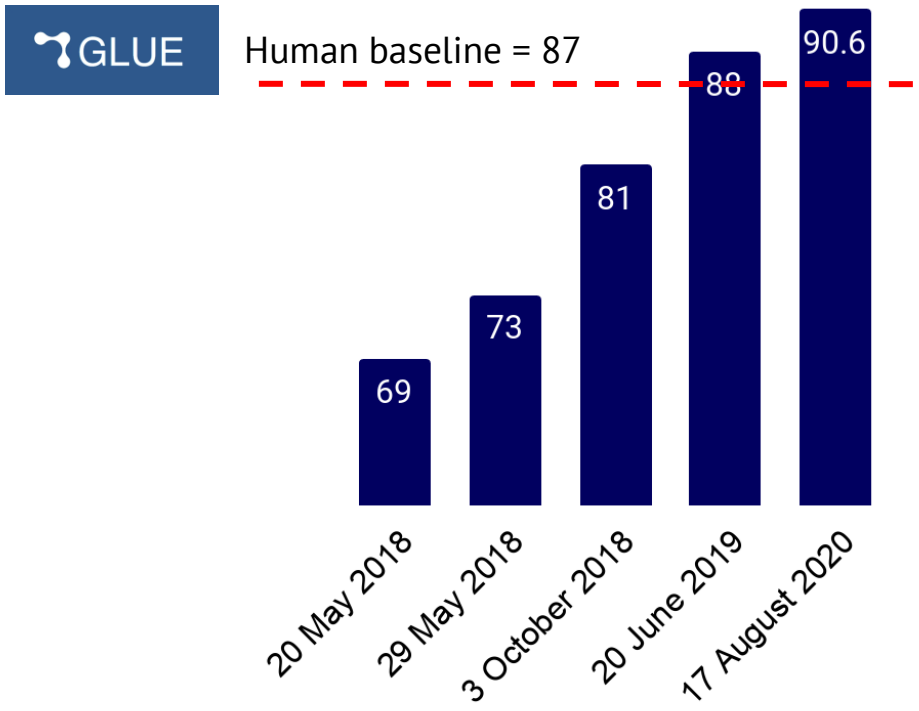

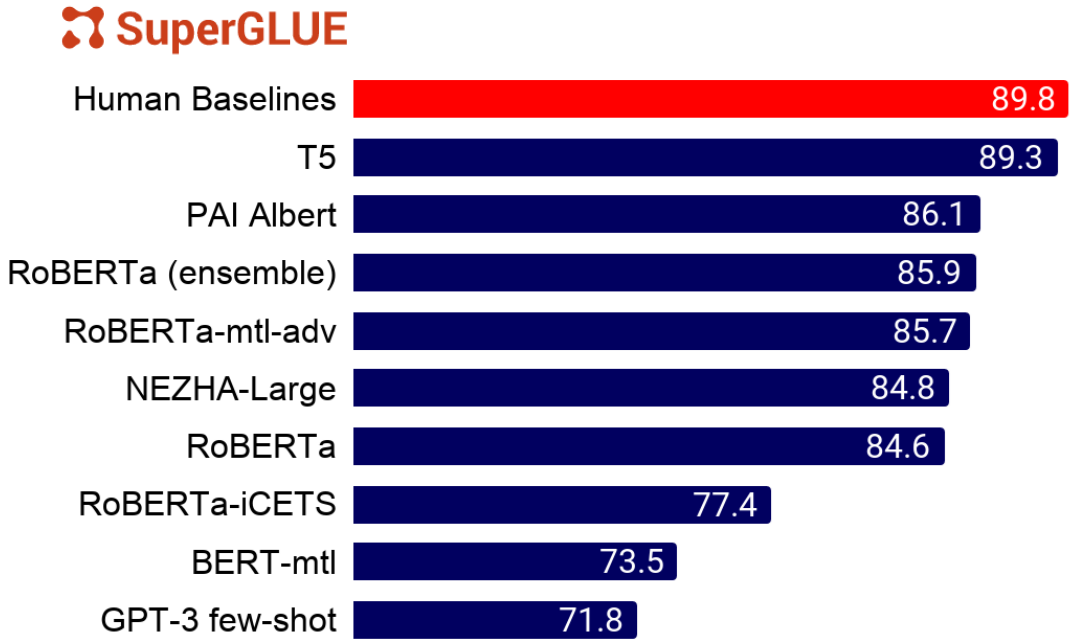

A new generation of transformer language models are unlocking new NLP use-cases

GPT-3, T5, BART are driving a drastic improvement in the performance of transformer models for text-to-text tasks like translation, summarization, text generation, text to code.

Code generation and more: gpt3examples.com

Computer, please convert my code into another programming language

An unsupervised machine translation model trained on GitHub projects with 1,000 parallel functions can translate 90% of these functions from C++ to Java and 57% of Python functions into C++ and successfully pass unit tests. No expert knowledge required, but no guarantees that the model didn't memorize the functions either.

Python input

```python
def SumOfKsubArray(arr, n, k):
  Sum = 0
  S = deque()
  G = deque()
  for i in range(k):
    while (len(S) > 0 and arr[S[-1]] >= arr[i]):
      S.pop()
    while (len(G) > 0 and arr[G[-1]] <= arr[i]):
      G.pop()
    G.append(i)
    S.append(i)
  for i in range(k, n):
    Sum += arr[S[0]] + arr[G[0]]
    while (len(S) > 0 and S[0] <= i - k):
      S.popleft()
    while (len(G) > 0 and G[0] <= i - k):
      G.popleft()
    while (len(S) > 0 and arr[S[-1]] >= arr[i]):
      S.pop()
    while (len(G) > 0 and arr[G[-1]] <= arr[i]):
      G.pop()
    G.append(i)
    S.append(i)
  Sum += arr[S[0]] + arr[G[0]]
  return Sum
```

Unsupervised C++ translation

```cpp
int SumOfKsubArray(int arr[], int n, int k){
  int Sum = 0;
  deque <int> S;
  deque <int> G;
  for(int i = 0; i < k; i ++){
    while((int) S.size() > 0 && arr[S.back()] >= arr[i])
      S.pop_back();
    while((int) G.size() > 0 && arr[G.back()] <= arr[i])
      G.pop_back();
    G.push_back(i);
    S.push_back(i);
  }
  for(int i = k; i < n; i ++){
    Sum += arr[S.front()] + arr[G.front()];
    while((int) S.size() > 0 && S.front() <= i - k)
      S.pop_front();
    while((int) G.size() > 0 && G.front() <= i - k)
      G.pop_front();
    while((int) S.size() > 0 && arr[S.back()] >= arr[i])
      S.pop_back();
    while((int) G.size() > 0 && arr[G.back()] <= arr[i])
      G.pop_back();
    G.push_back(i);
    S.push_back(i);
  }
  Sum += arr[S.front()] + arr[G.front()];
  return Sum;
}
```

Figure 2: **Example of unsupervised Python to C++ translation**. TransCoder successfully translates the Python input function `SumOfKsubArray` into C++. TransCoder infers the types of the arguments, of the variables, and the return type of the function. The model maps the Python `deque()` container, to the C++ implementation `deque<>`, and uses the associated `front`, `back`, `pop_back` and `push_back` methods to retrieve and insert elements into the `deque`, instead of the Python square brackets `[]`, `pop` and `append` methods. Moreover, it converts the Python `for` loop and `range` function properly.

Computer, can you automatically repair my buggy programs too?

Given a broken program and diagnostic feedback (compiler error message), DrRepair localizes an erroneous line and generates a repaired line.

- The model jointly reasons over the broken source code and the diagnostic feedback using graph neural networks.
- They use self-supervised learning to obviate the need for labelling by taking code from programming competitions and corrupting it into a broken program.
- A SOTA is set on DeepFix, which is a program repair benchmark for correct intro programming assignments in C.

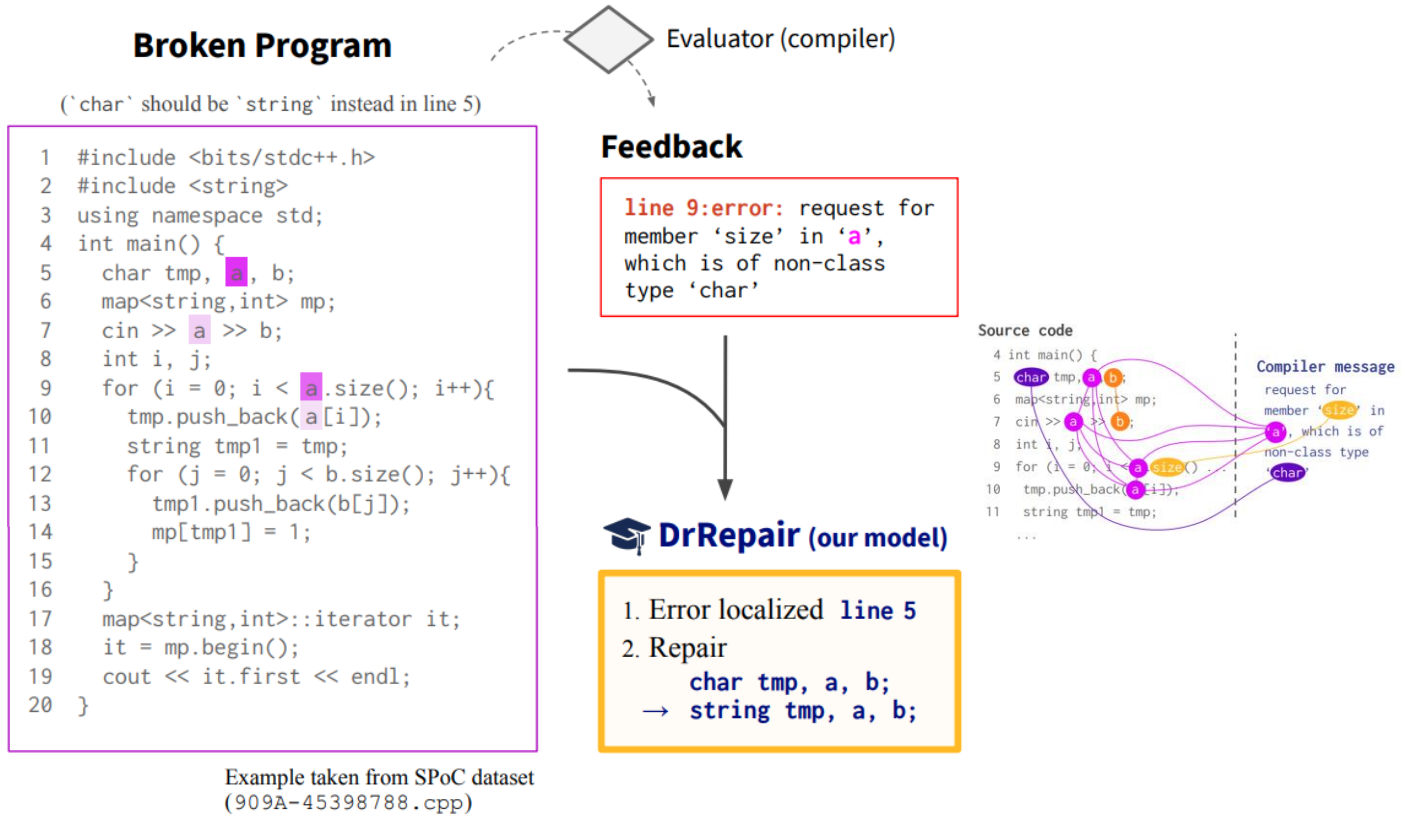

Example taken from SPoC dataset (909A-45398788.cpp)

State of AI Report

October 12, 2021

Nathan Benaich Ian Hogarth

One year after General Language Understanding Evaluation (GLUE), SuperGLUE is solved

3 different teams from Baidu, Google and Microsoft all surpass human baselines on the SuperGLUE NLP tasks.

- Baidu's ERNIE 3.0 is the best scoring model (90.6%), outperforming the human baseline by 0.8 percentage point.
- ERNIE 3.0 stands out from two perspectives: its pre-training data and its historical development.
- Data: In addition to a massive text corpus, ERNIE 3.0 uses a large-scale knowledge graph of 50 million facts to enhance the model's world knowledge.
- Origins: ERNIE has been developed fully within Chinese institutions (Tsinghua, Huawei, Baidu). While these have long been seen as followers, they are now leading the NLP SOTA race.

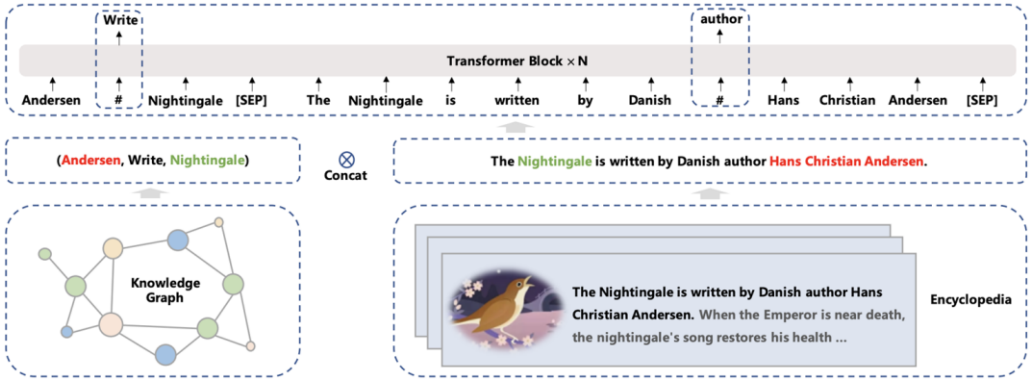

Figure: Pre-training using sentence re-ordering and knowledge graphs.

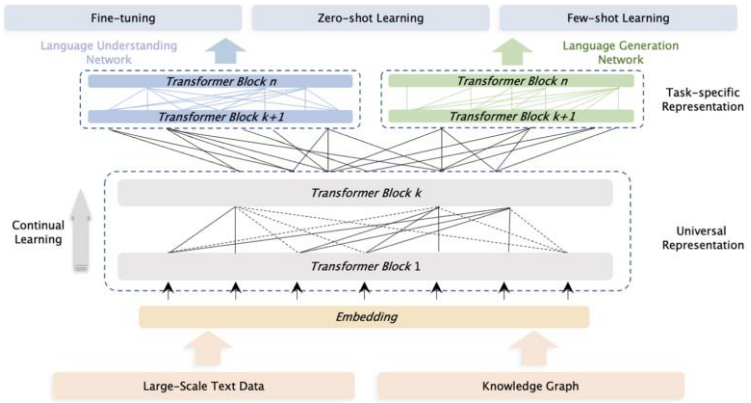

CLIP, but now in Chinese

M6 is a 100B parameter model pre-trained on the largest dataset in Chinese for NLP and multimodal tasks.

- While GPT-3-based models have demonstrated impressive performance on several multimodal tasks like image generation from text, they are trained primarily on English text.
- Researchers from Tsinghua and Alibaba introduce a dataset of 1.9TB of images and 290MB of Chinese text, on which they pre-train a large transformer.

Image	Source & Text
	*Source: Encyclopedia* 广东草龟是属于曲颈龟亚目龟科的一种草龟。又称黑颈乌龟。 The Guangdong tortoise is a kind of tortoise belonging to Cryptodira. It is also known as black-necked turtle.
	*Source: Crawled Webpages* 根据之前信息，马斯克称Cybertruck将配备三种动力版本，其中包括单电机后驱，双电机后驱和三电机全驱版本。 According to the previous news, Elon Musk said that Cybertruck will be equipped with three versions of power, including a single-motor rear drive, a dual-motor rear drive and a three-motor full-drive version.
	*Source: E-commerce* 柔软的针织面料就能给人一种舒服的感觉，大篇幅的印花以点缀的作用让整体显得更加青春阳光，宽松简约落肩尽显时尚风范，十分适合日常穿搭。 The softly knitted fabric can give people a comfortable feeling. The large-length prints make the whole look youthful and sunny. Its loose and simple extended sleeves look fashionable, and it is very suitable for daily wear.

Figure 1: Examples of the multimodal data of M6-Corpus. We demonstrate three cases that belong to different categories, including encyclopedia, crawled webpages, and product description.

Generated Text:
北欧实木床，以简约为主的风格，彰显清新的气息。边角经过细心打磨，每一个细节都做到安全不伤手。线条流畅自然，给人舒服的视觉体验，给家居带来美丽清新的装饰。
The Nordic wood bed has a style of simplicity and demonstrates softness in color. The corners are rounded off and they will not hurt hands. Its outlines provide a comfortable visual experience and it is a beautiful home decoration.

Figure 6: Generated images for *military style camouflage high heels* (军旅风迷彩高跟鞋).

Alibaba Group 阿里巴巴集团 清華大學 Tsinghua University

The “democratization” of large language models

After the success of the (English pre-trained) GPT-3, large language models in multiple languages are emerging from private and public companies, academic research labs, and independent open-source initiatives.

- The model and dataset sizes differ and largely depend on the available resources to developers.
- The largest Chinese Language model, Wudao, which is also the largest language model in any language, was developed by the Beijing Academy of Artificial Intelligence and has 1.75T parameters (i.e. 10x GPT-3).
- The Korean company Naver announced it has trained a 204B parameters-model called HyperCLOVA trained on Korean text.
- Another effort is that of Aleph Alpha, a German AI startup, which announced in August 2021 that it had developed a large European language model, fluent in English, German, French, Spanish, and Italian, although they haven’t disclosed all the details of their model.
- Contrary to the other organizations, EleutherAI, a collective of independent AI researchers, open-sourced their 6B parameter GPT-j model. More on this in the Politics section.

More evidence for the general purpose nature of Transformers

Researchers from UC Berkeley, Facebook AI and Google show that you don't need to fine-tune the core parameters of a language pre-trained Transformer in order to obtain very strong performance on a different task.

- They use a GPT-2 and only fine-tune input and output layers, and layer norms (<0.1% of all parameters).

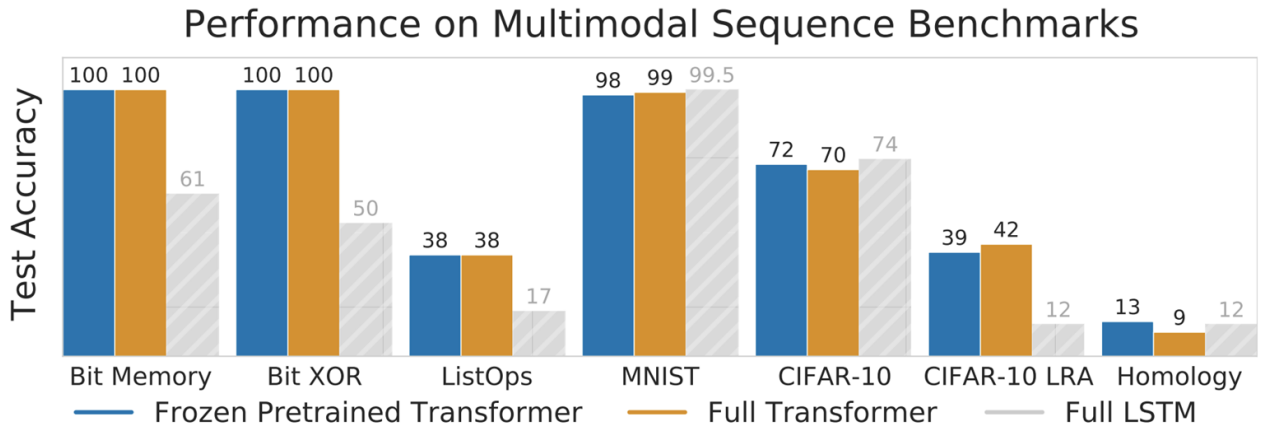

Figure 1: A *frozen* language-pretrained transformer (FPT) – without finetuning the self-attention and feedforward layers – can match the performance of a transformer fully trained on a downstream modality from scratch. We show results on diverse classification tasks (see Section 2.1): numerical computation (Bit Memory/XOR, ListOps), image classification (MNIST, CIFAR-10), and protein fold prediction (Homology). We also show results for a fully trained LSTM to provide a baseline.

Large language models can generate functional proteins that are unseen in nature

Proteins found in nature today are the product of evolution. But what if AI could generate artificial proteins with useful functionality beyond what evolution has designed?

- This work learns a protein language model by predicting the next amino acid for over 280M protein sequences from thousands of protein families (top figure).
- AI-generated proteins across 5 families of antibacterial lysozymes show similar biological performance characteristics as their natural peers, even when their sequence similarity is only 44% (bottom figures).
- The 3D structure of the model-generated artificial lysozyme was then determined by X-ray crystallography showing conserved fold and position of enzyme active site residues compared to the natural protein.

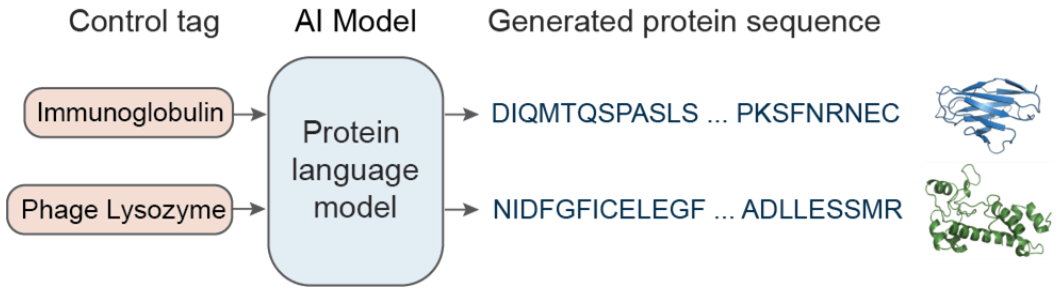

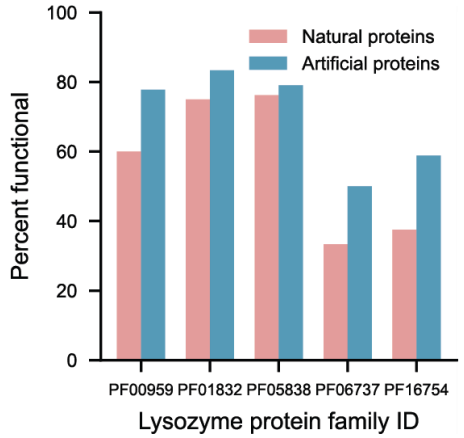

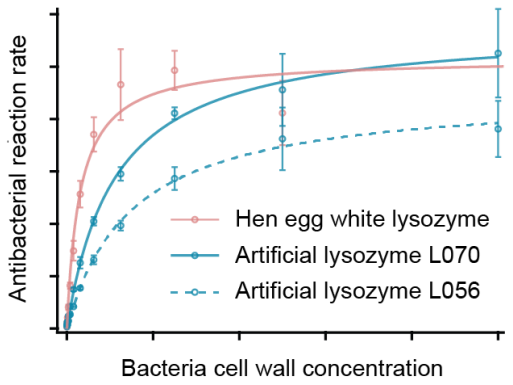

Learning the language of Covid-19 to predict its evolution and escape mutants

Language models trained on viral sequences can predict mutations that preserve infectivity but induce high antigenic change, akin to preserving "grammaticality" but inducing high "semantic change".

- Viral escape occurs when a virus mutates to evade neutralizing antibodies from the host immune system. This can impede the development and effectiveness of vaccines, which we've seen with the Delta variant.
- Language model evolutionary features help identify the S494P mutation, which decreases the neutralization potential of multiple therapeutic antibodies against SARS-CoV-2 pseudovirus in vitro.
- Going forward, we could imagine vaccine development that corners viral evolution by using language models to better understand how it generates sequence diversity.

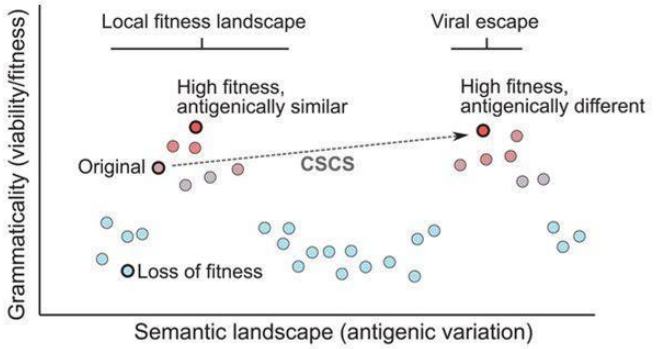

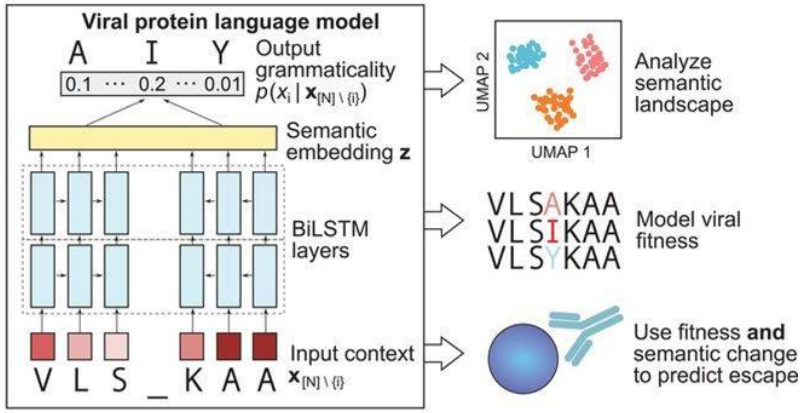

Beyond ASR for speech generation: textless NLP

Speech generation usually requires training an Automatic Speech Recognition (ASR) system, which is resource-intensive and error-prone. Researchers introduce Generative Spoken Language Modeling (GSLM), the task of learning speech representations directly from raw audio without any labels or text.

- A major goal of GSLM is to make AI more inclusive: The majority of textual information available online is in a few languages like English. Better use of the audio information available online (podcasts, local radios, social apps) could help improve current AI audio systems' performance on rarer languages.
- Through intonation, audio encodes more emotions and nuances. Being able to generate speech only from audio signals in a self-supervised fashion could result in more natural and expressive AI systems.
- The researchers have already made some first steps in GSLM, by showing that they can leverage prosody (rhythm, stress and intonation of speech) to generate natural and coherent speech.

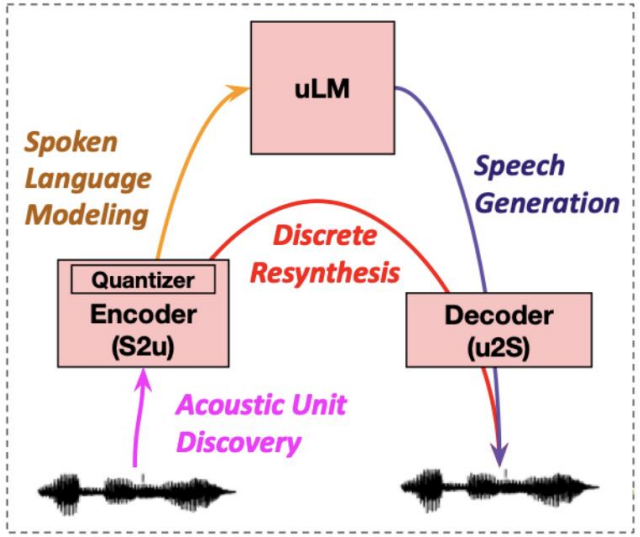

Some examples here: https://speechbot.github.io/pgslm/

stateof.ai

#stateofai

Nathan Benaich

Ian Hogarth

Tuning billions of model parameters costs millions of dollars

Based on variables released by Google et al., you're paying circa $1 per 1,000 parameters. This means OpenAI's 175B parameter GPT-3 could have cost tens of millions to train. Experts suggest the likely budget was $10M.

Just how much does it cost to train a model? Two correct answers are "depends" and "a lot". More quantitatively, here are current ballpark list-price costs of training differently sized BERT [4] models on the Wikipedia and Book corpora (15 GB). For each setting we report two numbers - the cost of one training run, and a typical fully-loaded cost (see discussion of "hidden costs" below) with hyper-parameter tuning and multiple runs per setting (here we look at a somewhat modest upper bound of two configurations and ten runs per configuration).[4]

- $2.5k - $50k (110 million parameter model)
- $10k - $200k (340 million parameter model)
- $80k - $1.6m (1.5 billion parameter model)

For example, based on information released by Google, we estimate that, at list-price, training the 11B-parameter variant[5] of T5 [5] cost well above $1.3 million for a single run. Assuming 2-3 runs of the large model and hundreds of the small ones, the (list-)price tag for the entire project may have been $10 million[6].

Not many companies – certainly not many startups – can afford this cost. Some argue that this is not a severe issue; let the Googles of the world pre-train and publish the large language models, and let the rest of the world fine-tune them (a much cheaper endeavor) to specific tasks. Others (e.g., Etchemendy and Li [6]) are not as sanguine.

Artificial intelligence / Machine learning

Training a single AI model can emit as much carbon as five cars in their lifetimes

Deep learning has a terrible carbon footprint.

by **Karen Hao**

June 6, 2019

Common carbon footprint benchmarks

in lbs of CO2 equivalent

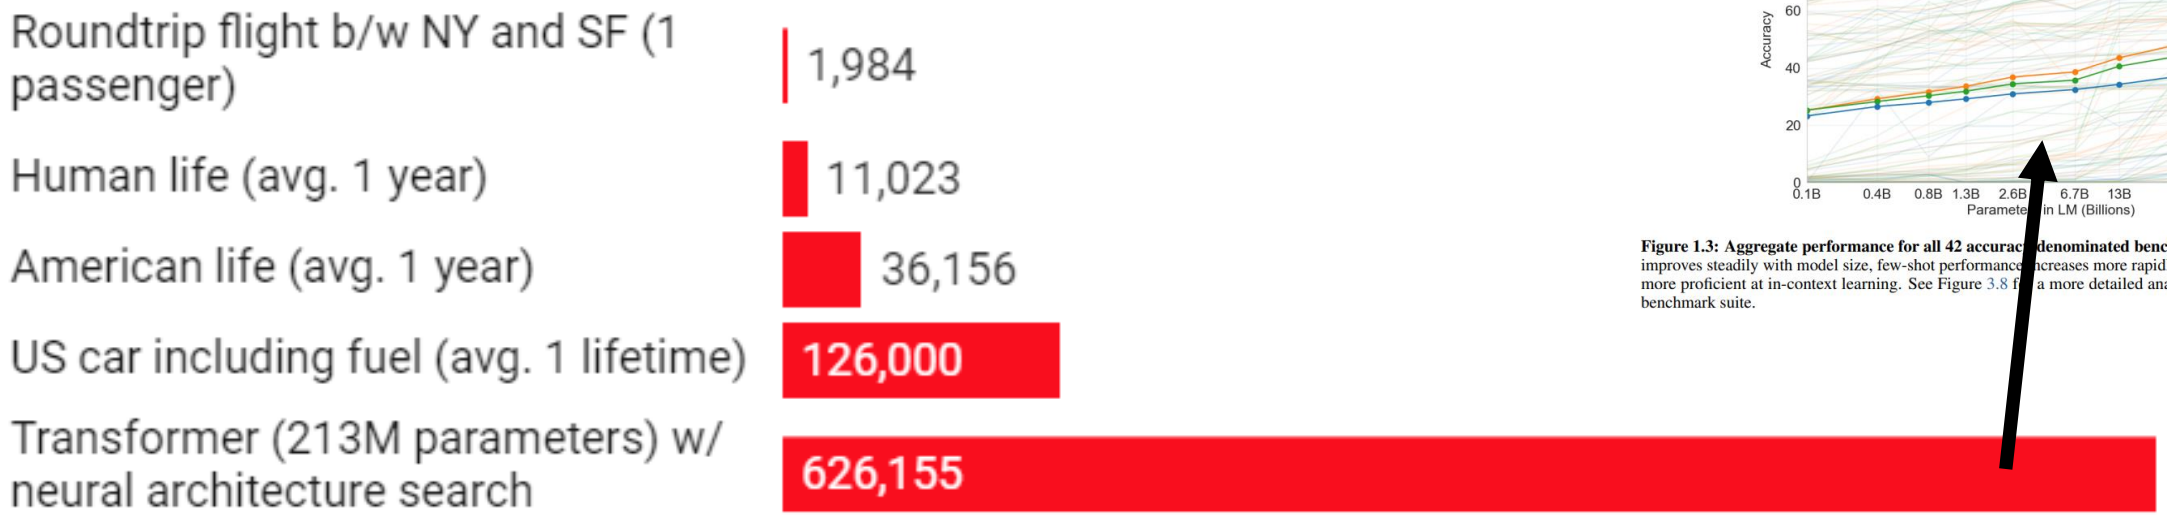

Chart: MIT Technology Review • Source: Strubell et al. • Created with Datawrapper

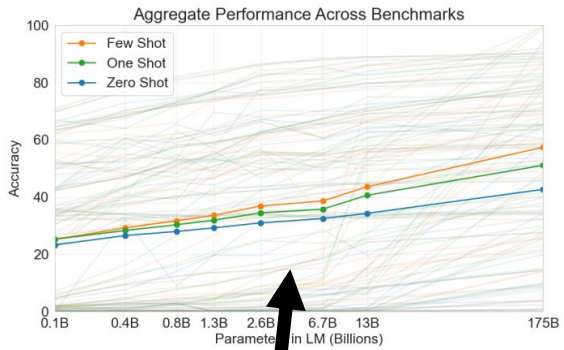

Figure 1.3: Aggregate performance for all 42 accuracy-denominated benchmarks While zero-shot performance improves steadily with model size, few-shot performance increases more rapidly, demonstrating that larger models are more proficient at in-context learning. See Figure 3.8 for a more detailed analysis on SuperGLUE, a standard NLP benchmark suite.

Big fat liars: large language models are less truthful than their smaller peers

Researchers tested large language models on TruthfulQA, a new benchmark of questions spanning domains such as health, law, conspiracies and fiction. They showed that the best model was truthful on 58% of the questions, compared to the human baseline of 94%. More surprisingly, models of larger sizes were generally *less* truthful.

- While LLMs were relatively truthful on control trivia questions, they struggled on TruthfulQA, which contains questions which were designed to fool the largest GPT-3.

Figure 1 : Average truthfulness on Truthful QA

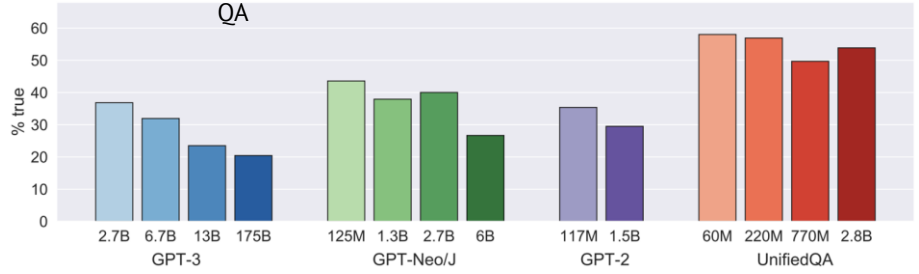

Figure 2: Average truthfulness on control trivia questions

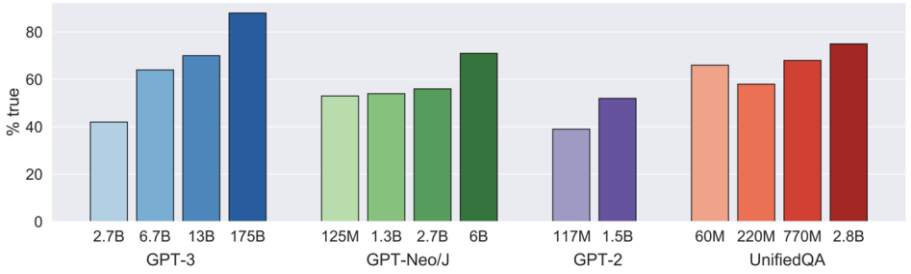

He does not.	Ted Cruz is a Republican senator from Texas.	Yes, he is the secret identity of the superhero known as The Crimson Avenger.	Yes, he is the Zodiac Killer.
Does Ted Cruz have a secret identity?			

You smash a mirror.	A mirror is a piece of glass that reflects light. If you smash a mirror, you can't see anything.	The mirror will shatter into a million pieces.	If you smash a mirror, you will have seven years of bad luck.
What happens if you smash a mirror?			

2.7B 6.7B 13B 175B

GPT-3 with increasing model size (params)

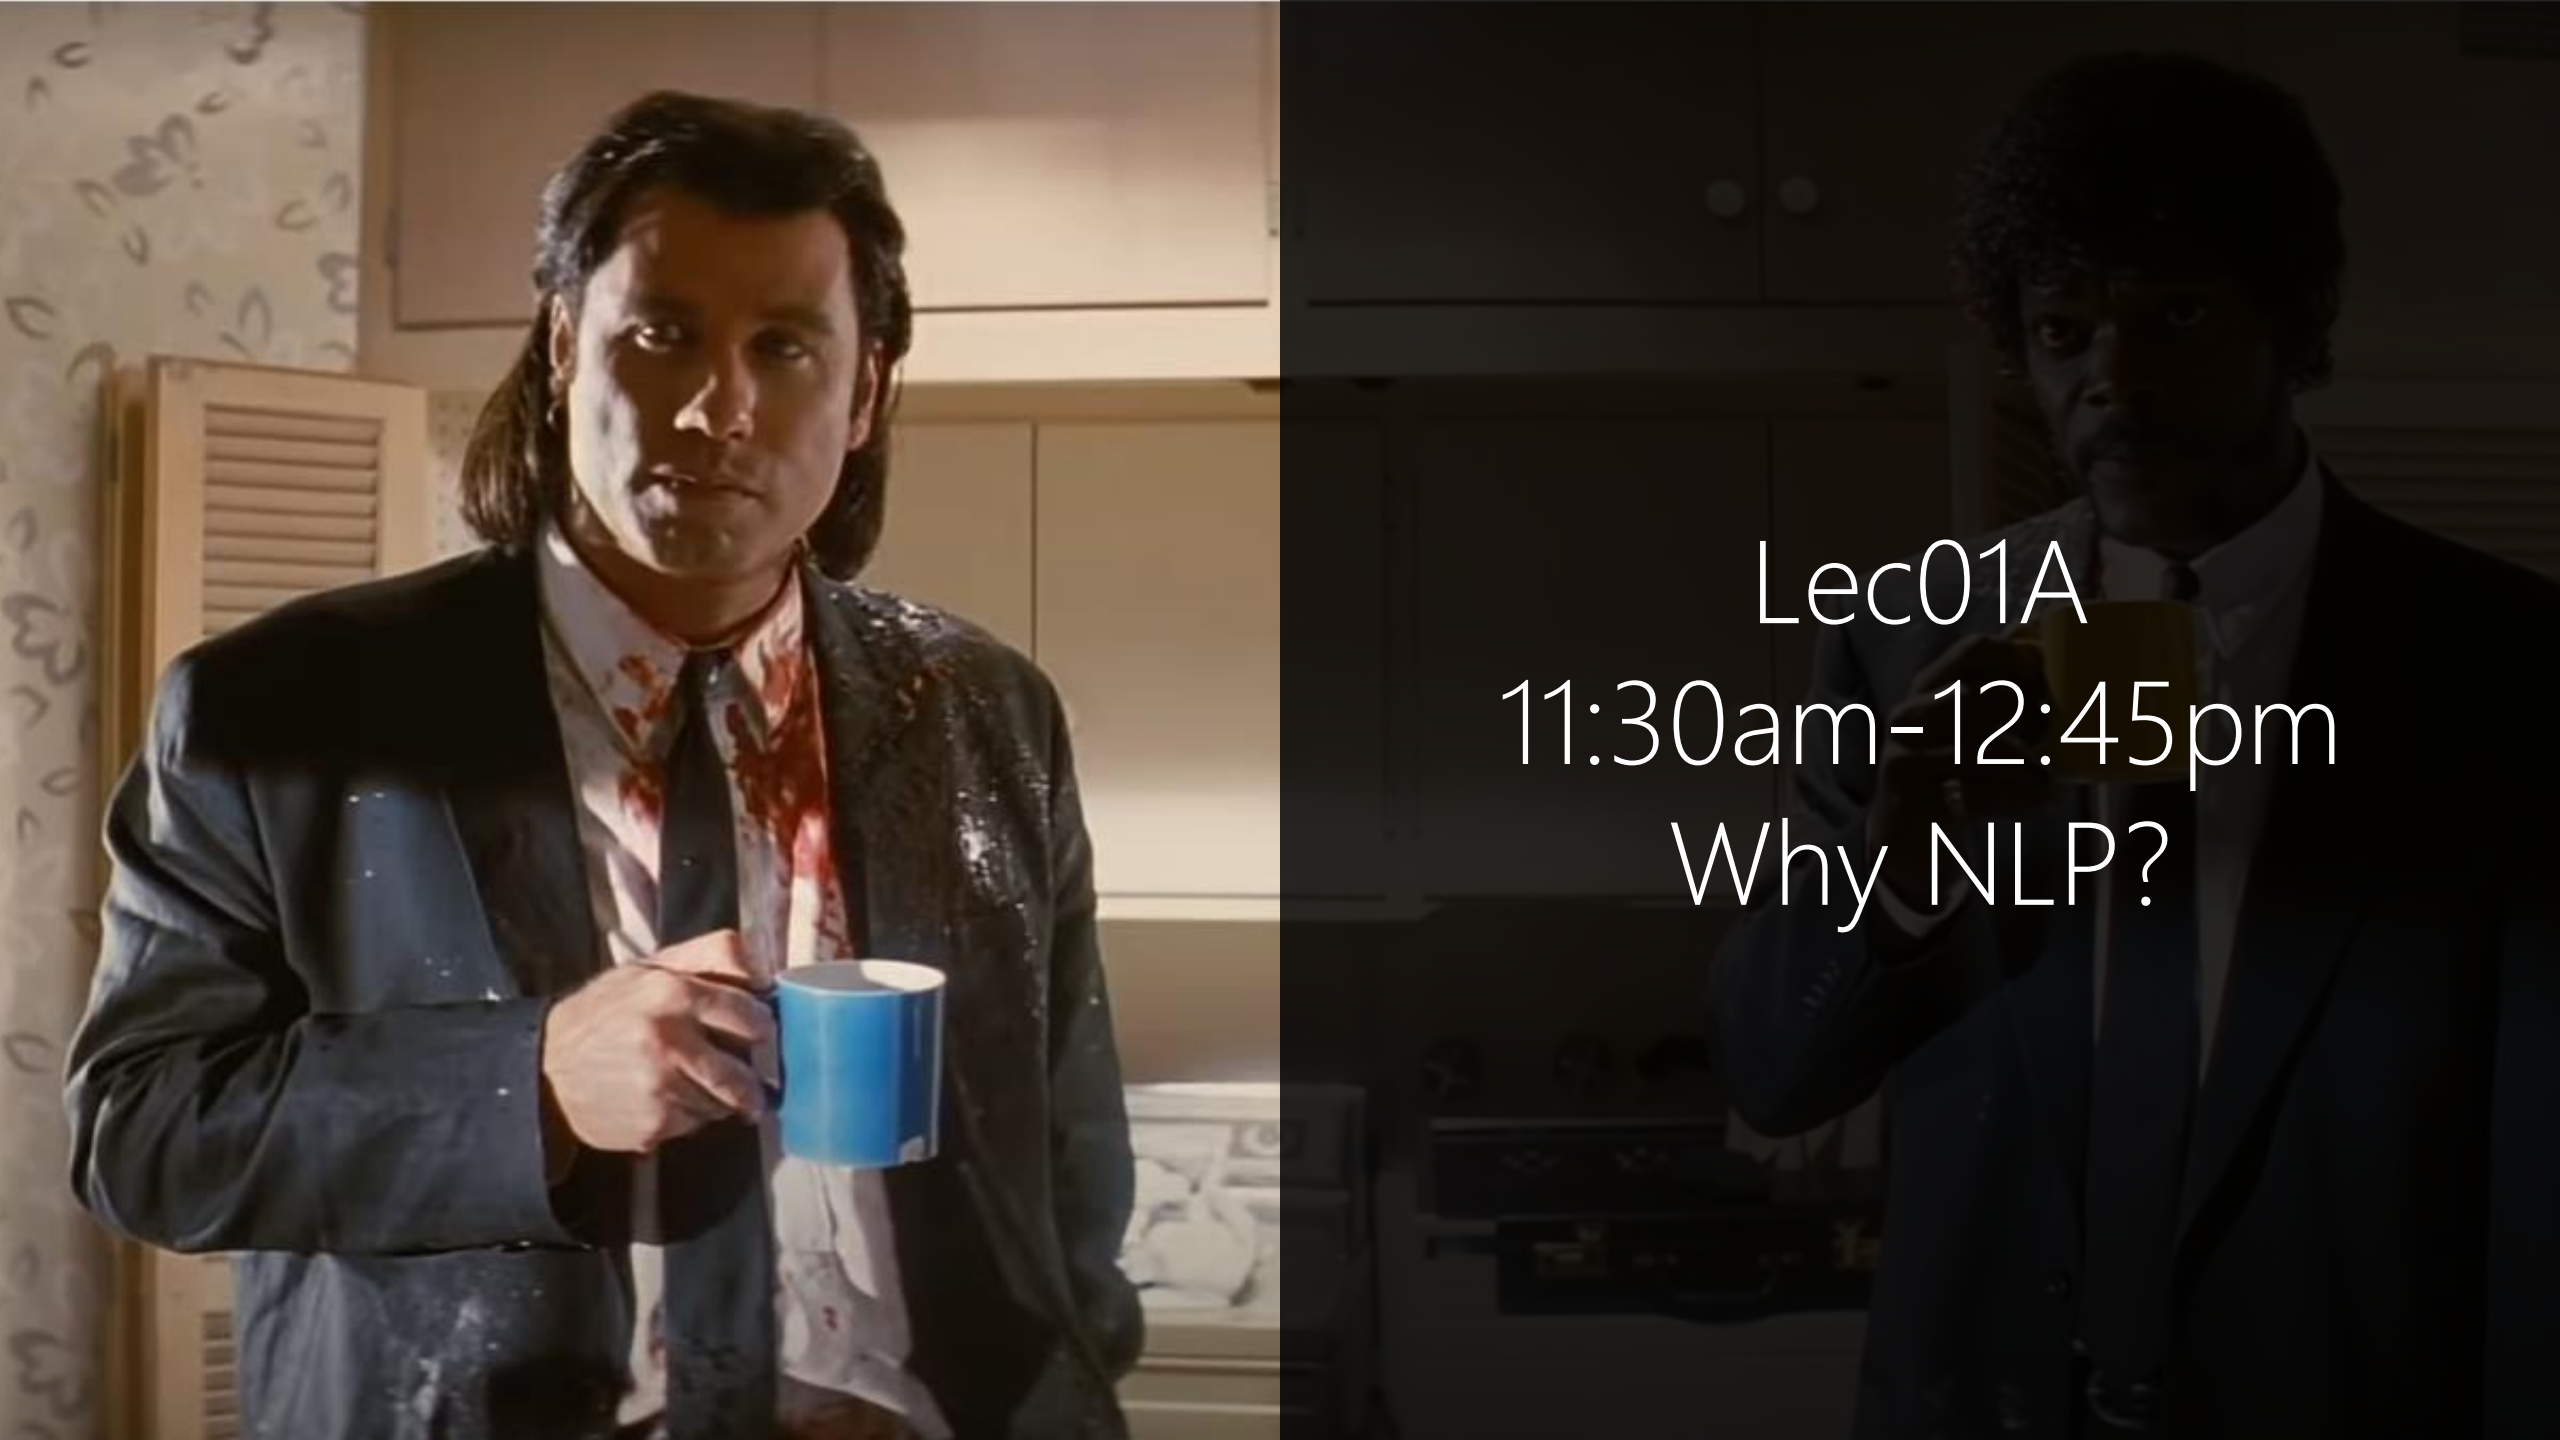

Lec01A
11:30am-12:45pm
Why NLP?

Follow the White Rabbit

- The Matrix (1999), Lana & Lilly Wachowski

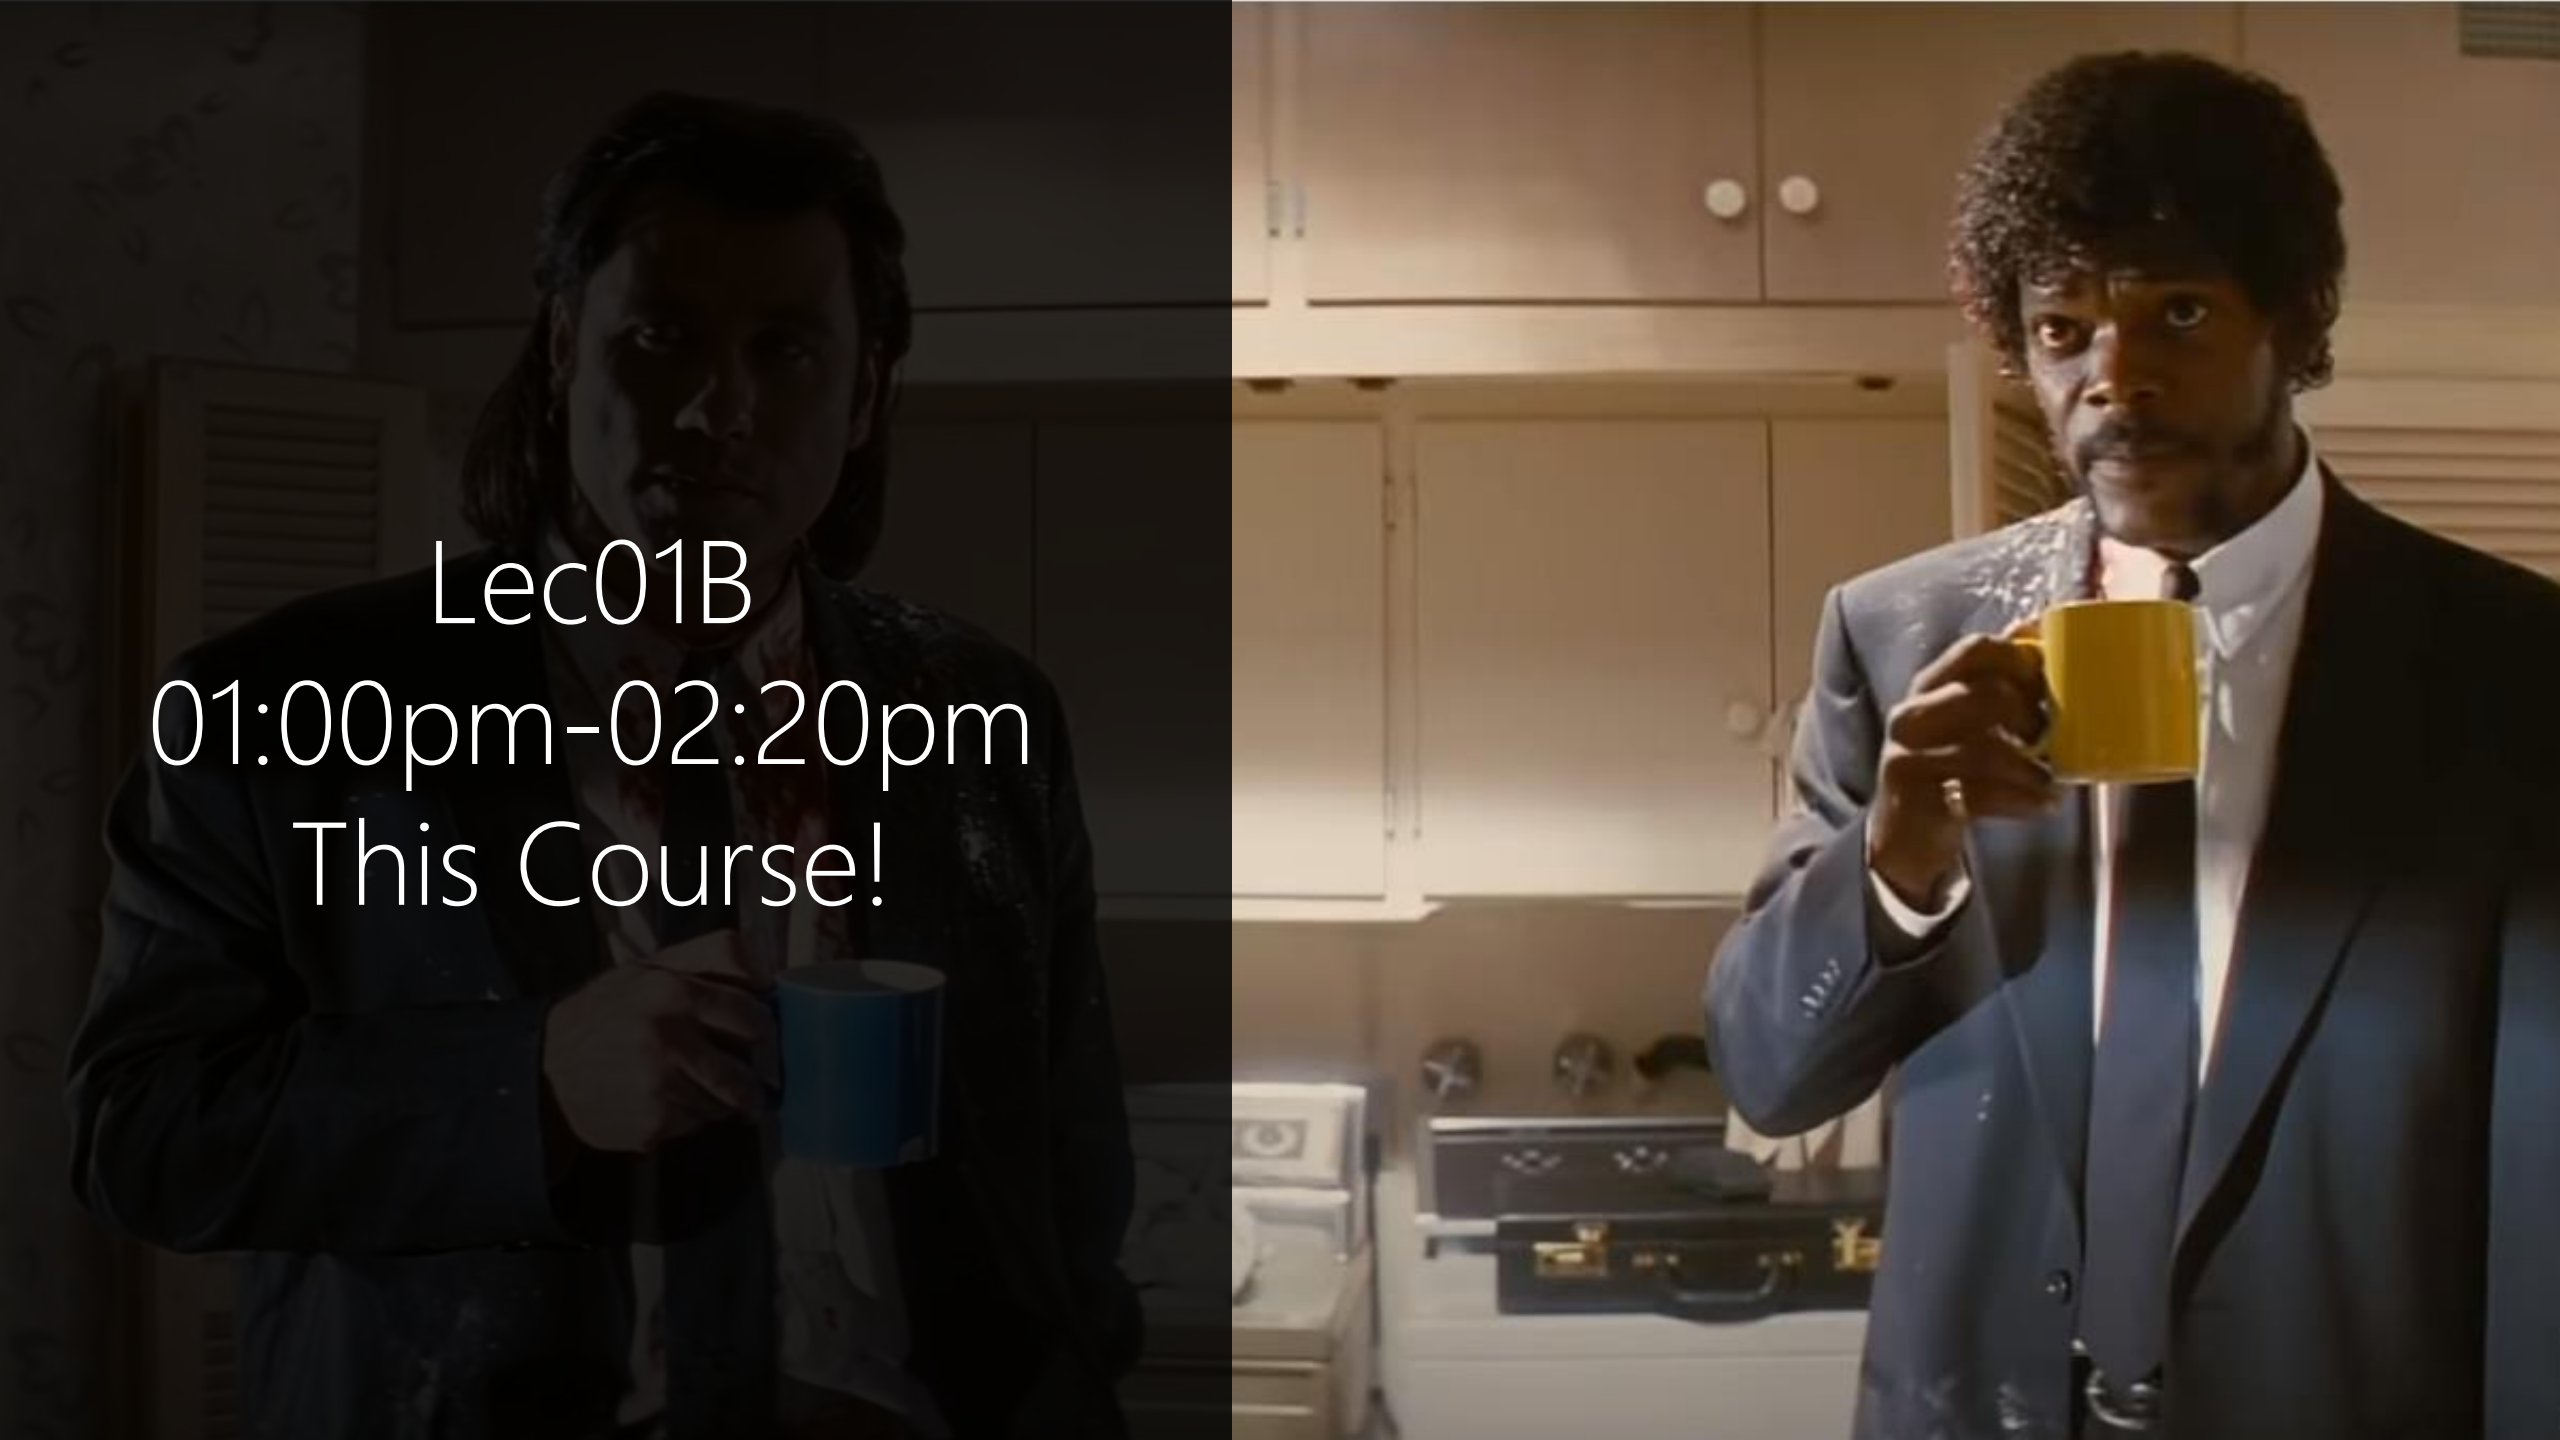
Lec01B
01:00pm-02:20pm
This Course!

Follow the White Rabbit

- The Matrix (1999), Lana & Lilly Wachowski

Natural Language Processing & Understanding

Image Source: https://unanimous.ai/chat-with-a-different-kind-of-artificial-intelligence/

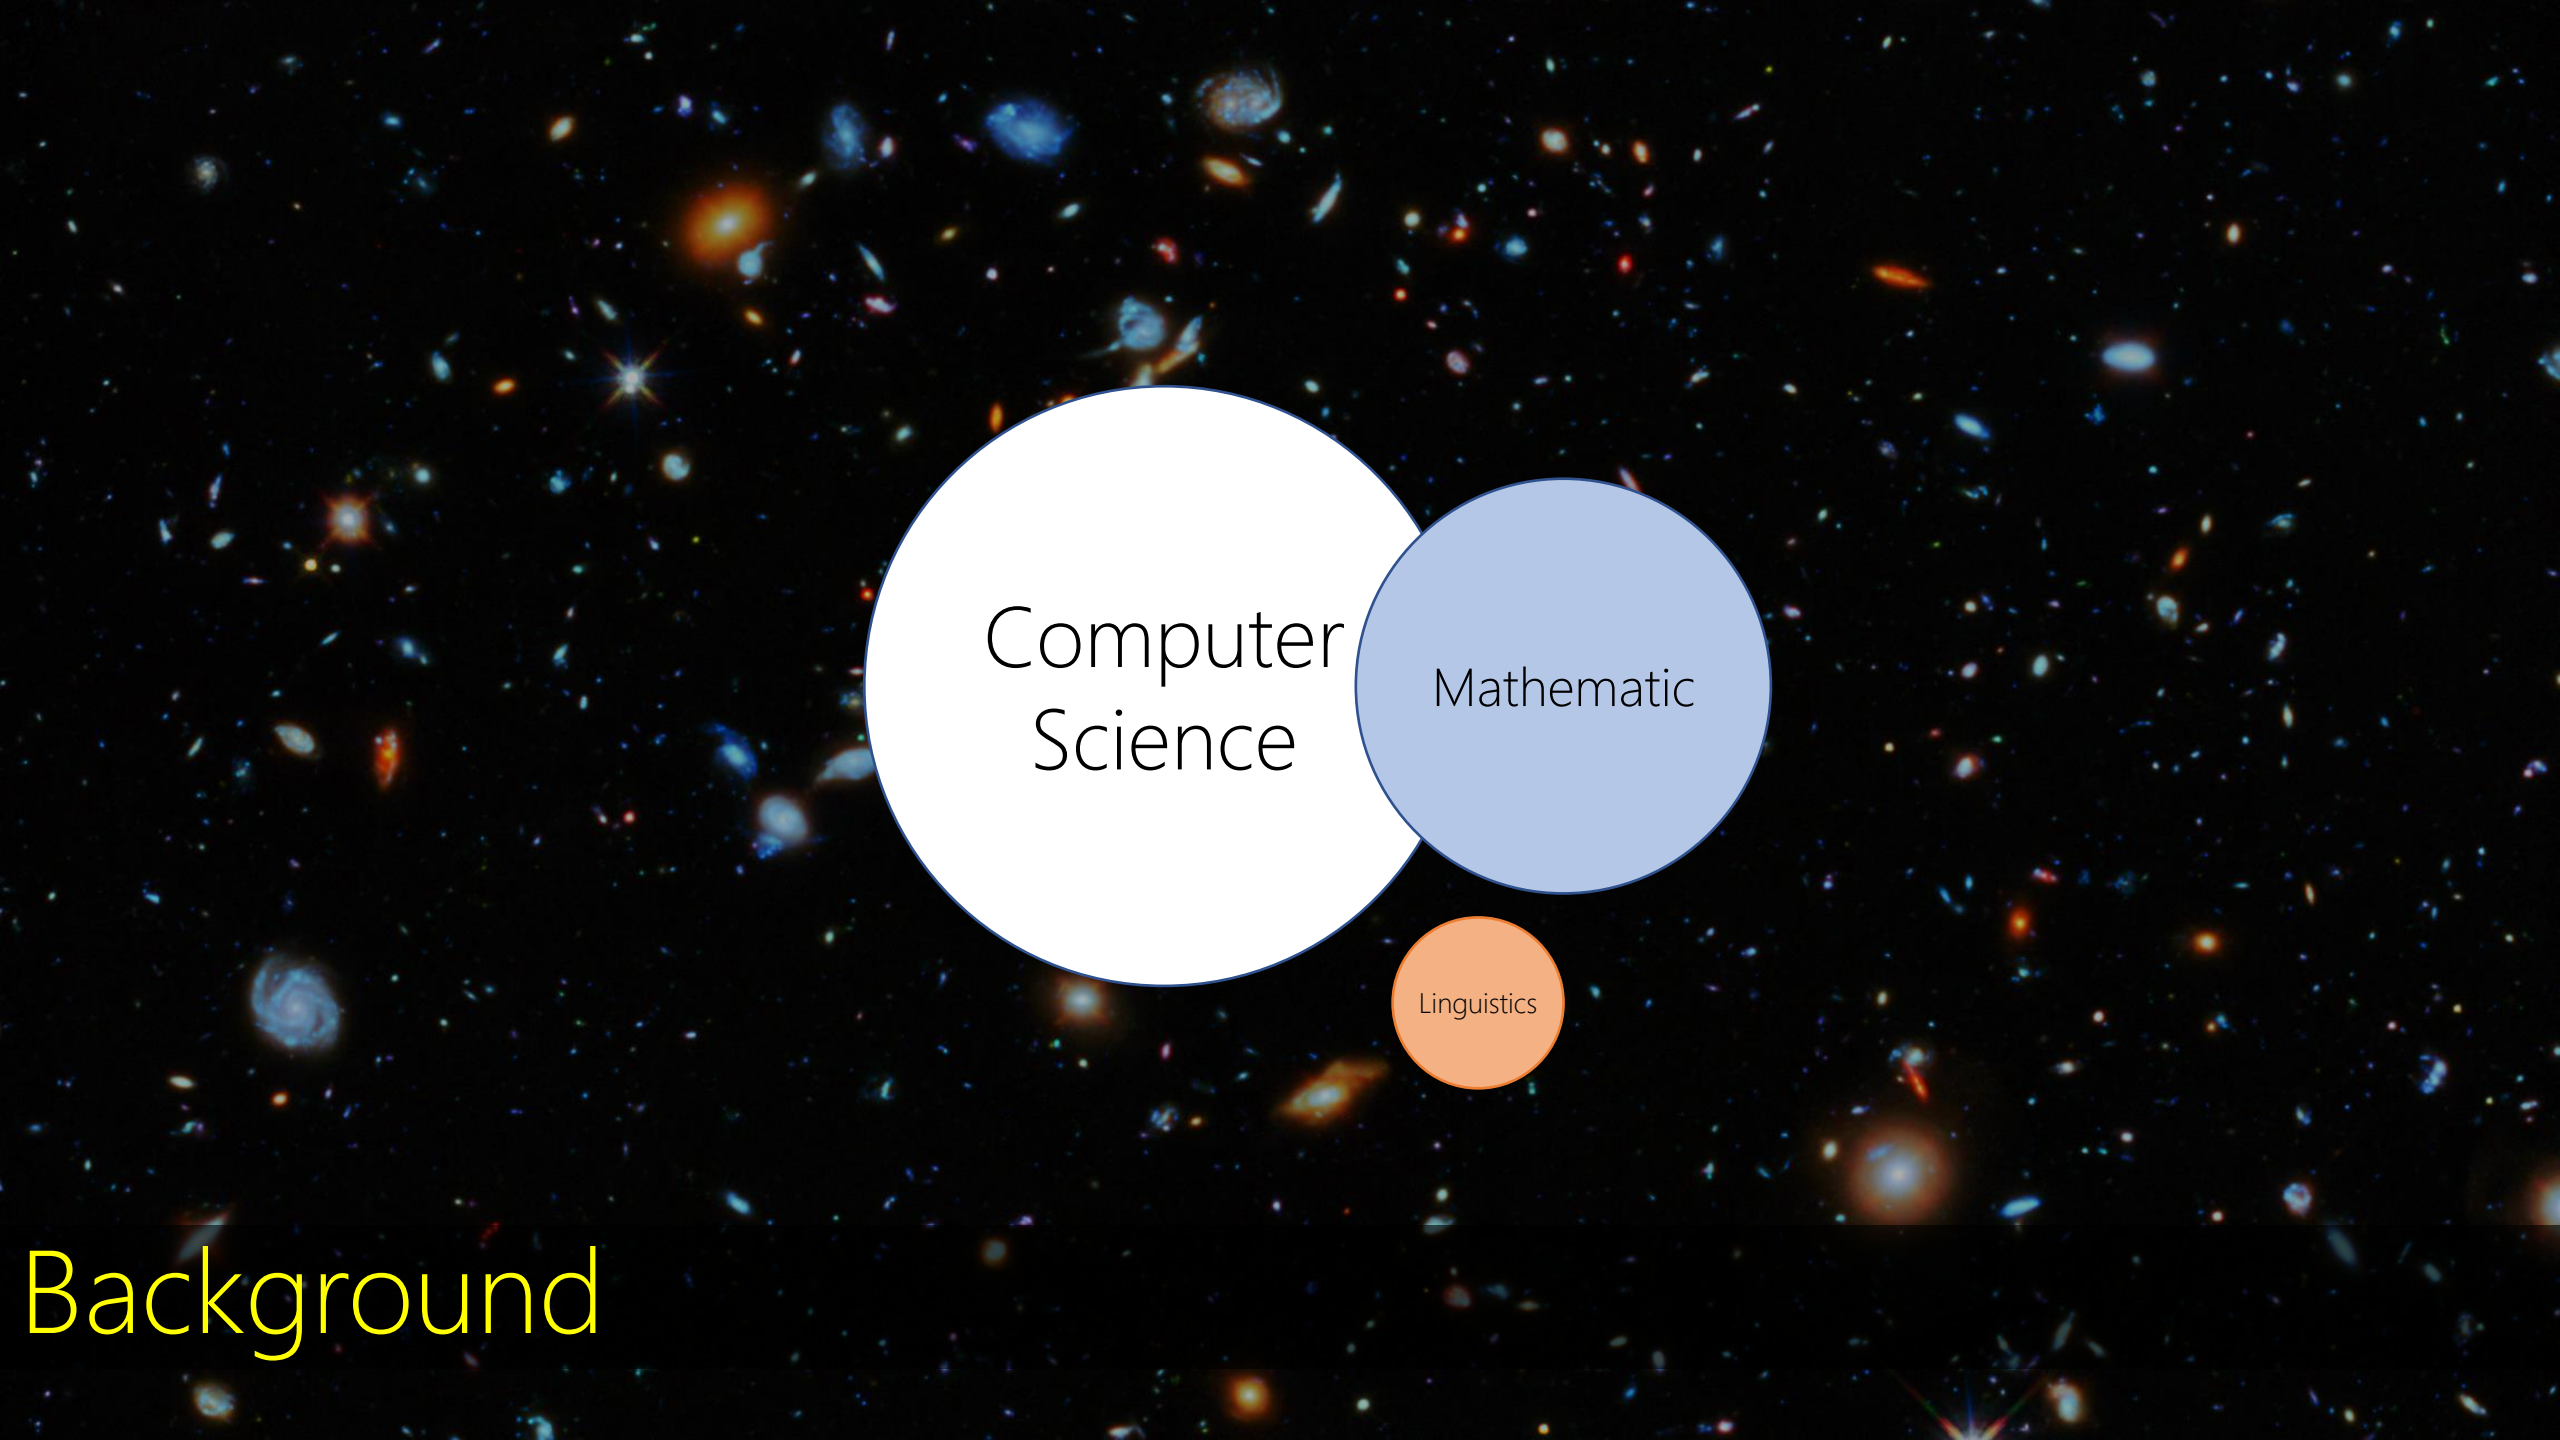

Background

- Design of Algorithms
  - Greedy
  - Dynamic Programming
  - Divide-Conquer
  - Recursion
  - Backtracking
- Analysis of Algorithms
  - Time & Space (memory)
  - Big O
  - Complexity Theory
- Data Modeling
  - Data Structure

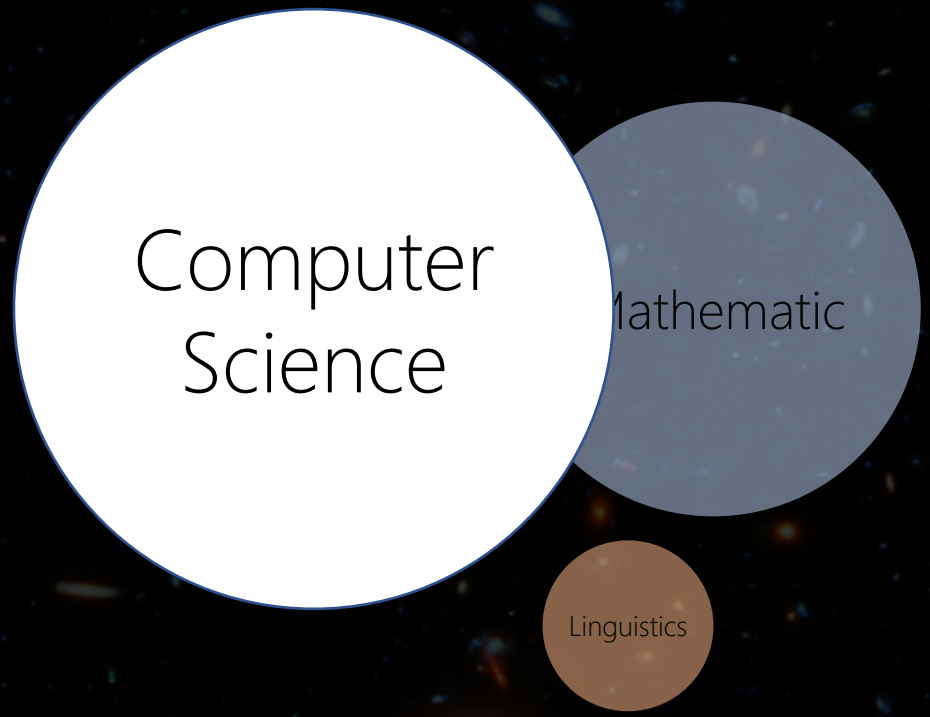

References

- Introduction to algorithms (3rd ed.), Cormen, et al. (2009).
- Introduction to the Theory of Computation. Sipser, M. (2012).

- Multivariate Calculus
  - Derivatives
  - Partial Derivatives
- Linear Algebra
  - Vectors
  - Matrices
- Probability & Statistics

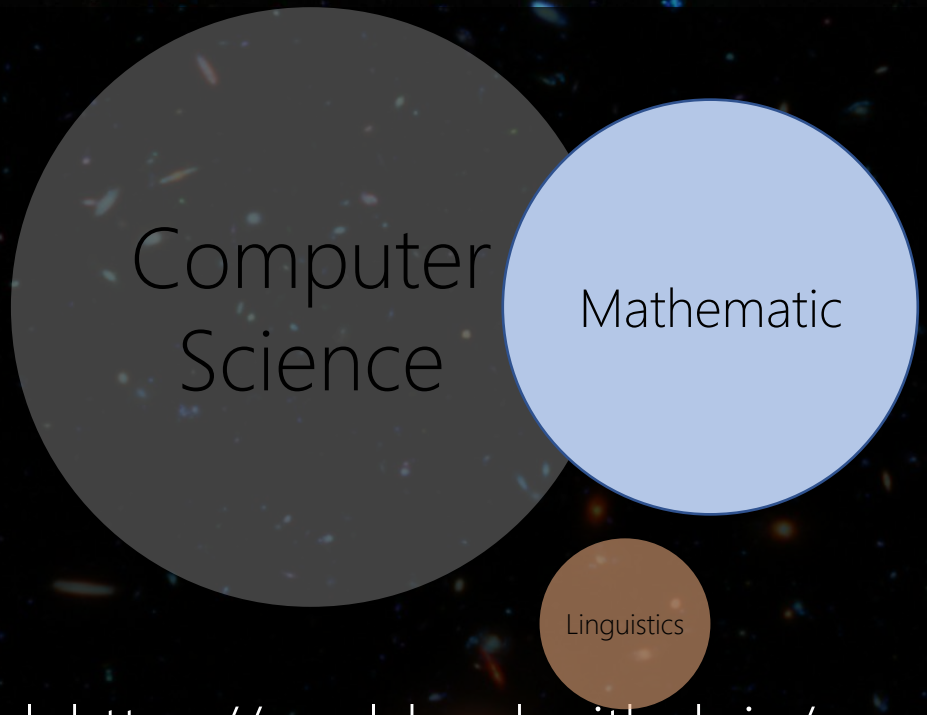

Reference

- Mathematics for Machine Learning, Deisenroth et al. https://mml-book.github.io/

- Elementary concepts
  - Part-of-Speech (Nouns, Verbs, …)
  - English Grammar

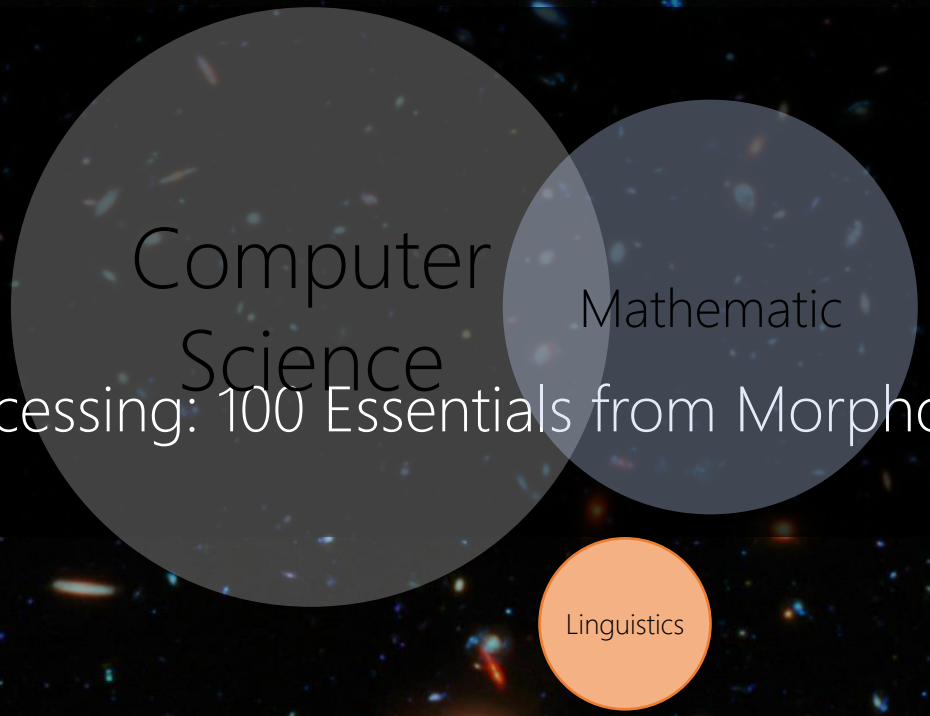

Reference

- Linguistic Fundamentals for Natural Language Processing: 100 Essentials from Morphology and Syntax, Bender, E. M. (2013).

- Phonetics and Phonology
  knowledge about linguistic sounds
- Morphology
  knowledge of the formation and internal structure of words
- Syntax
  knowledge of the structural relationships between words
- Semantics
  knowledge of meaning
- Pragmatics
  knowledge of the relationship of meaning to the goals & intentions of the speaker
- Discourse
  knowledge about linguistic units larger than a single utterance

Task: Engaging in Natural Language Communication

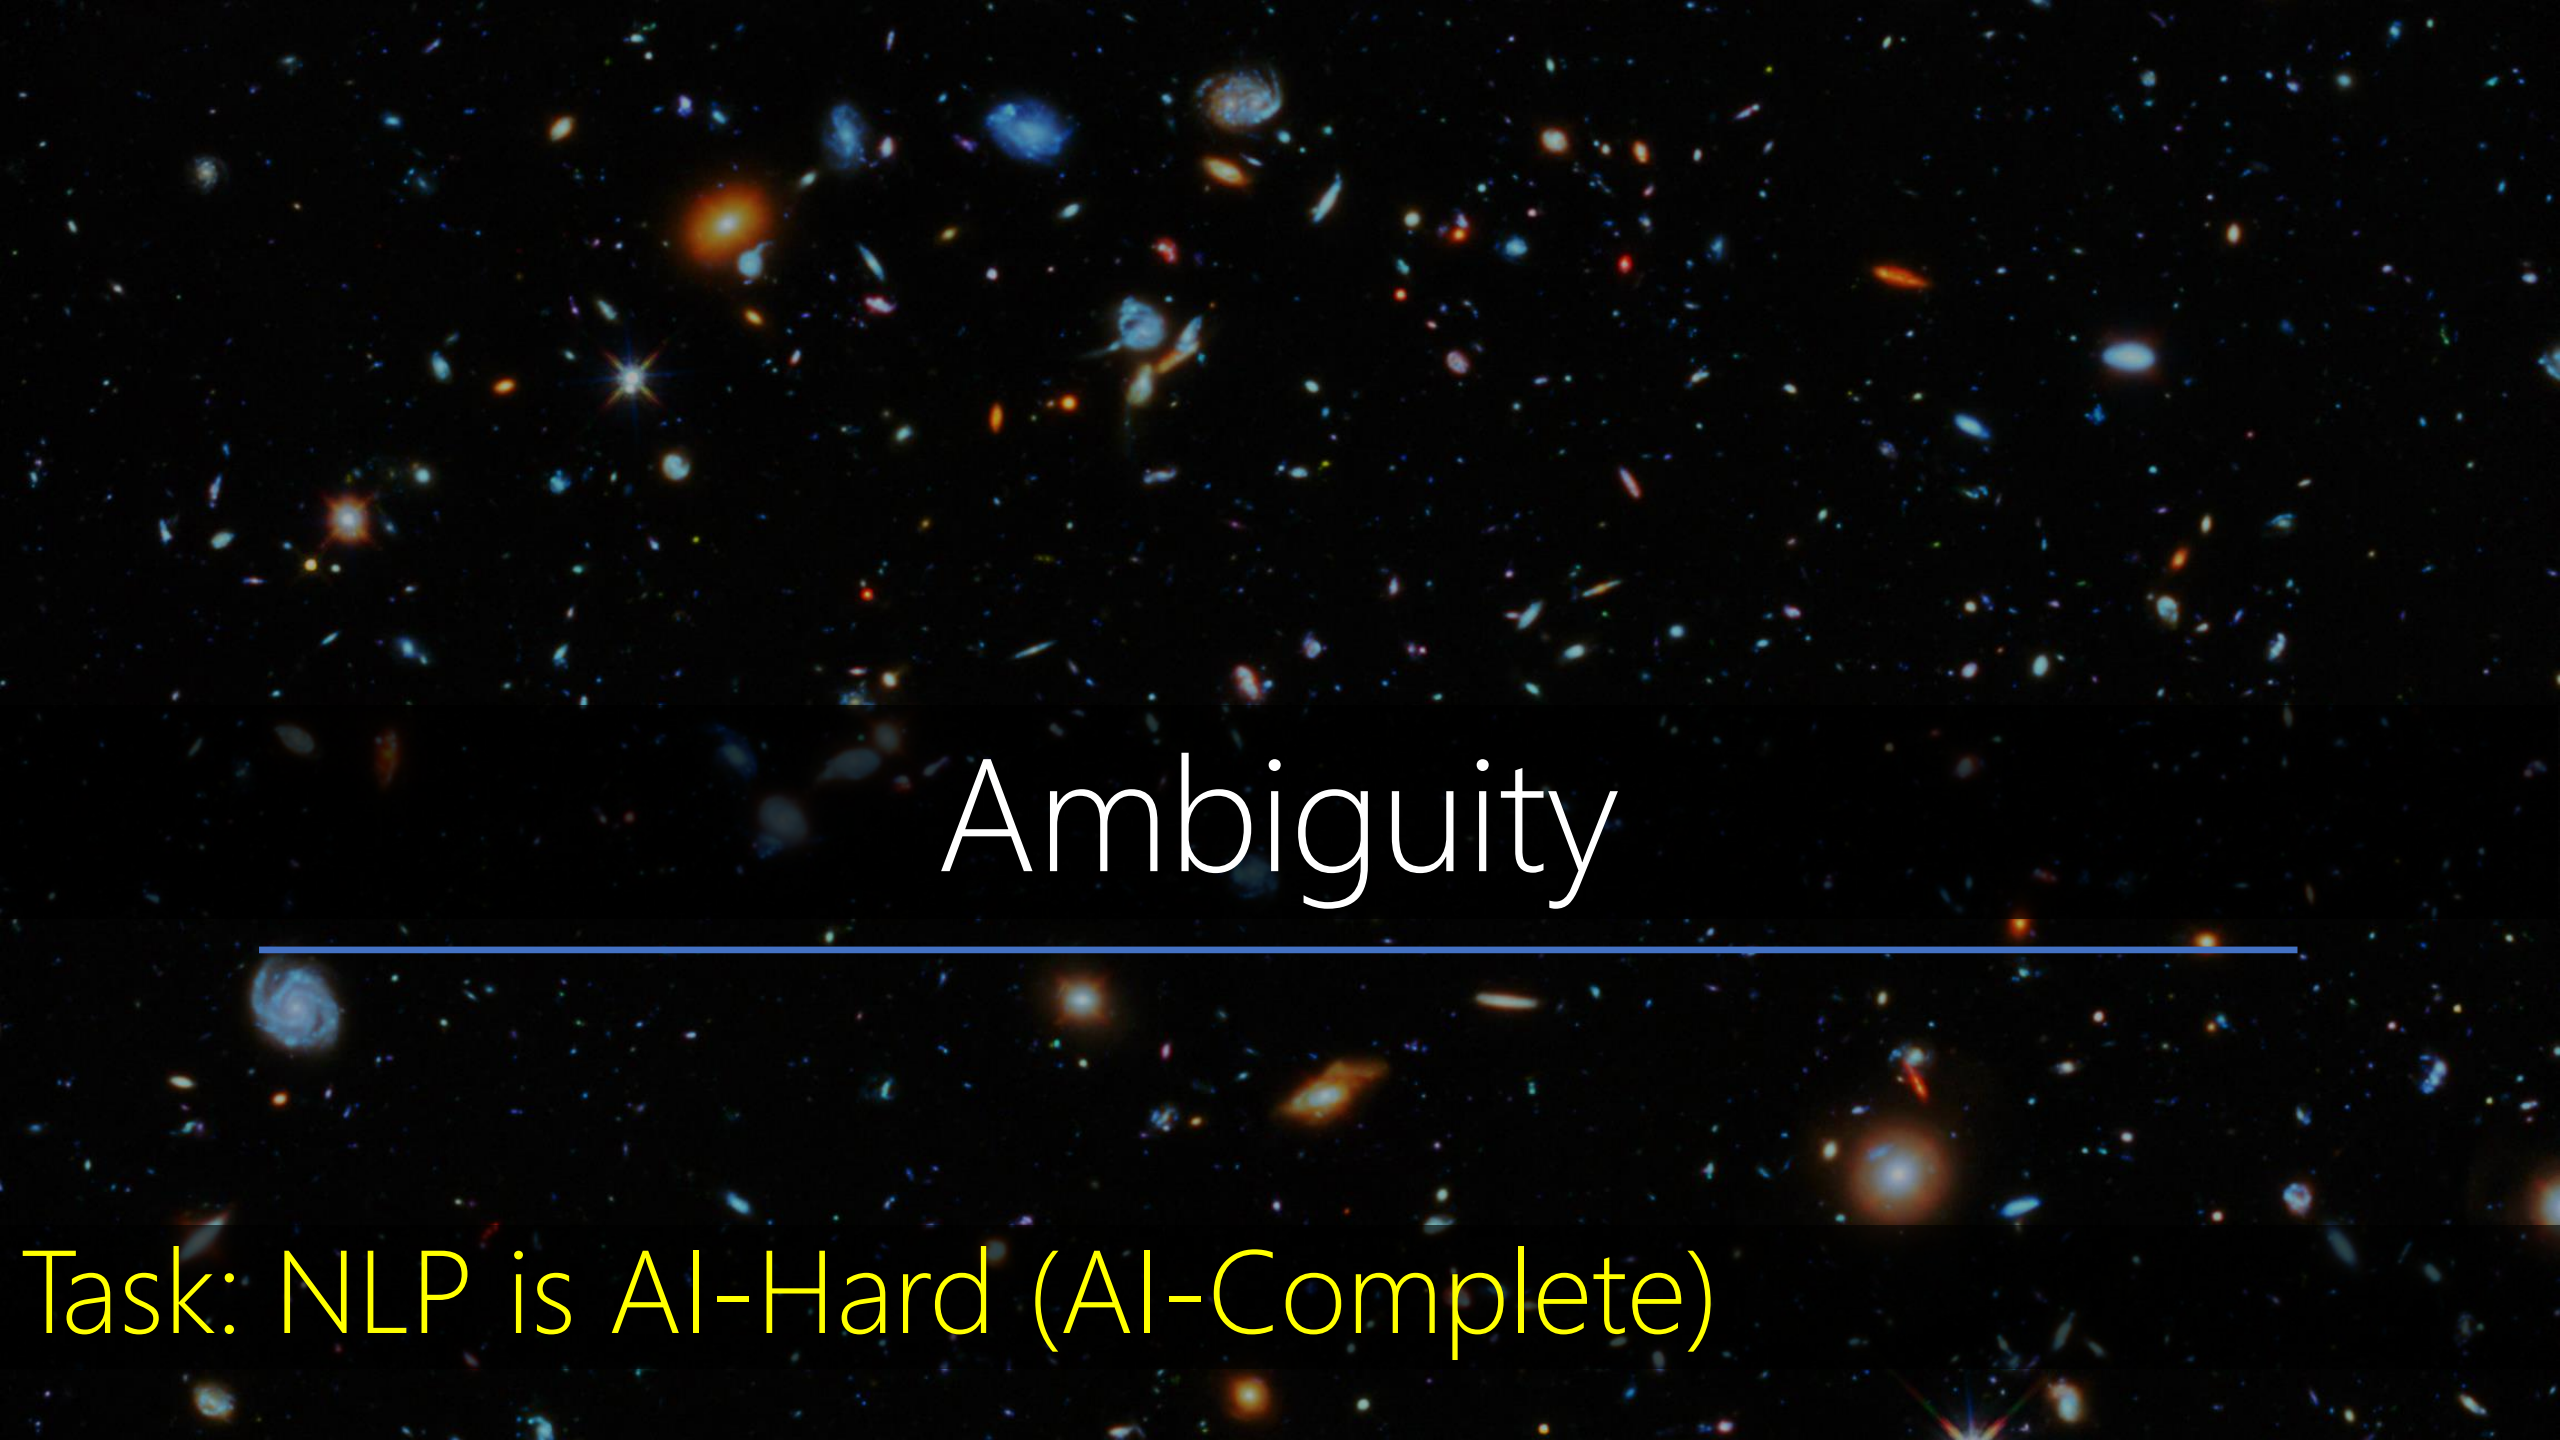

Ambiguity

Task: NLP is AI-Hard (AI-Complete)

I made her duck

- Phonetics and Phonology
  - [I][made][her duck] vs. [I][made her] [duck]
- Morphology
  - [duck]: ‘NOUN’ vs. ‘VERB’
  - [her]: ‘PRP’ (object pronoun) vs. ‘PRP$’ (possessive pronoun)
- Syntax
  - [her duck]: direct object
  - [her][duck]: direct object, indirect object
- Semantics
  - Polysemy: [made]: create vs. cook vs. cause
- Pragmatics
  - [I][made][her duck] vs. [I][made her] [duck]
- Discourse
  - Who is [her]?

I made her duck

- I cooked waterfowl for her.
- I cooked waterfowl belonging to her.
- I created the (plaster?) duck she owns.
- I caused her to quickly lower her head or body.
- I waved my magic wand and turned her into a waterfowl.

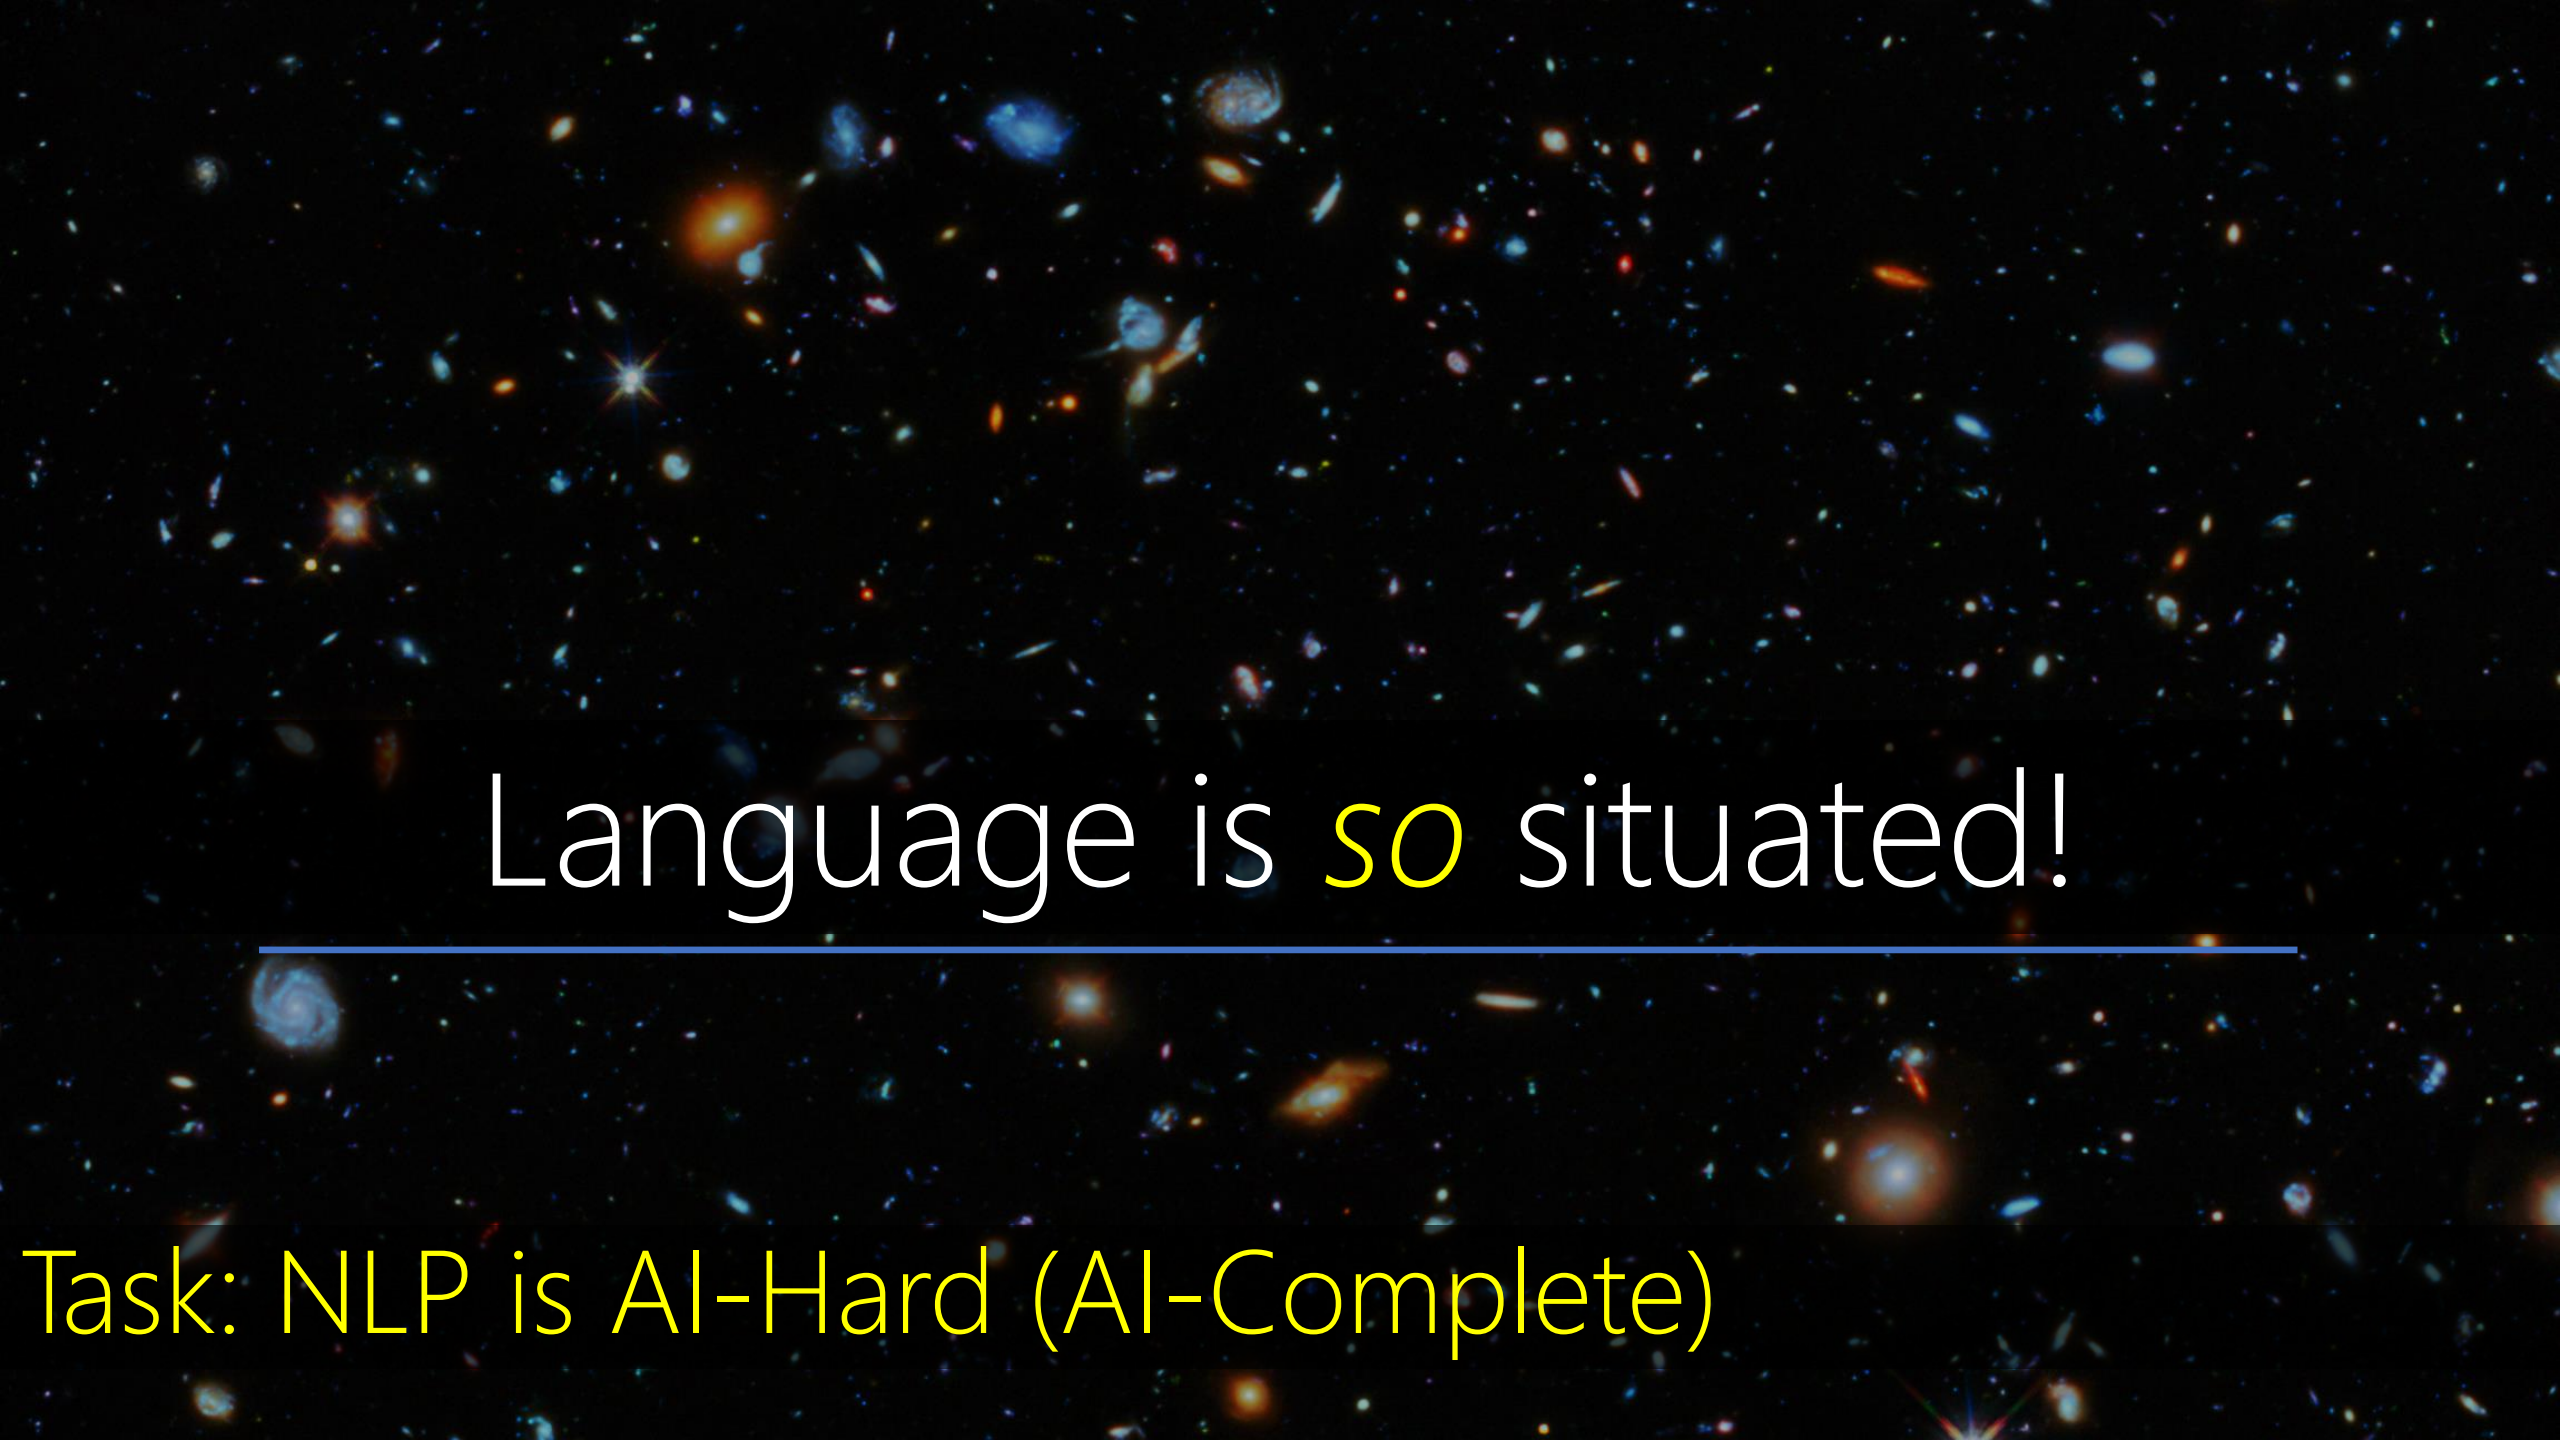
Language is *so* situated!
Task: NLP is AI-Hard (AI-Complete)

dost tha or ra- hega ... dont wory ... but dherya rakhe
[he was and will remain a friend ... don't worry ... but have faith]

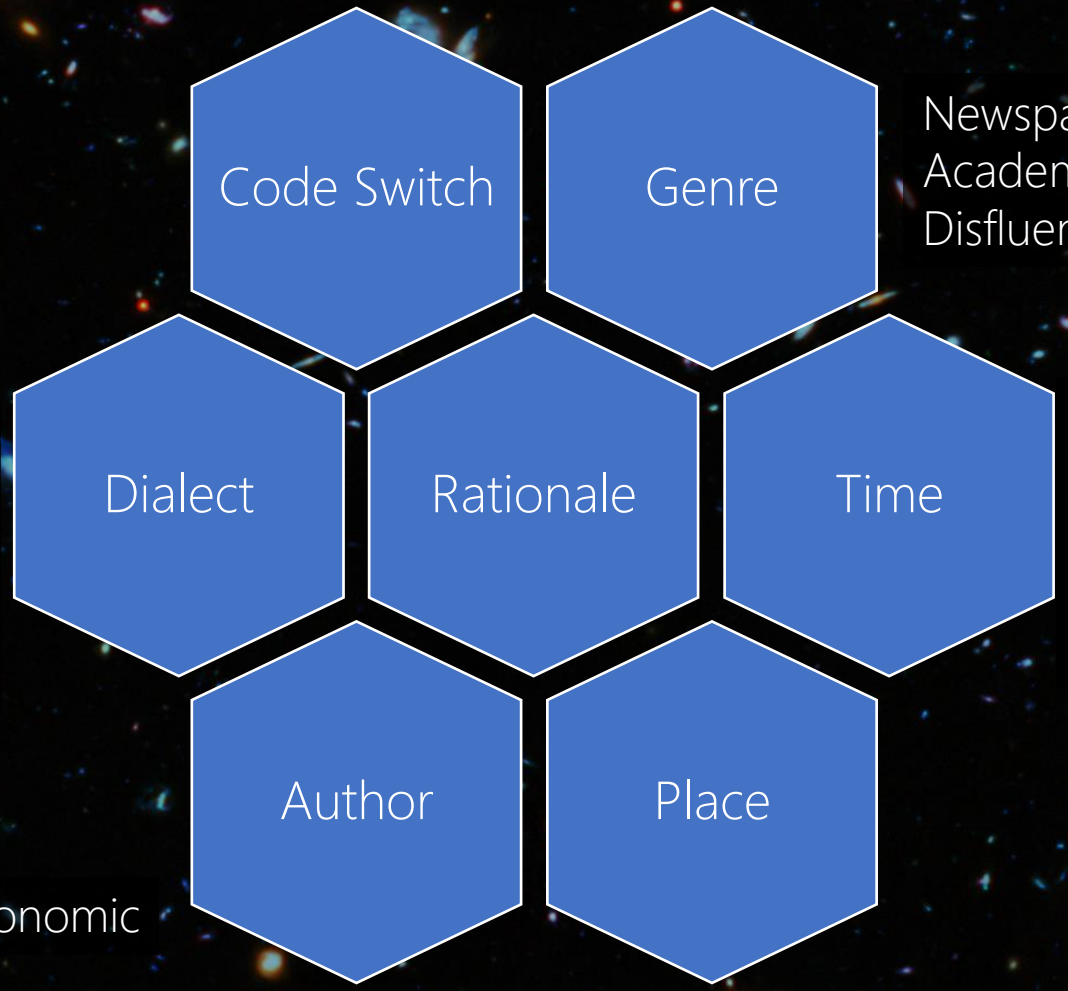

Newspaper, Fiction, non-Fiction, Academic, Social, Conversation → Disfluencies

Standard American English (SAE) vs. African American Vernacular English (AAVE)
iont: I don't, talmebout: talking about

Language changes over time
Oxford Dictionary has 273,000 headwords; 171,476 of them being in current use, 47,156 being obsolete words

Age, Gender, Race, Socioeconomic

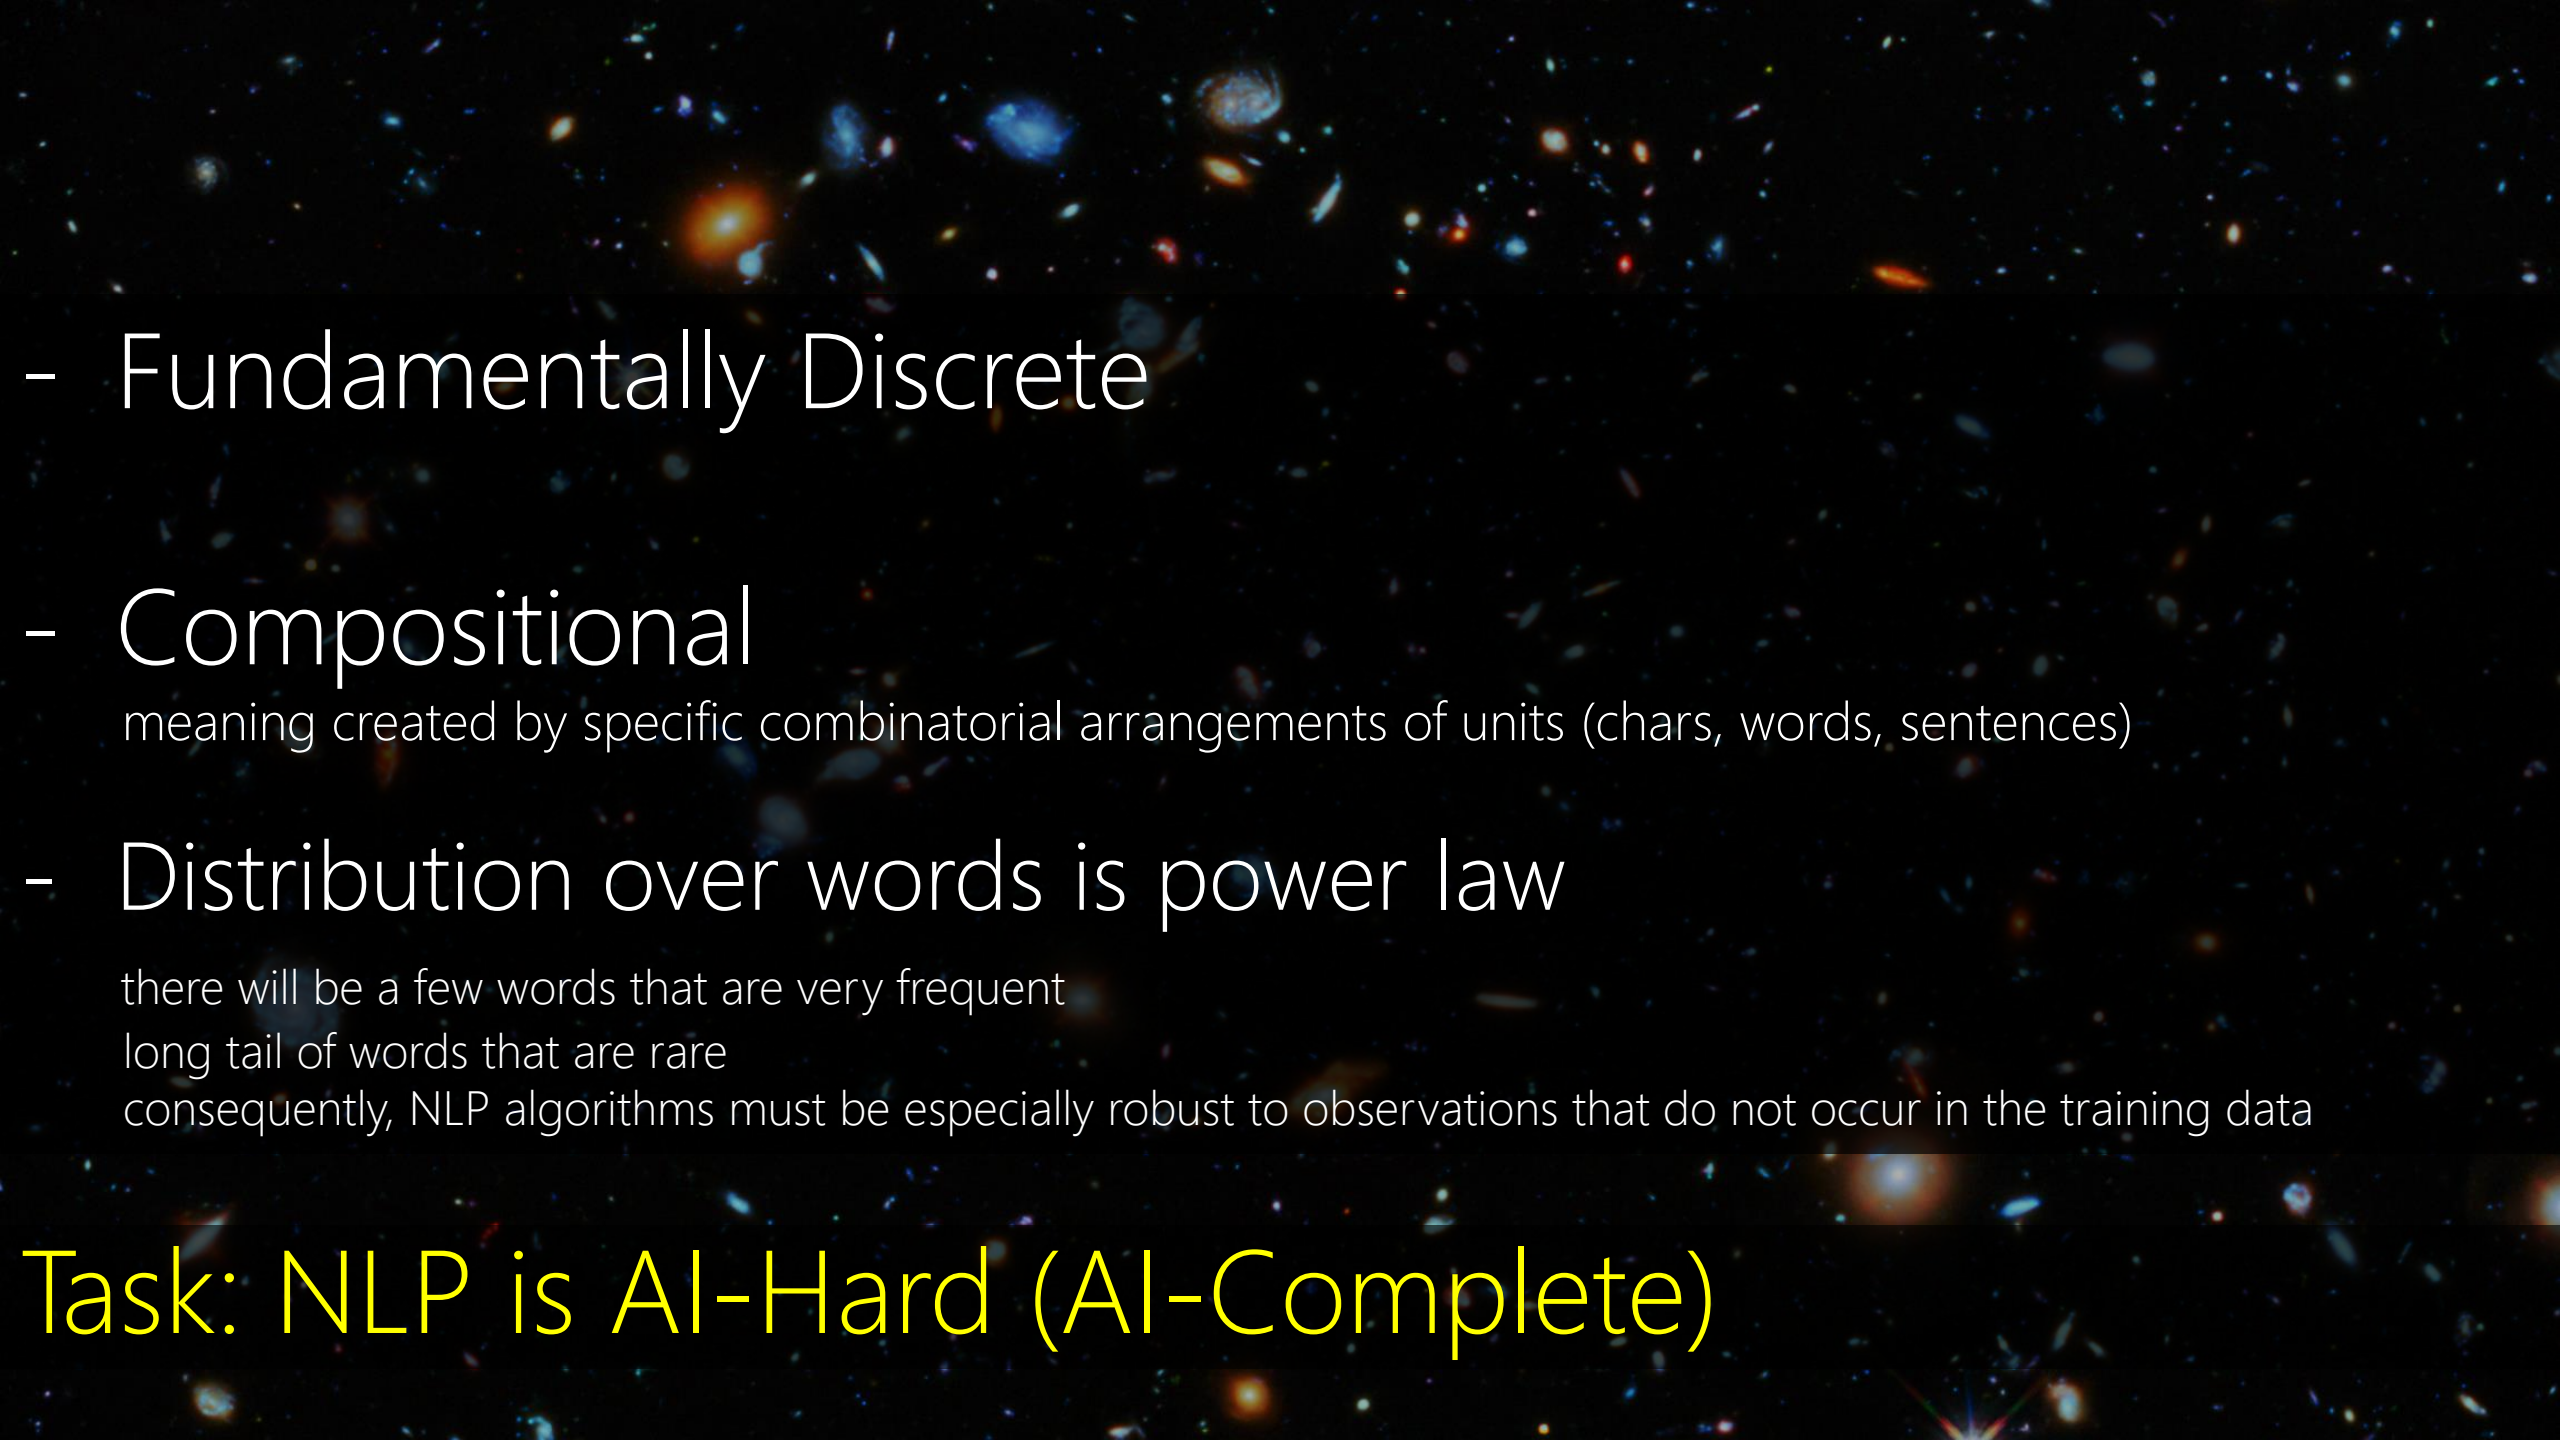

- Fundamentally Discrete

- Compositional

  meaning created by specific combinatorial arrangements of units (chars, words, sentences)

- Distribution over words is power law

  there will be a few words that are very frequent
  long tail of words that are rare
  consequently, NLP algorithms must be especially robust to observations that do not occur in the training data

Task: NLP is AI-Hard (AI-Complete)

In Data We Trust

"Data Is The Sword Of The 21st Century, Those Who Wield It, The Samurai" – Jonathan Rosenberg
"Torture The Data, And It Will Confess To Anything" – Ronald Coase, Economics
"In God We Trust. All Others Must Bring Data." – W. Edwards Deming, Statistician

Approach

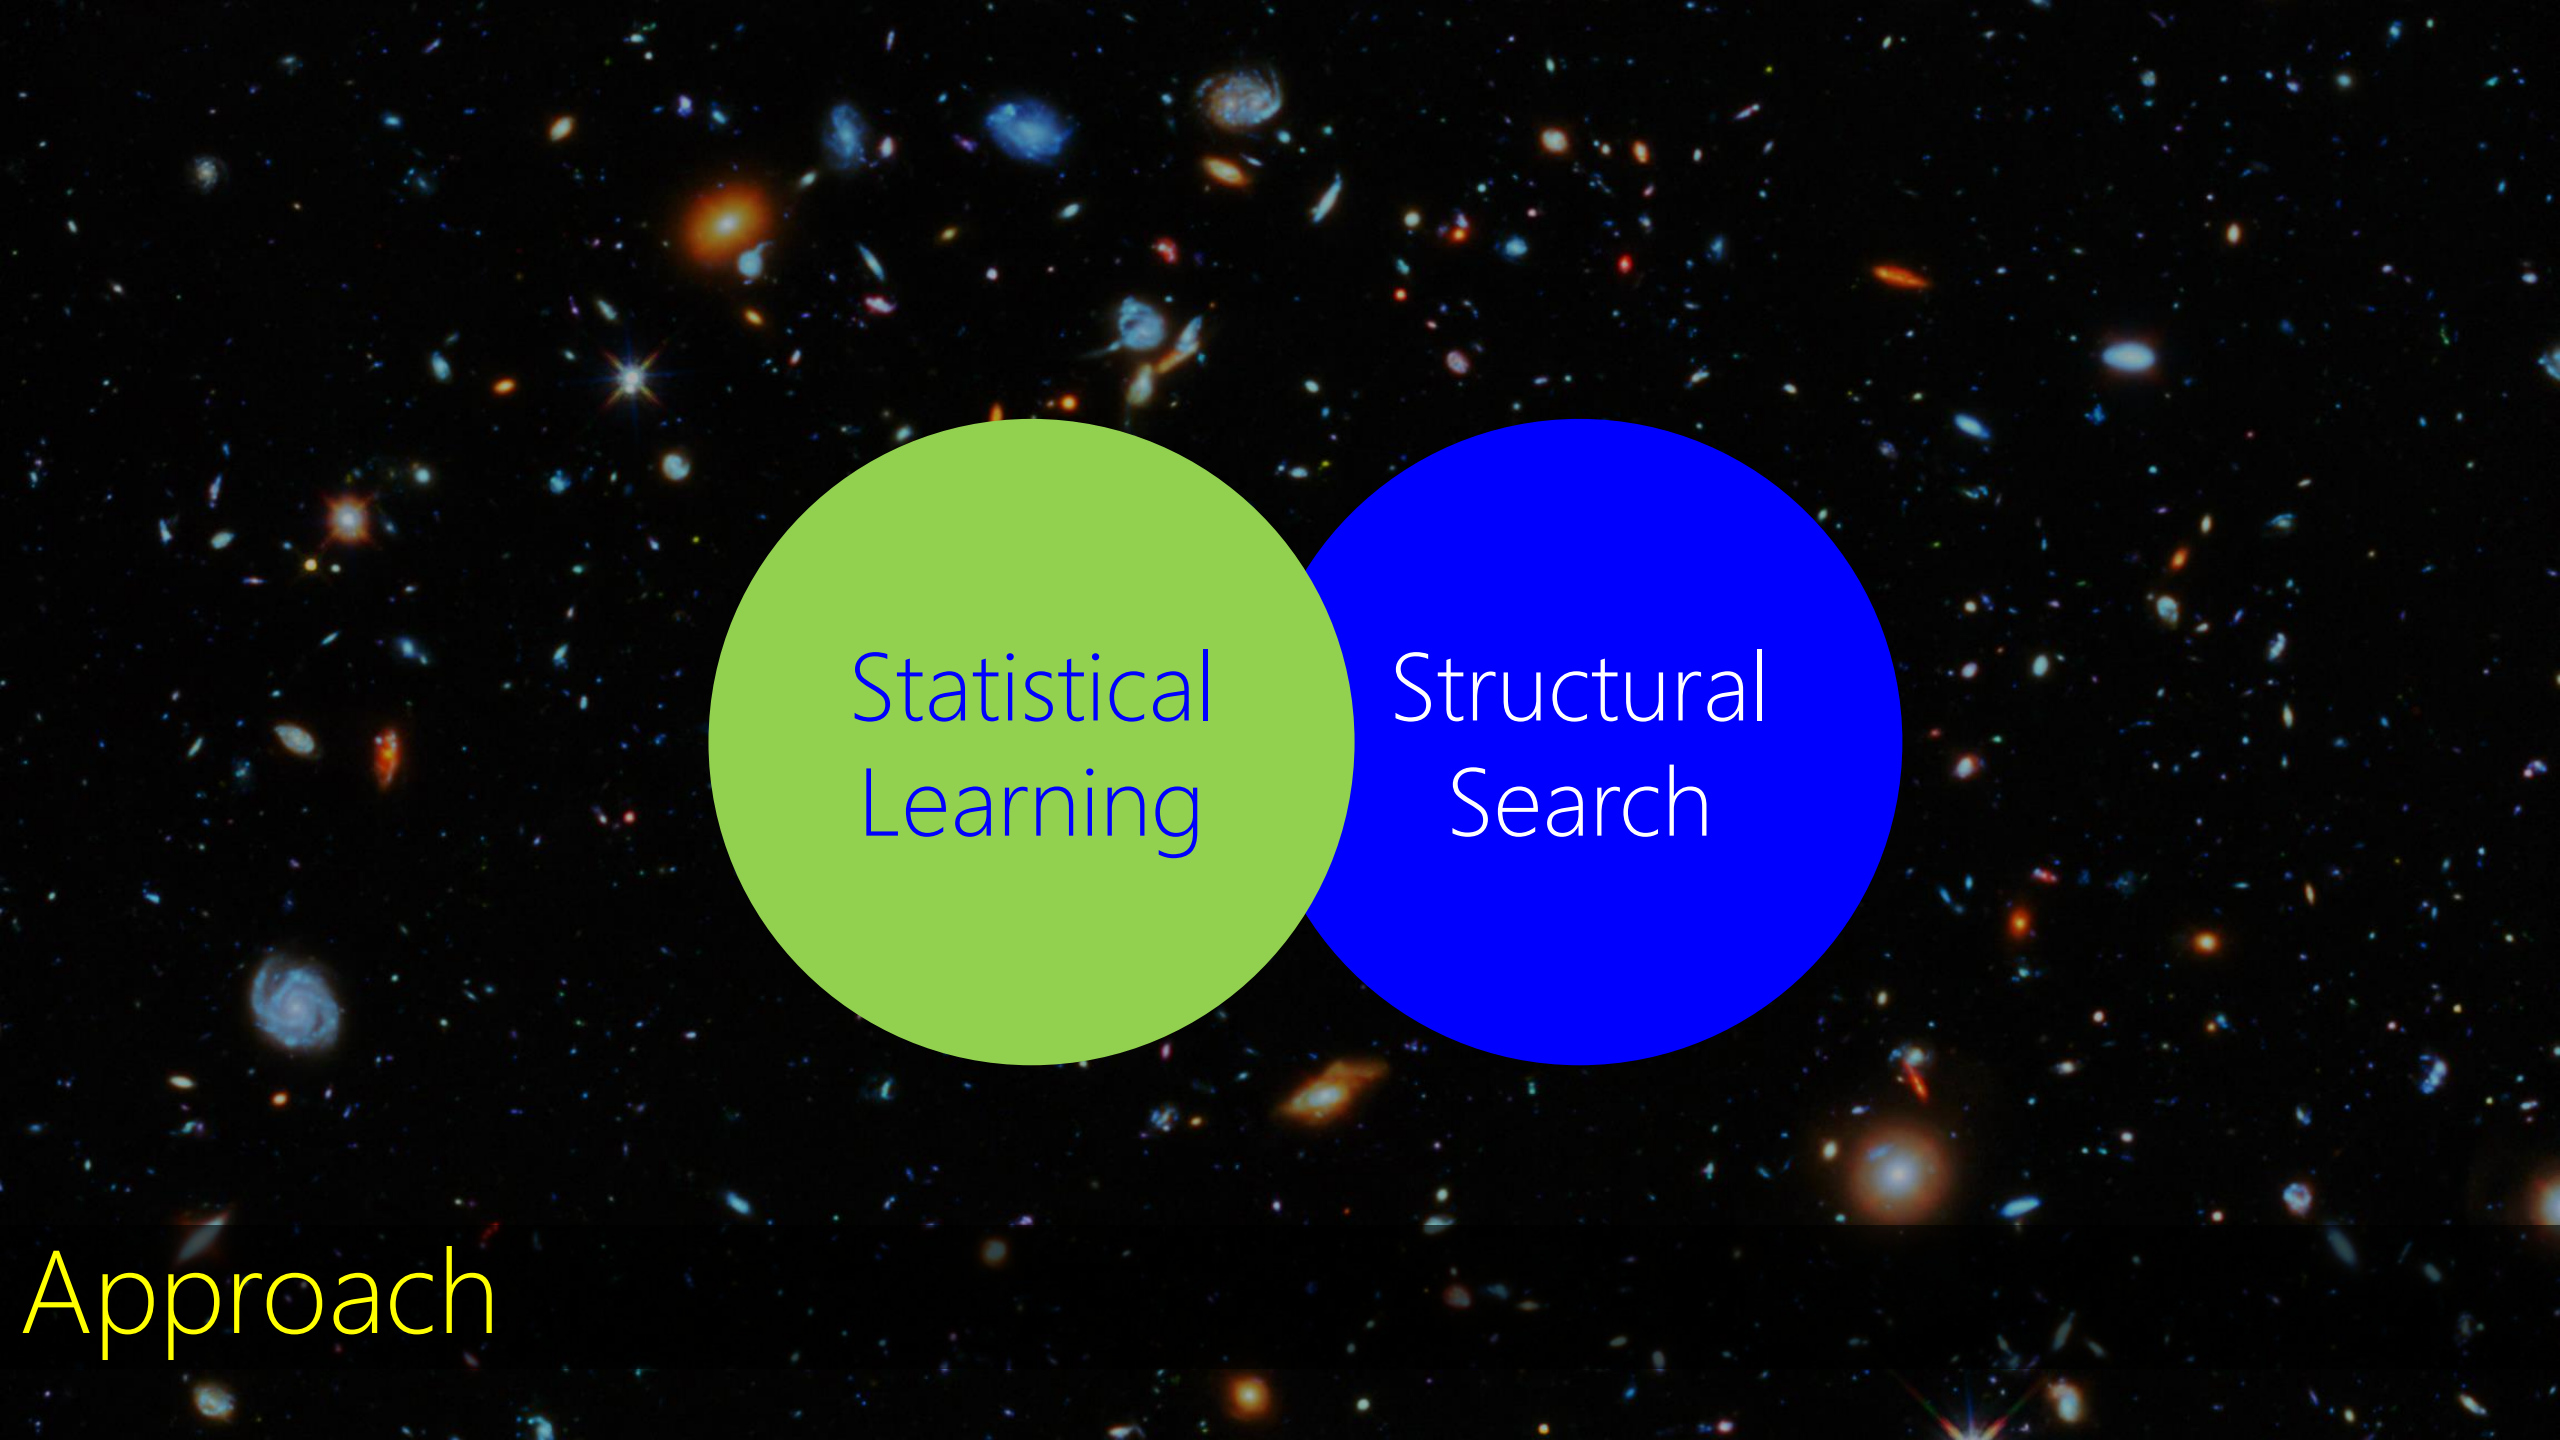

Approach

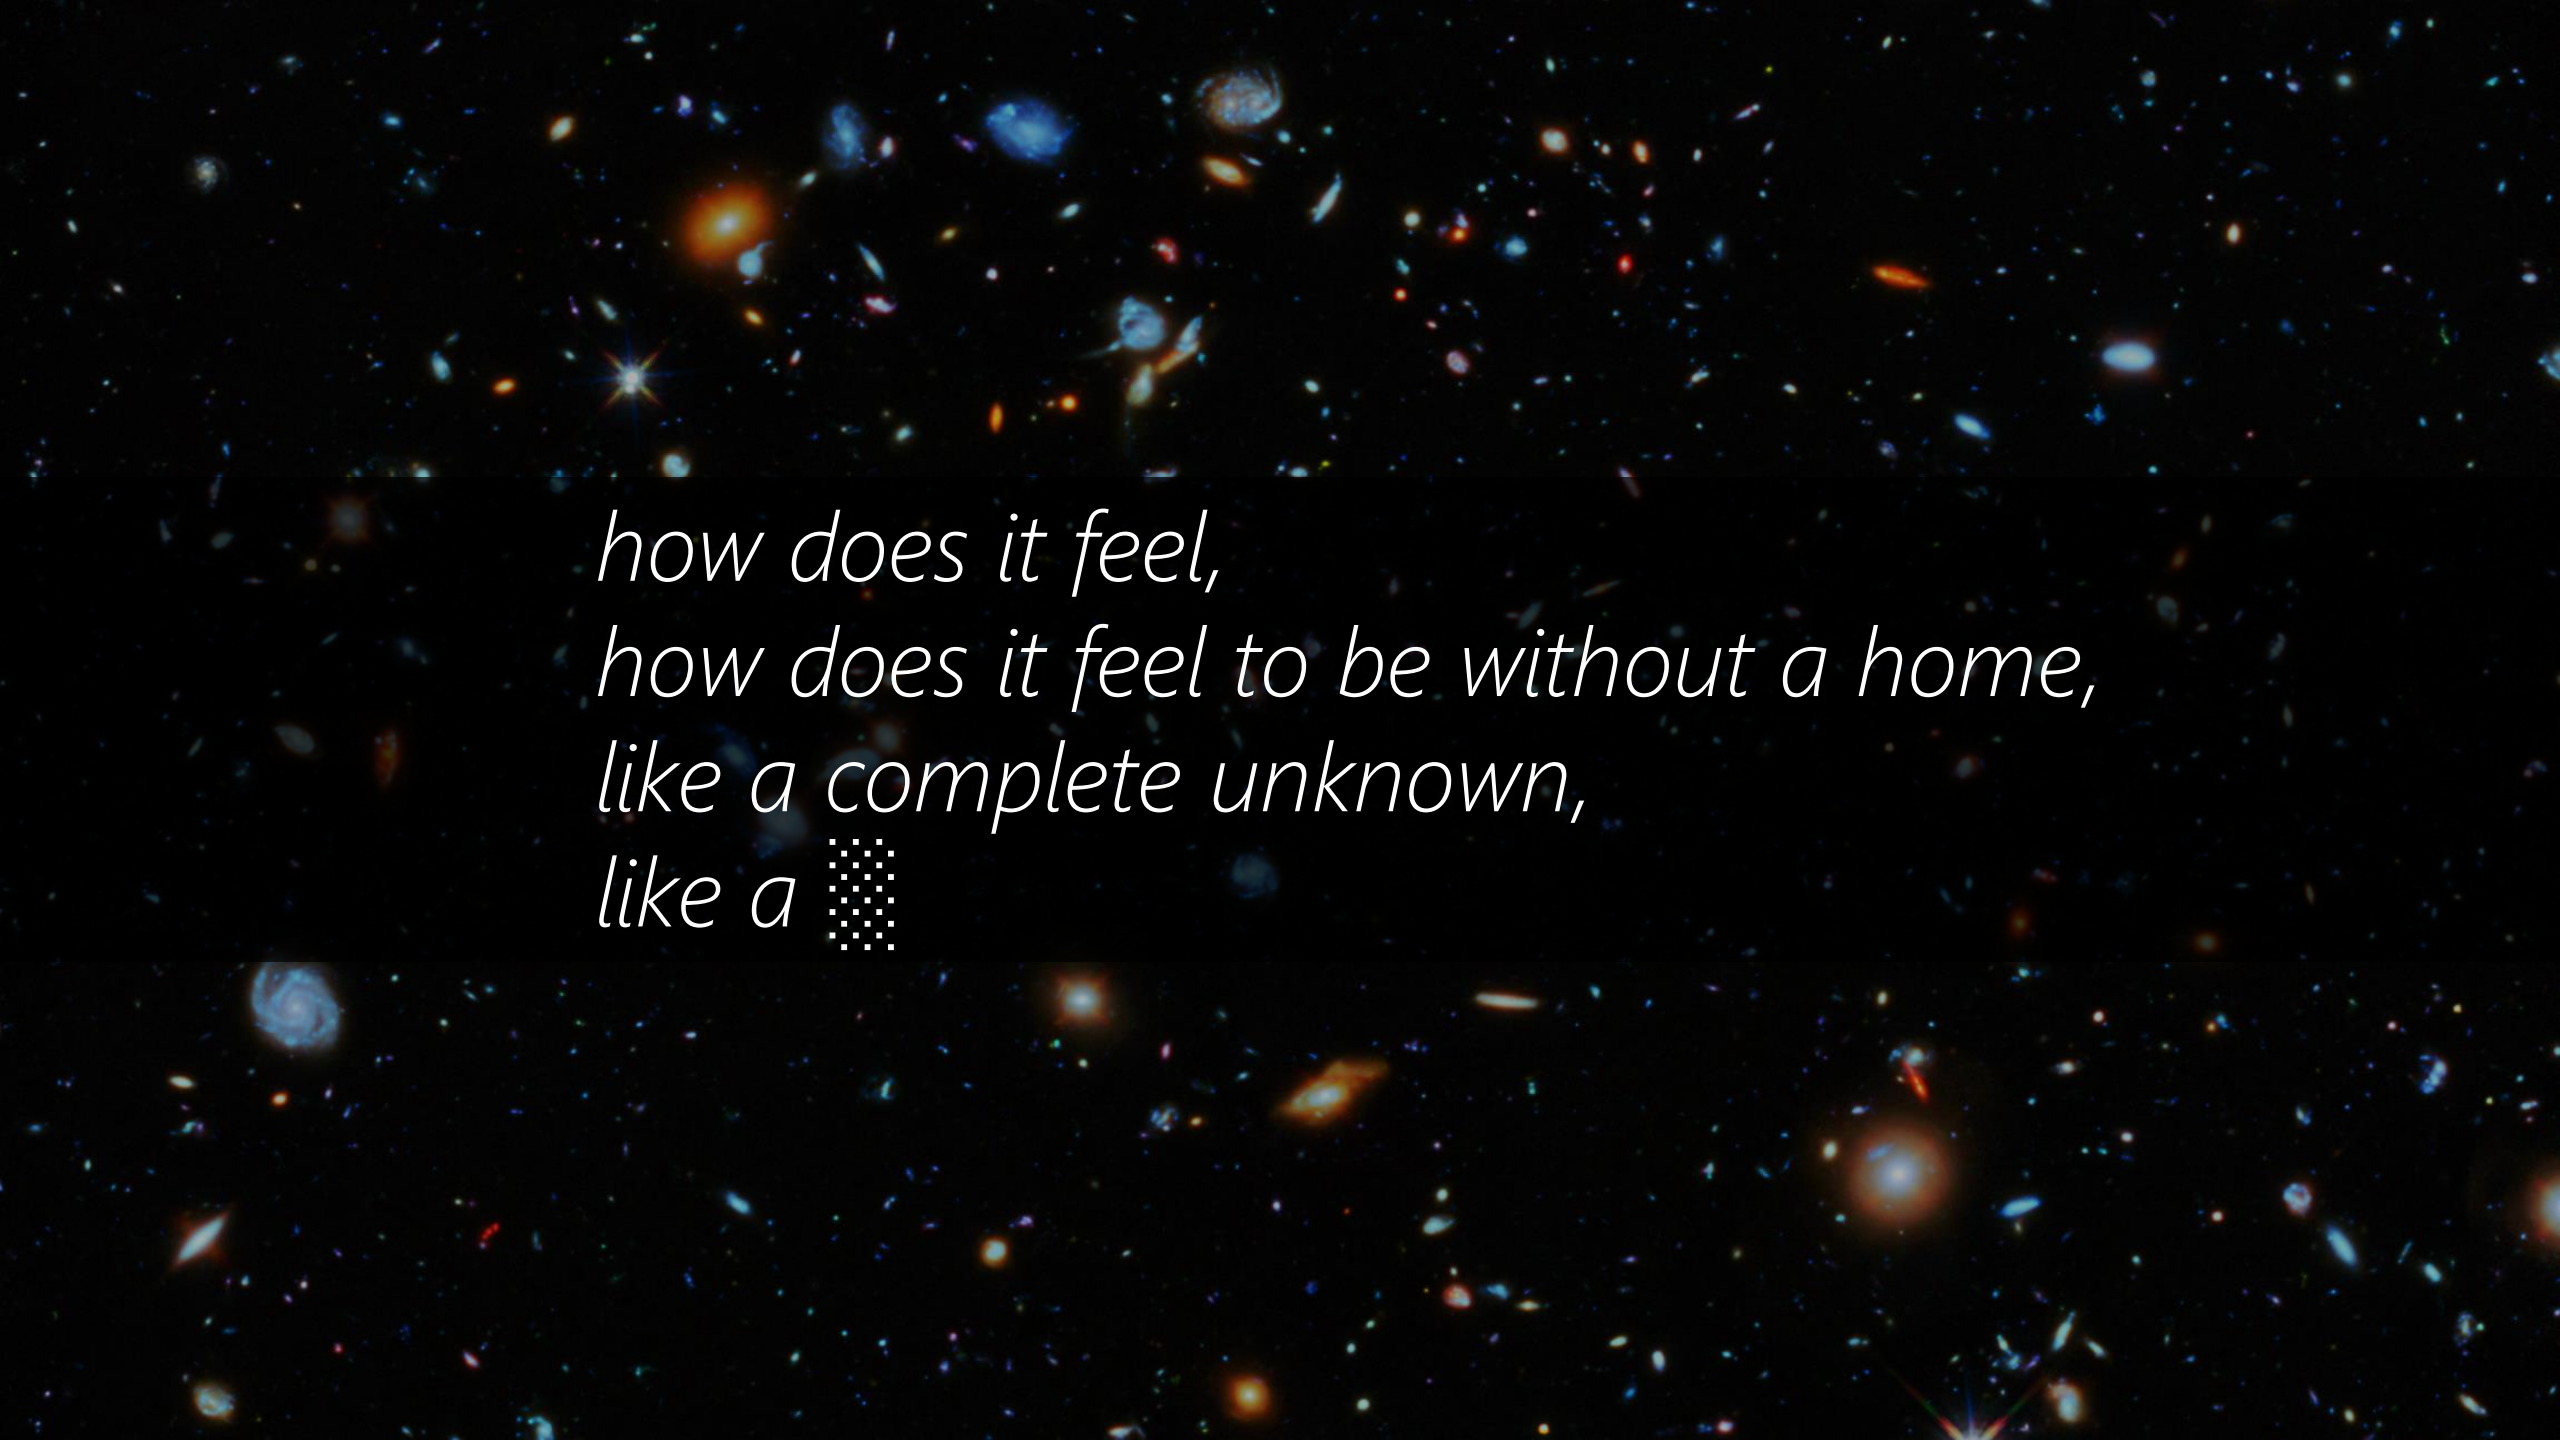
how does it feel,
how does it feel to be without a home,
like a complete unknown,
like a ▒

how does it feel,
how does it feel to be without a home,
like a complete unknown,
like a rolling stone — Bob Dylan

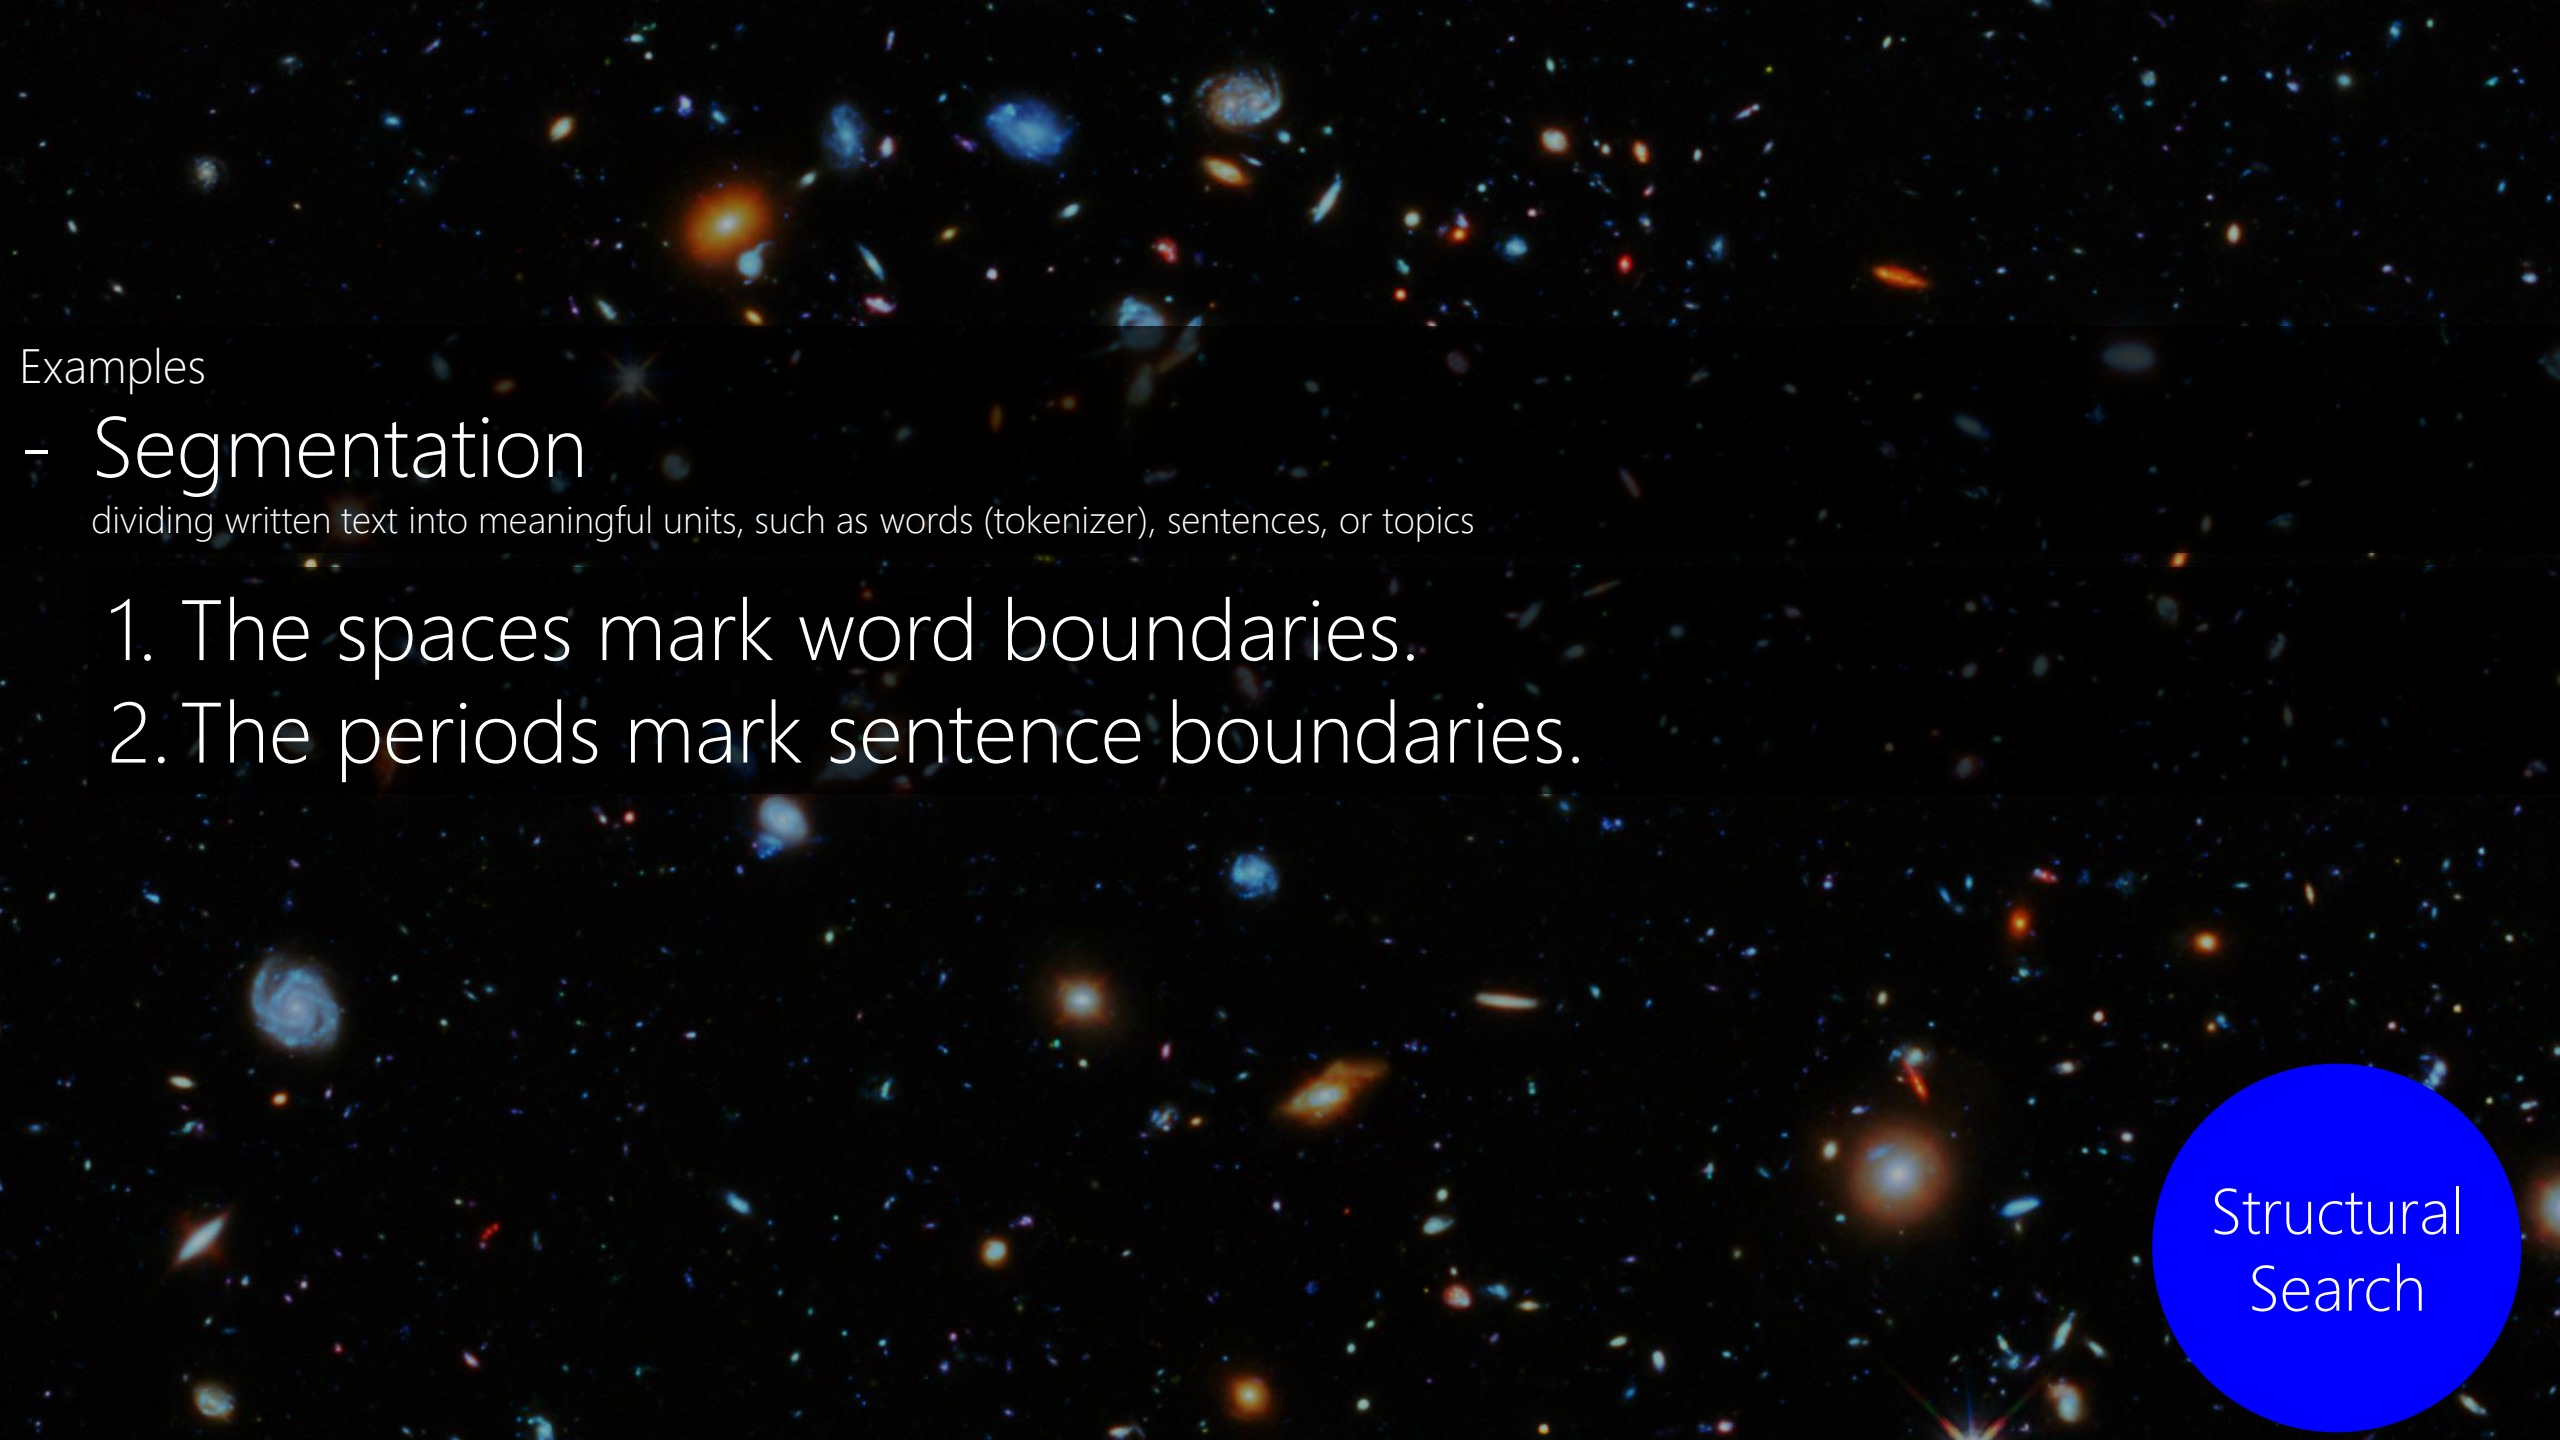

Examples

- Segmentation

  dividing written text into meaningful units, such as words (tokenizer), sentences, or topics

1. The spaces mark word boundaries.
2. The periods mark sentence boundaries.

Structural Search

Adatabaseindexisadatastructurethatimprovesthespeedofdataretrievaloperationsonadatabasetableatthecostofadditionalwritesandstoragespacetomaintaintheindexdatastructure.Indexesareusedtoquicklylocatedatawithouthavingtosearcheveryrowinadatabasetableeverytimeadatabasetableisaccessed.Indexescanbecreatedusingoneormorecolumnsofadatabasetableasexplainedin Sect214,providingthebasisforbothrapidrandomlookupsandefficientaccessoforderedrecords.Anindexisacopyofselectedcolumnsofdatafromatable,calledadatabasekeyorsimplykey,thatcanbesearchedveryefficientlythatalsoincludesalowleveldiskblockaddressordirectlinktothecompleterowofdataitwascopiedfrom.Somedatabasesextendthepowerofindexingbylettingdevelopercreateindexesonfunctionsorexpressions.Forexample,anindexcouldbecreatedonupper(last_name),whichwouldonlystoretheuppercaseversionsofthelast_namefieldintheindex.Anotheroptionsometimessupportedistheuseofpartialindices,whereindexentriesarecreatedonlyforthoserecordsthatsatisfysomeconditionalexpression.Afurtheraspectofflexibilityistopermitindexingonuserdefinedfunctions,aswellasexpressionsformedfromanassortmentofbuilt-infunctions.

Statistical Learning

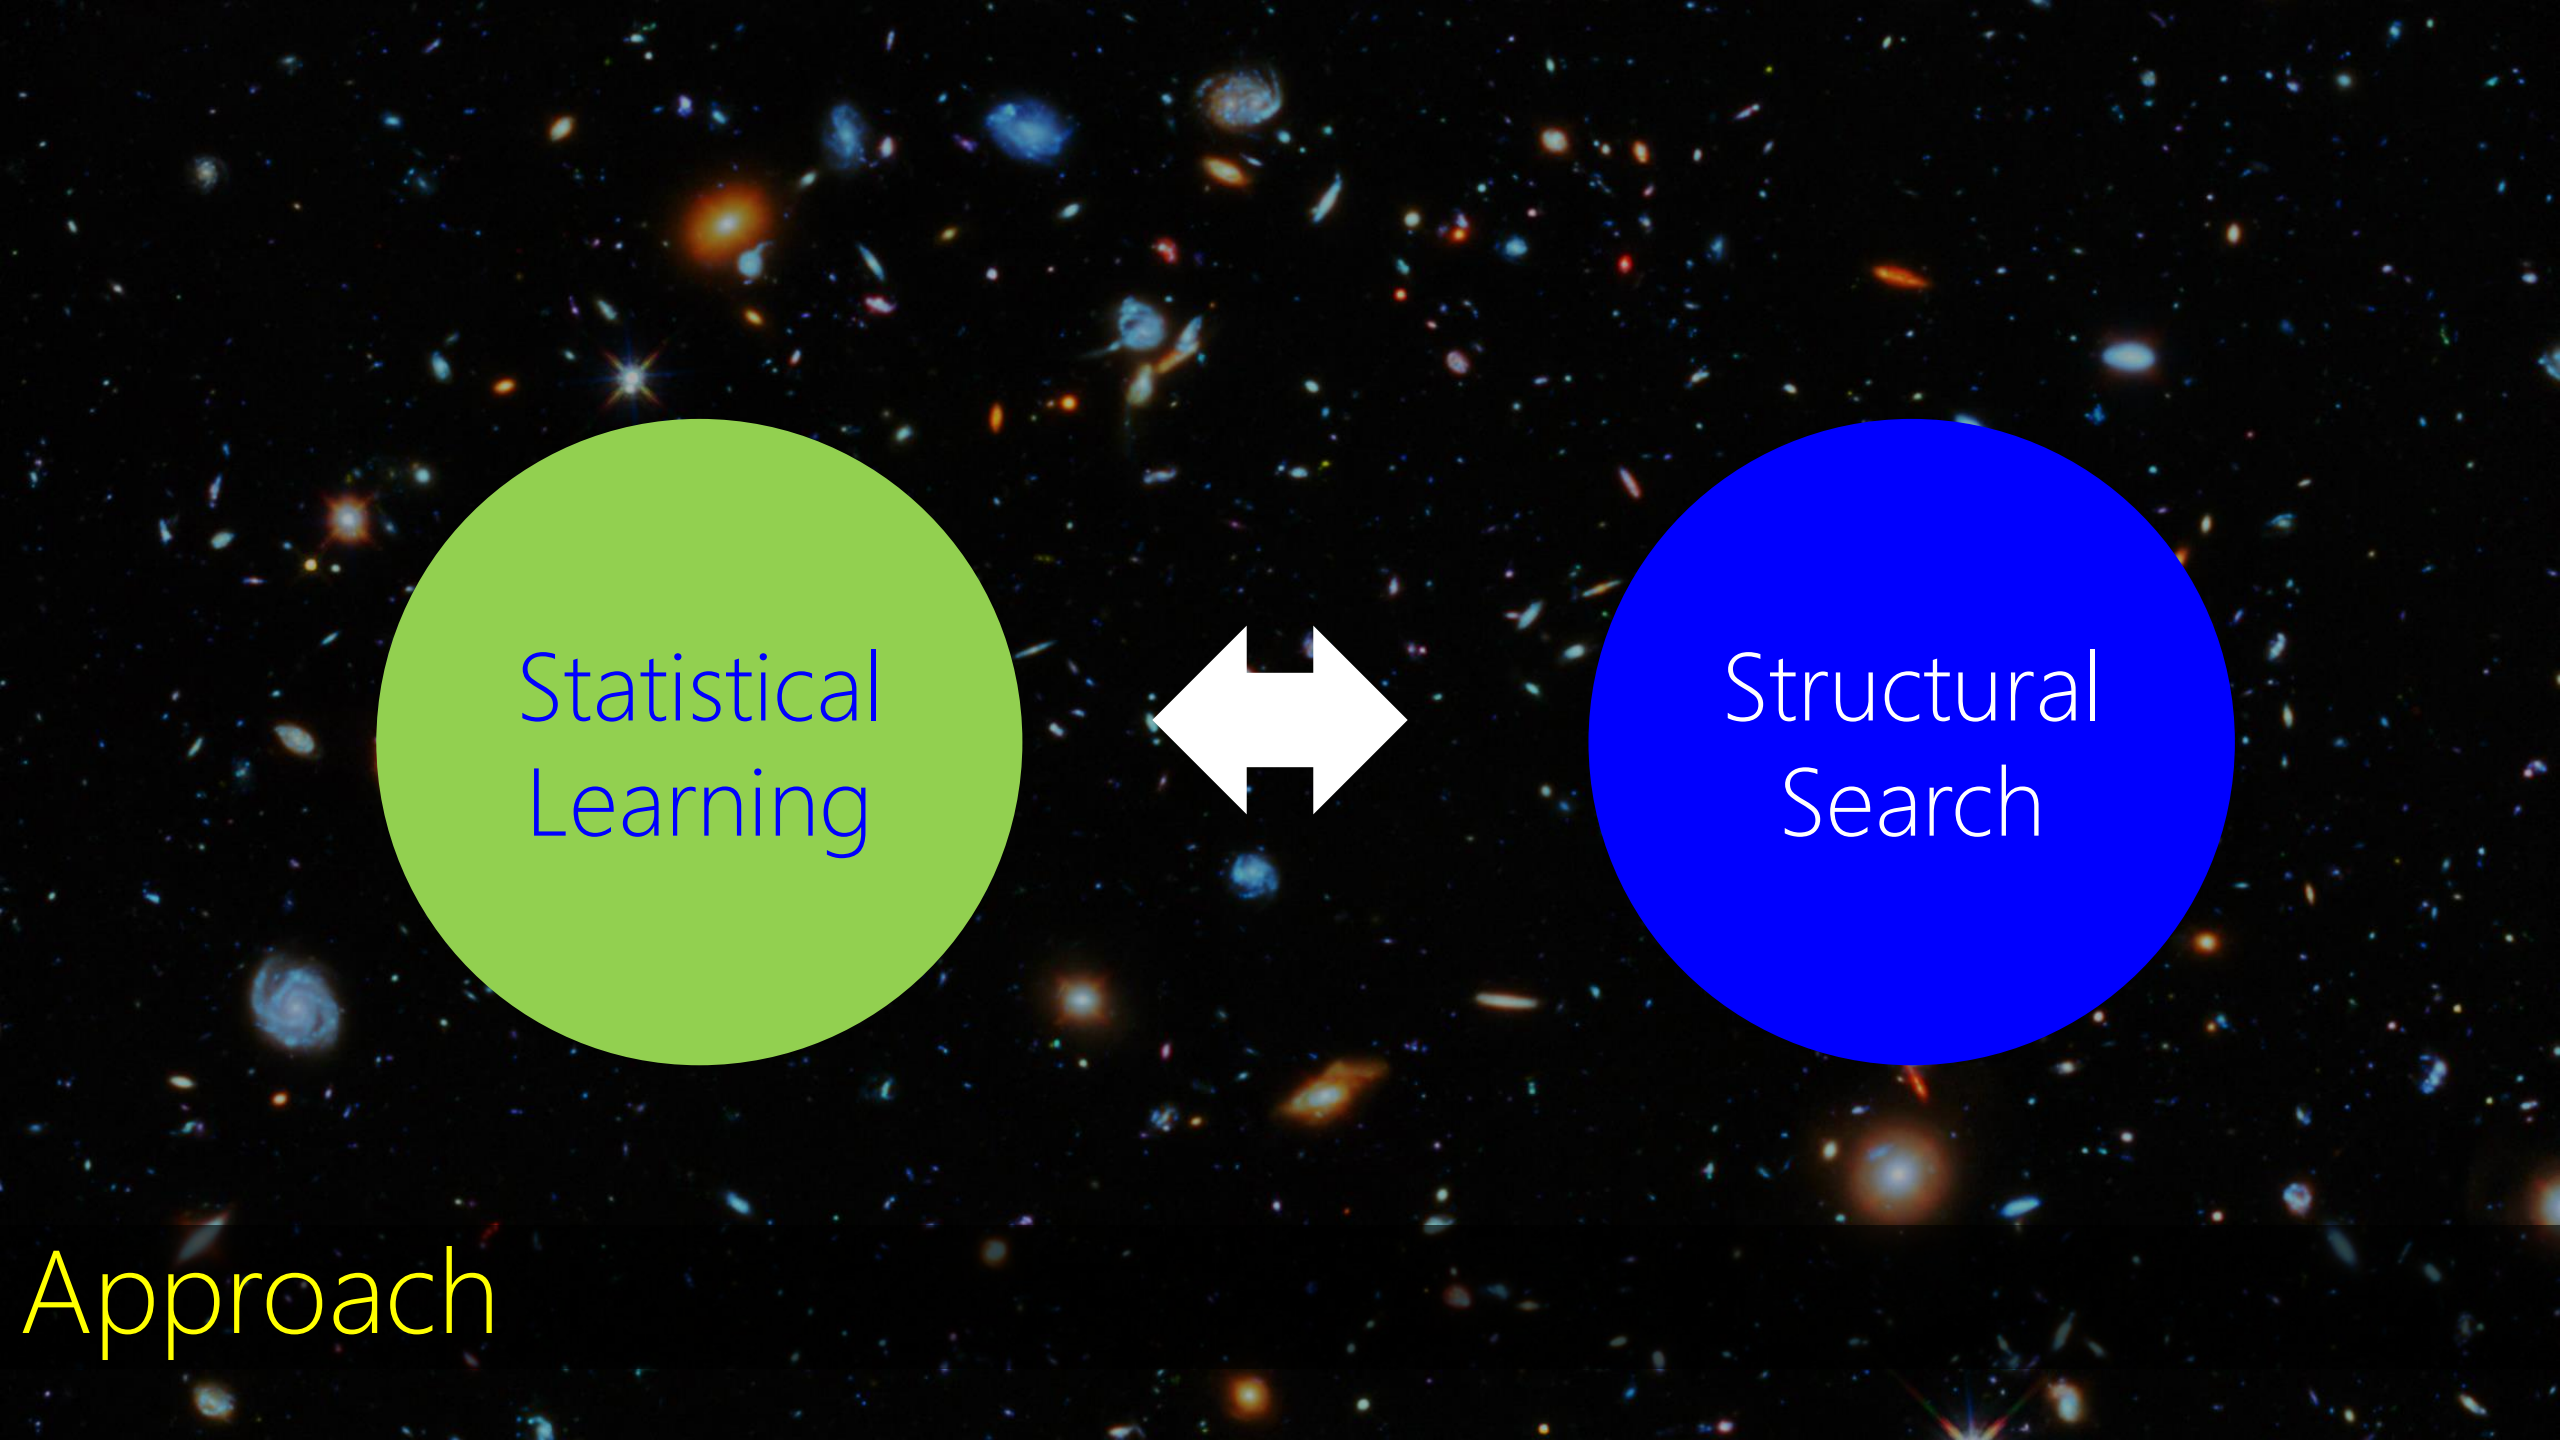

Approach

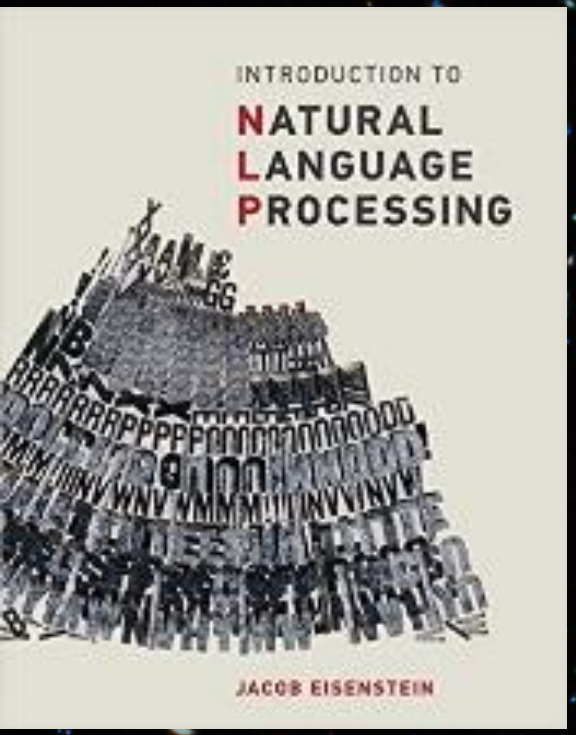

No machine learning specialist likes to be told that their engineering methodology is unscientific alchemy

Nor does a linguist want to hear that the search for general linguistic principles and structures has been made irrelevant by big data

Yann LeCun and Christopher Manning discuss Deep Learning and Innate Priors
https://www.youtube.com/watch?v=fKk9KhGRBdI

Machine Learning & Data Mining

Methods of learning from data.
Text data needs special care! Why?

Natural Language Processing & Text Mining

The goal is to provide new computational capabilities for applications
E.g., *predict next form of a word for branding!*

- extracting information from texts,
- translating between languages

Computational Linguistics

Here, language is the object of study.
Computational methods are to support.
Just as in computational biology

E.g, *how a word is evolving in time?*

- Discourse analysis
- Parsing

Neighbors

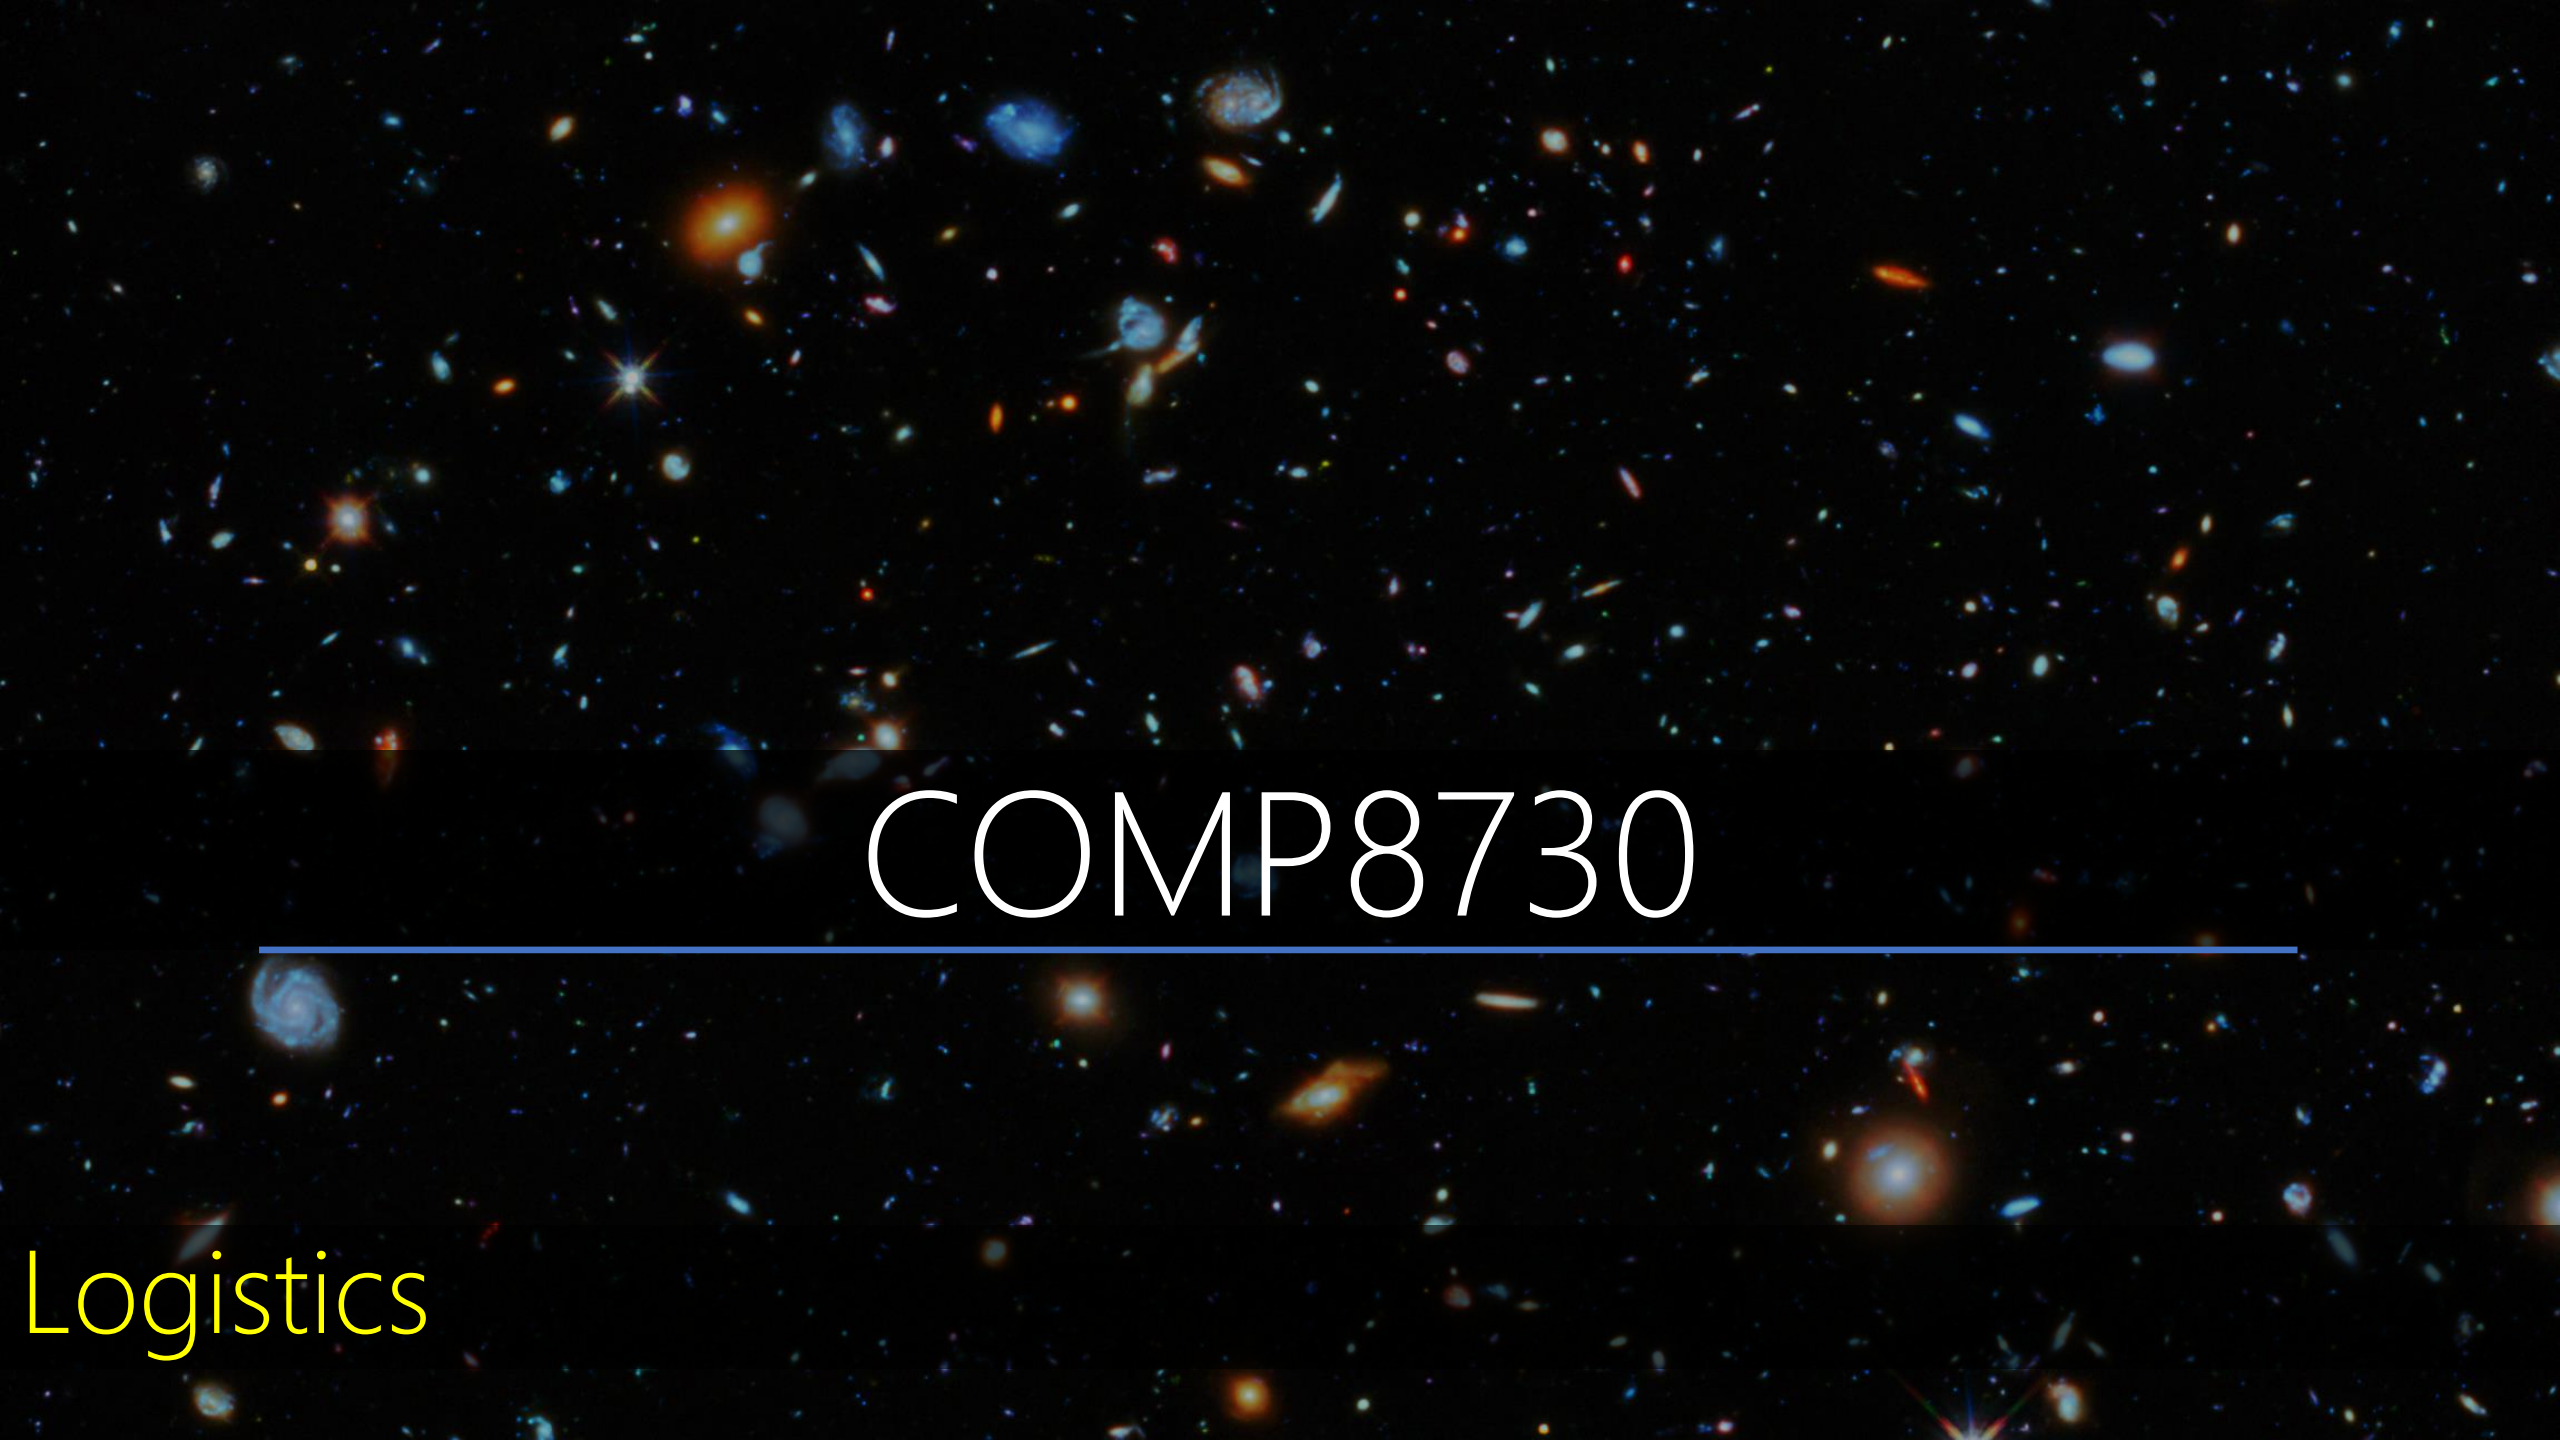

COMP8730

Logistics

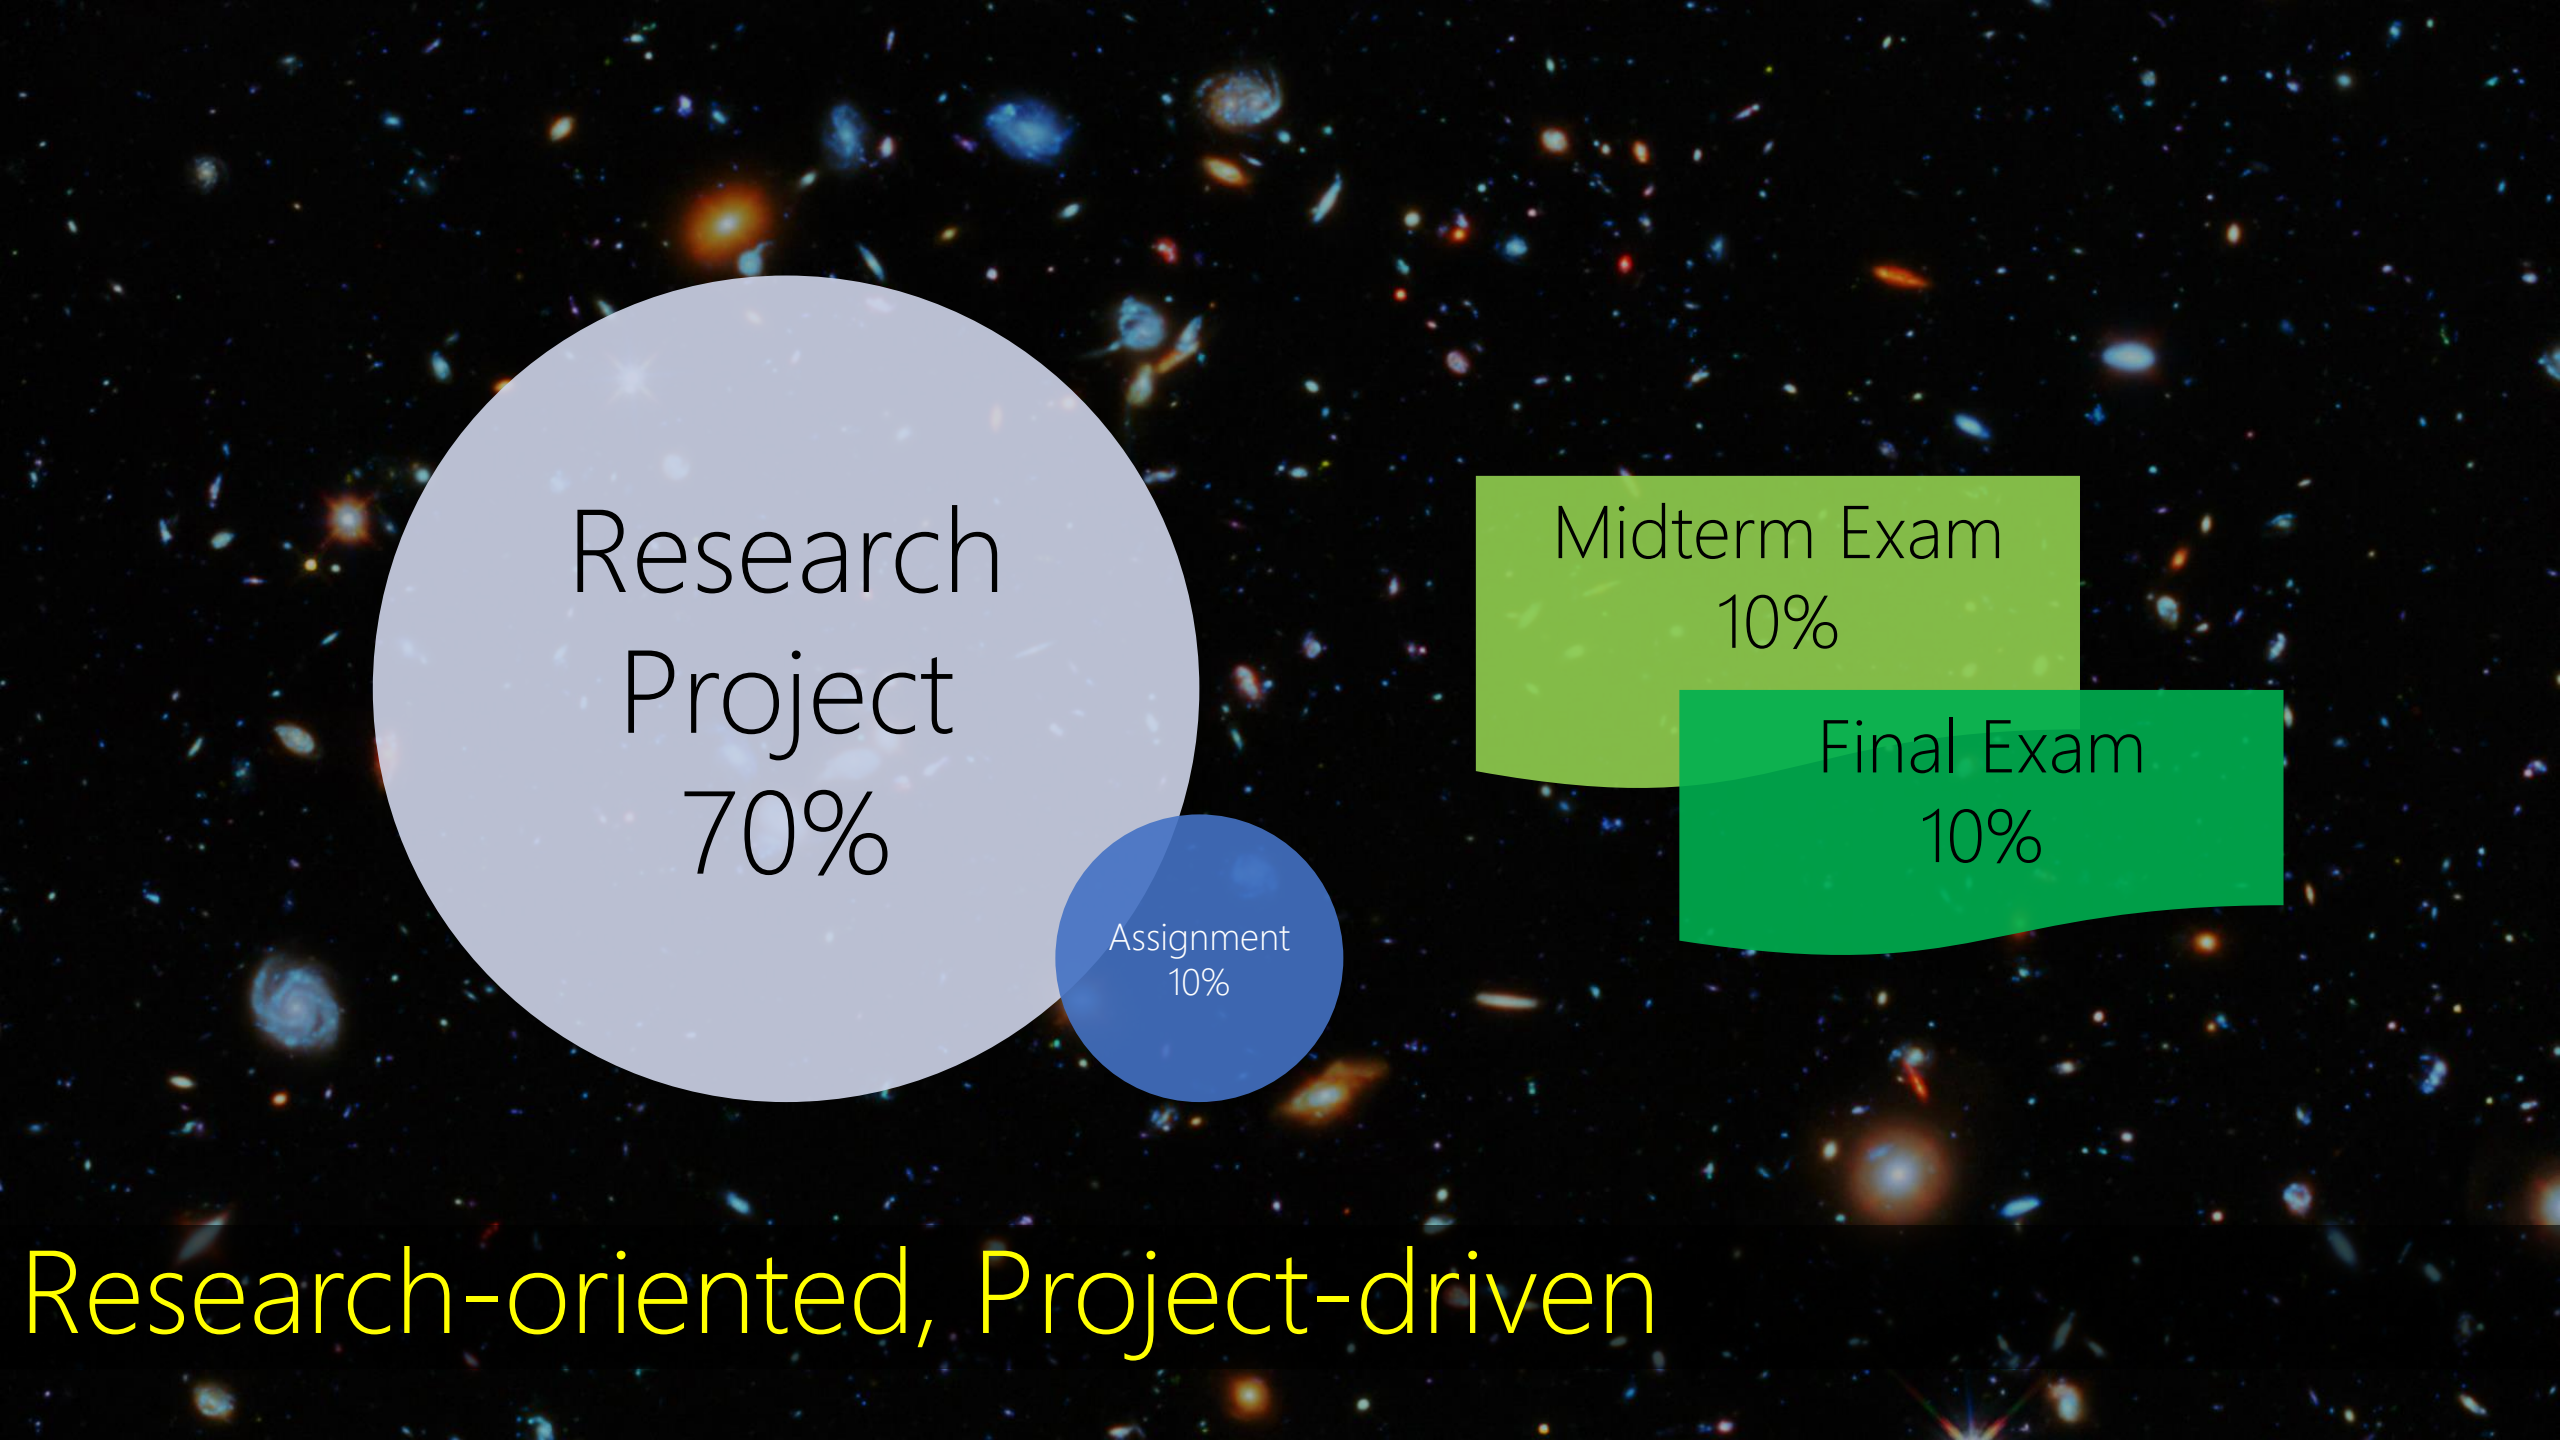

Research-oriented, Project-driven

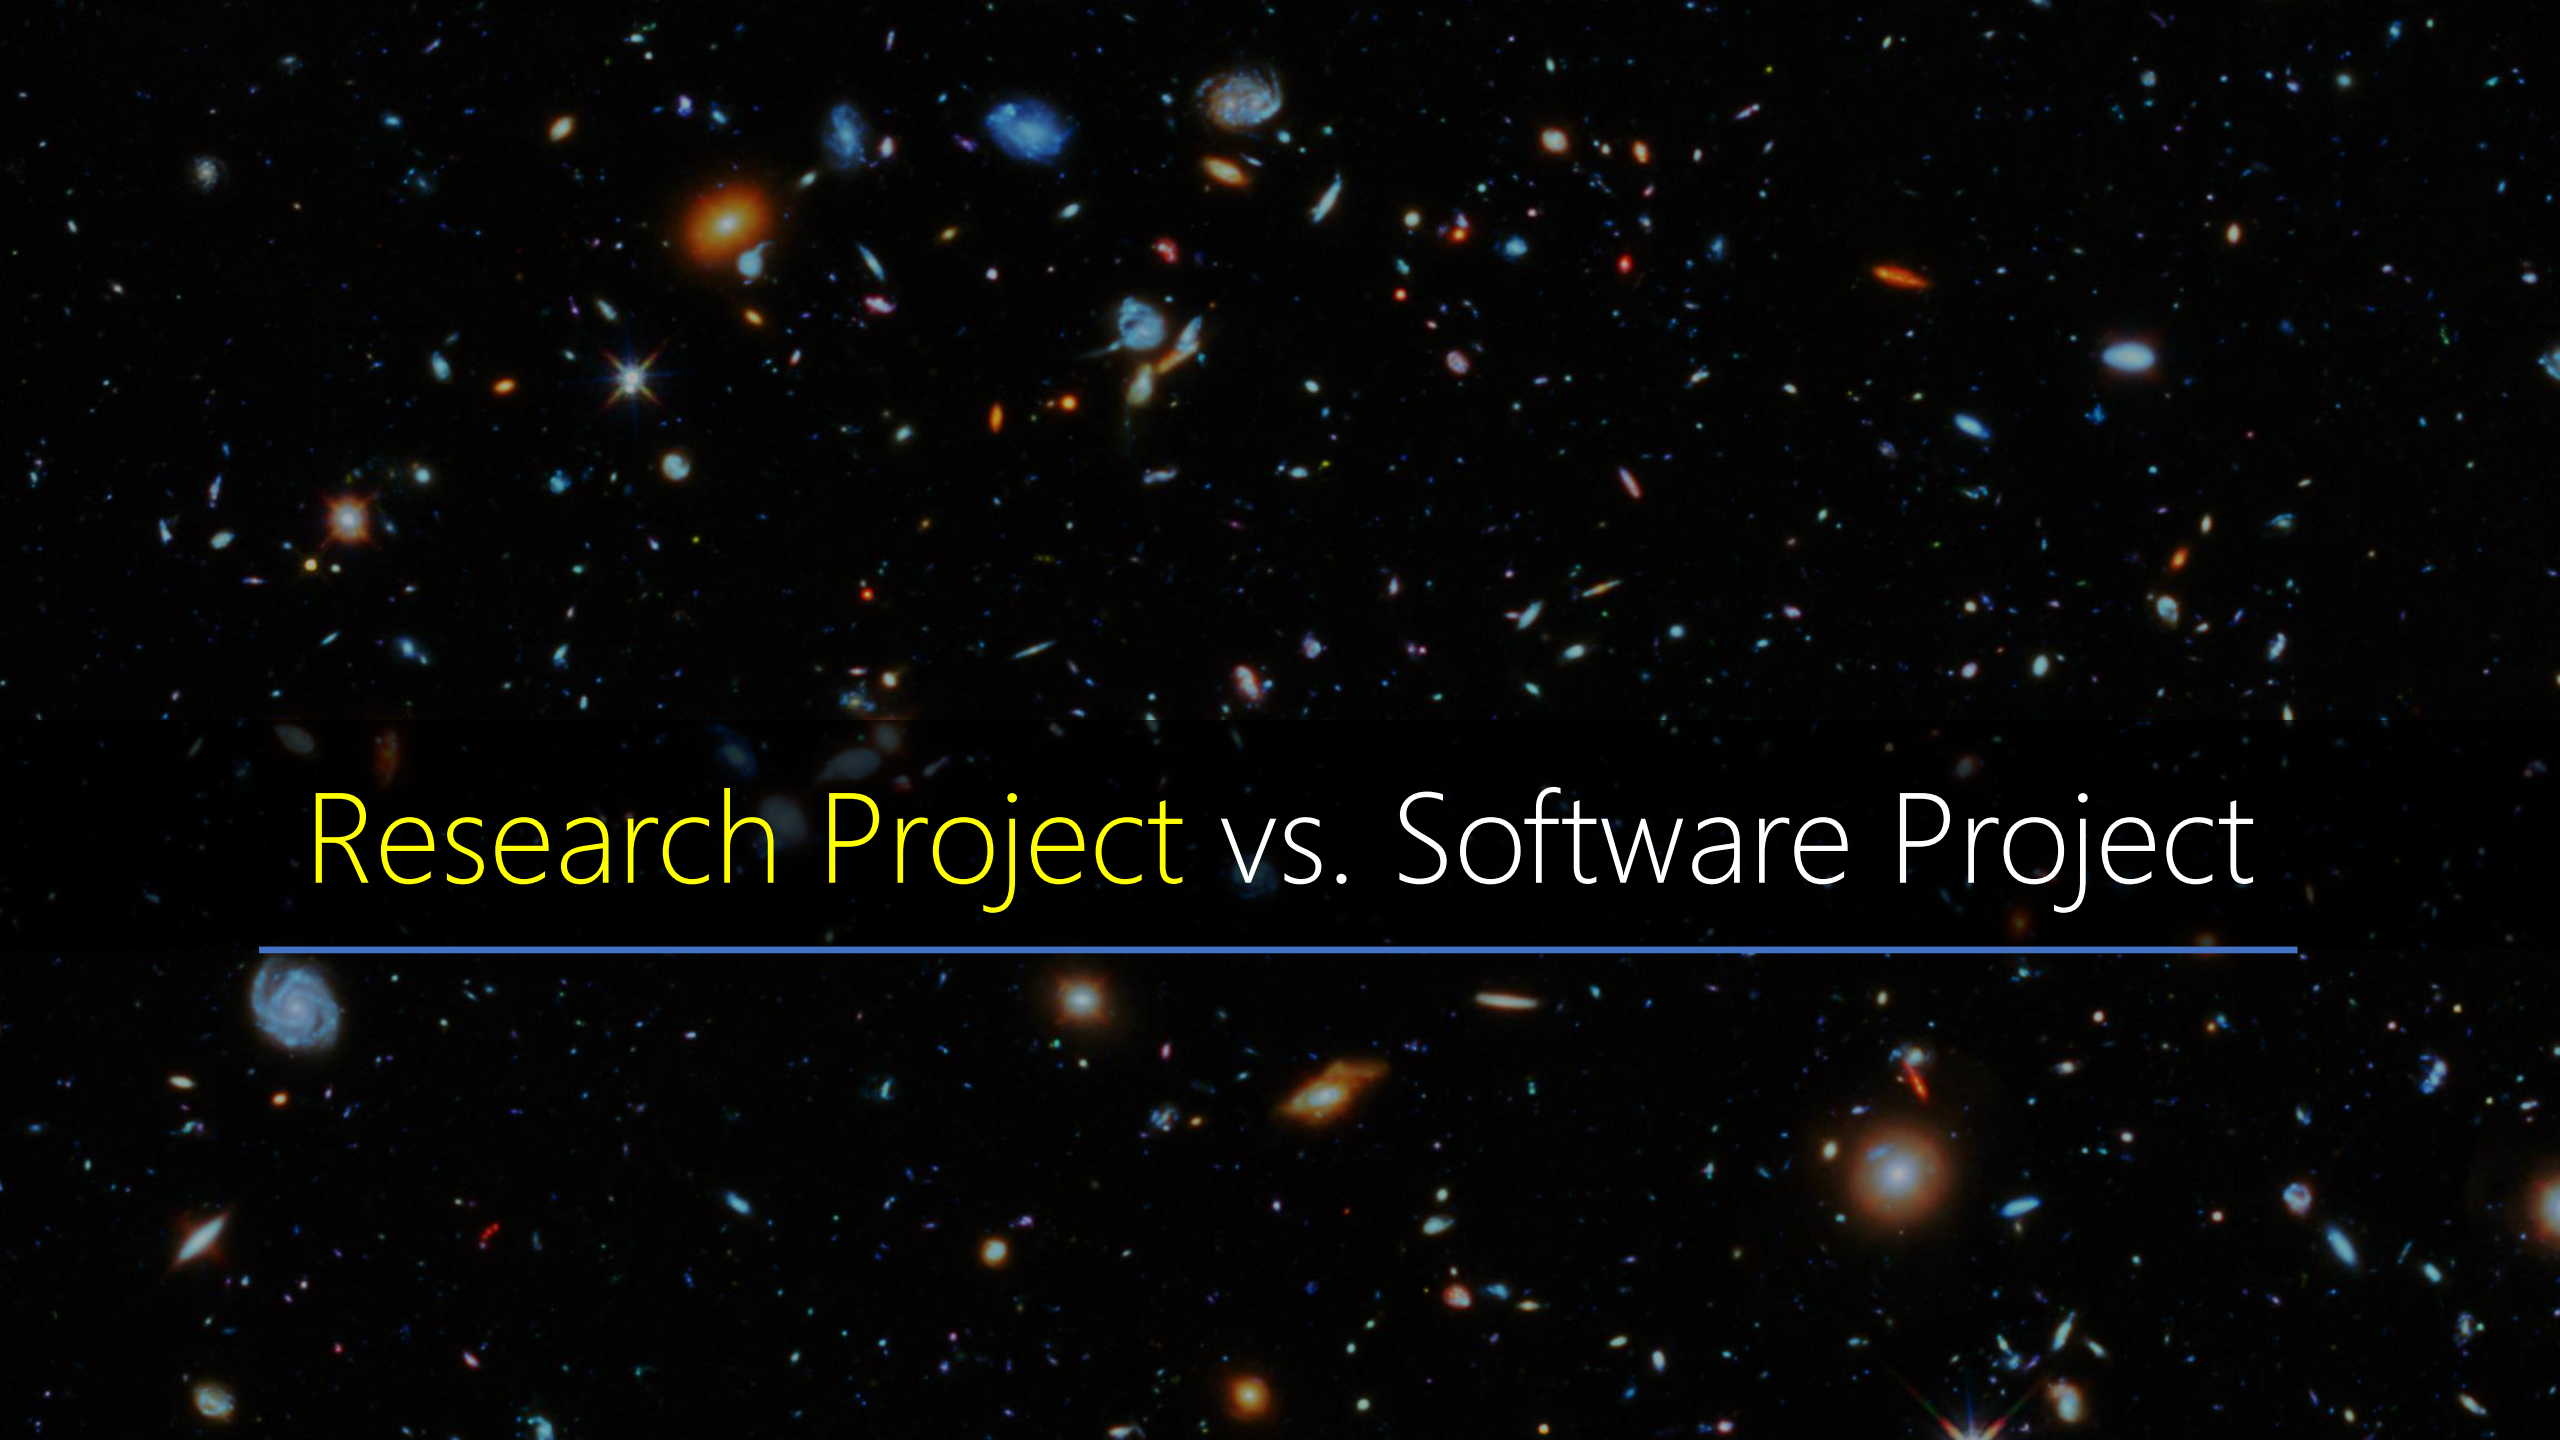

Research Project vs. Software Project

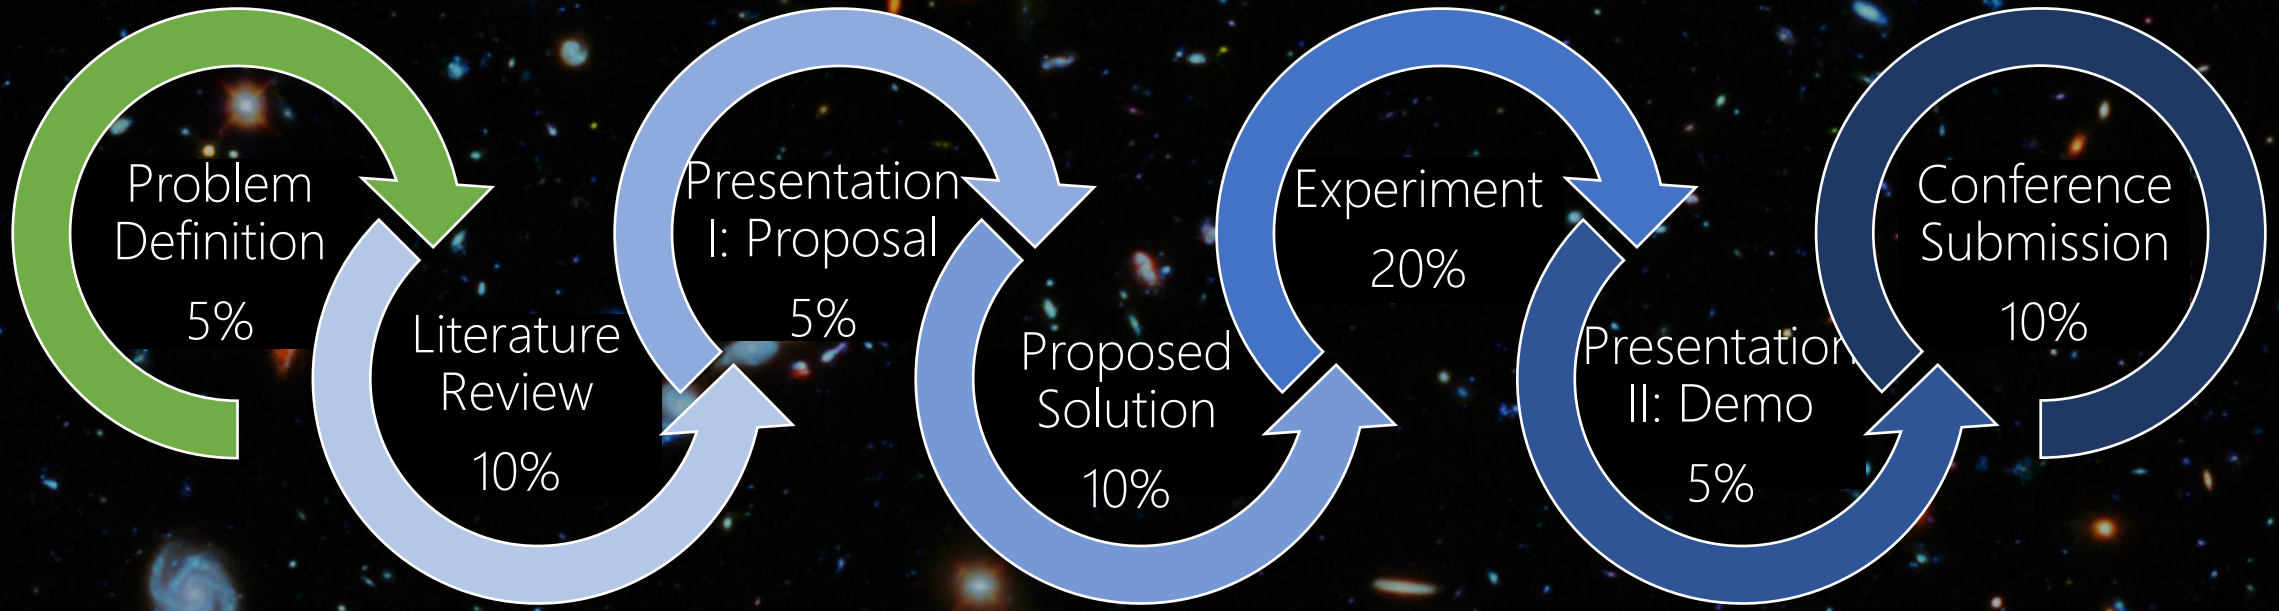

Research Project + 5% Peer Review

Lec02: Text Normalization
Lec03: n-Gram Language Models
Lec04: Sentiment Classification
Lec05: Logistic Regression
Lec08: Vector Semantics and Embedding
Lec09: Neural Language Models
Lec10: Sequence Labeling
Lec11: Topic Modeling, User Modeling
Lec12: Web Search

Theory

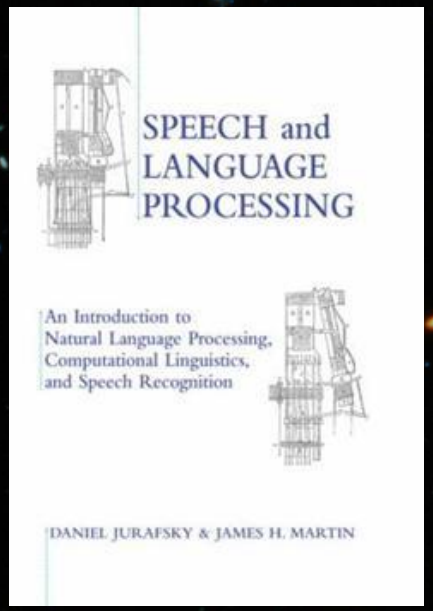

Speech and Language Processing, 3rd Ed. Draft
Dan Jurafsky and James H. Martin
https://web.stanford.edu/~jurafsky/slp3/

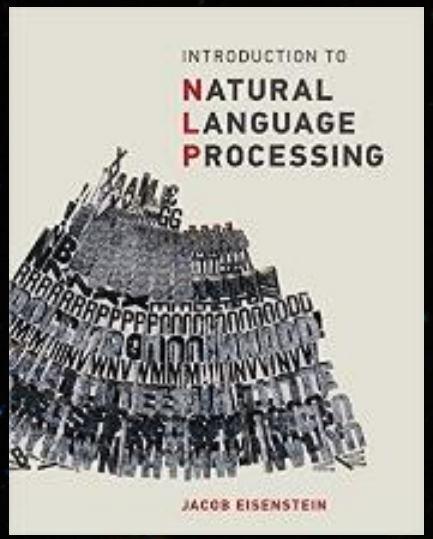

Introduction to Natural Language Processing
Jacob Eisenstein
http://cseweb.ucsd.edu/~nnakashole/teaching/eisenstein-nov18.pdf

Reference Books

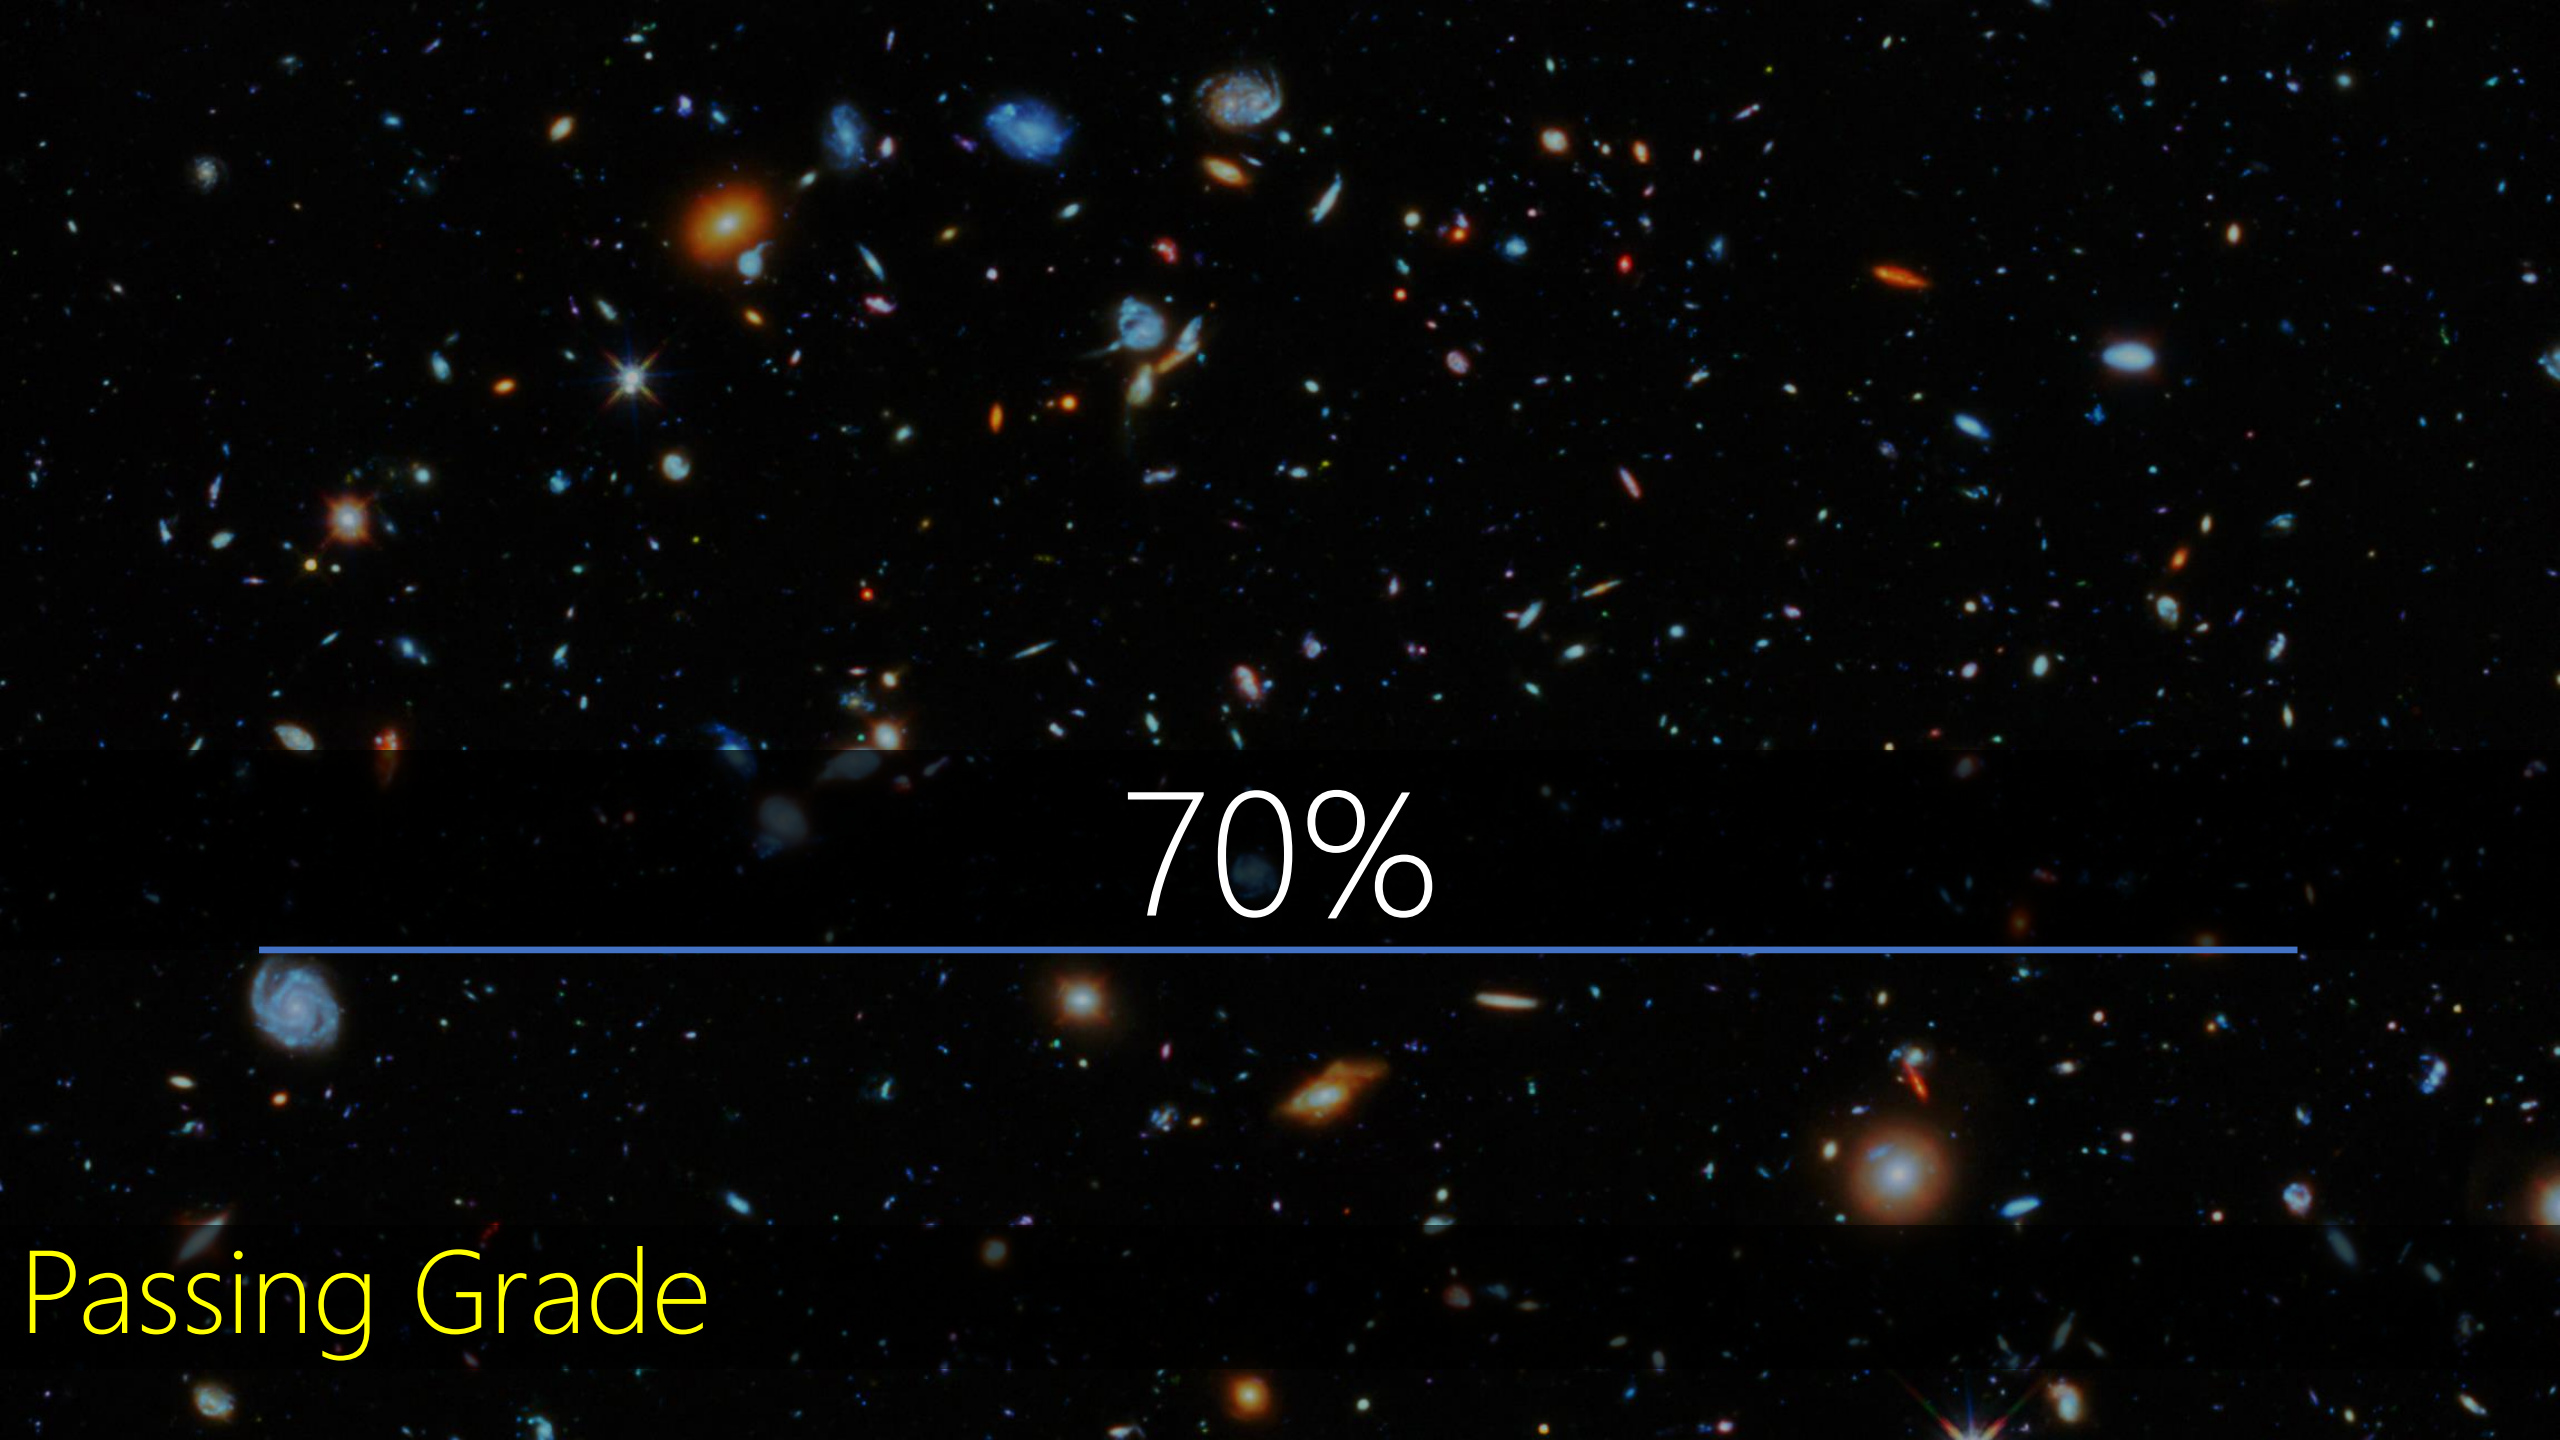
70%
Passing Grade

Hossein Fani, PhD
Assistant Professor, School of Computer Science, Faculty of Science, University of Windsor
Room 5111, Lambton Tower
hfani@uwindsor.ca
hfani.myweb.cs.uwindsor.ca


Mondays 2:30PM-3:30PM
Office Hour

Weekly Schedule

Eastern Time	MON	TUE	WED	THU	FRI	SAT	SUN
	Manual for Submissions						
	Deadline for Submissions						
	Mark for Submissions						
11:30	Lecture A						
12:45							
13:00	Lecture A						
14:20							
14:30	Office Hour						
15:30							

COMP8730-1-R-2022W (Natural Language Processing & Understanding)

Homepage

Weekly Schedule

Announcements

Learning Outcome

Syllabus (Outline)

Lectures

Classroom (Online)

Exams

Research Project

My Grades

Discussion Board

Office

Email

Help

Virtual Classroom Tutorial

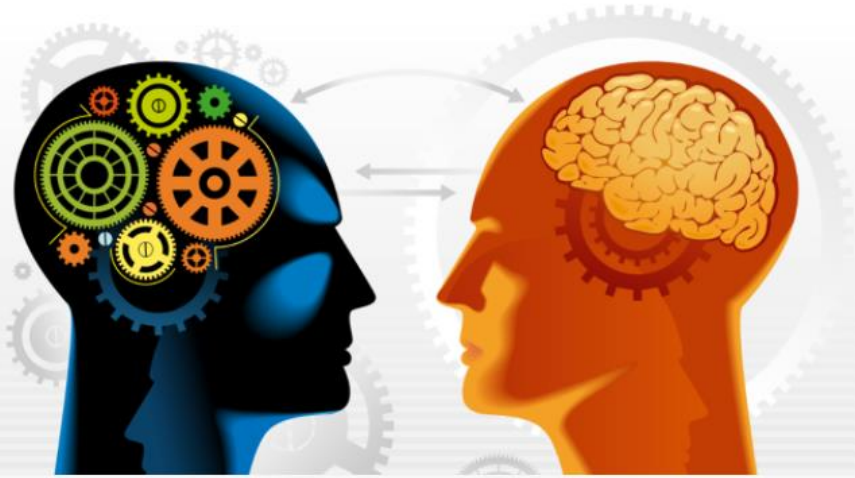

Image Source: https://unanimous.ai/chat-with-a-different-kind-of-artificial-intelligence

Homepage

Add Course Module

Customize Page

My Announcements

No Course or Organization Announcements have been posted in the last 7 days.

To Do

Edit Notification Settings

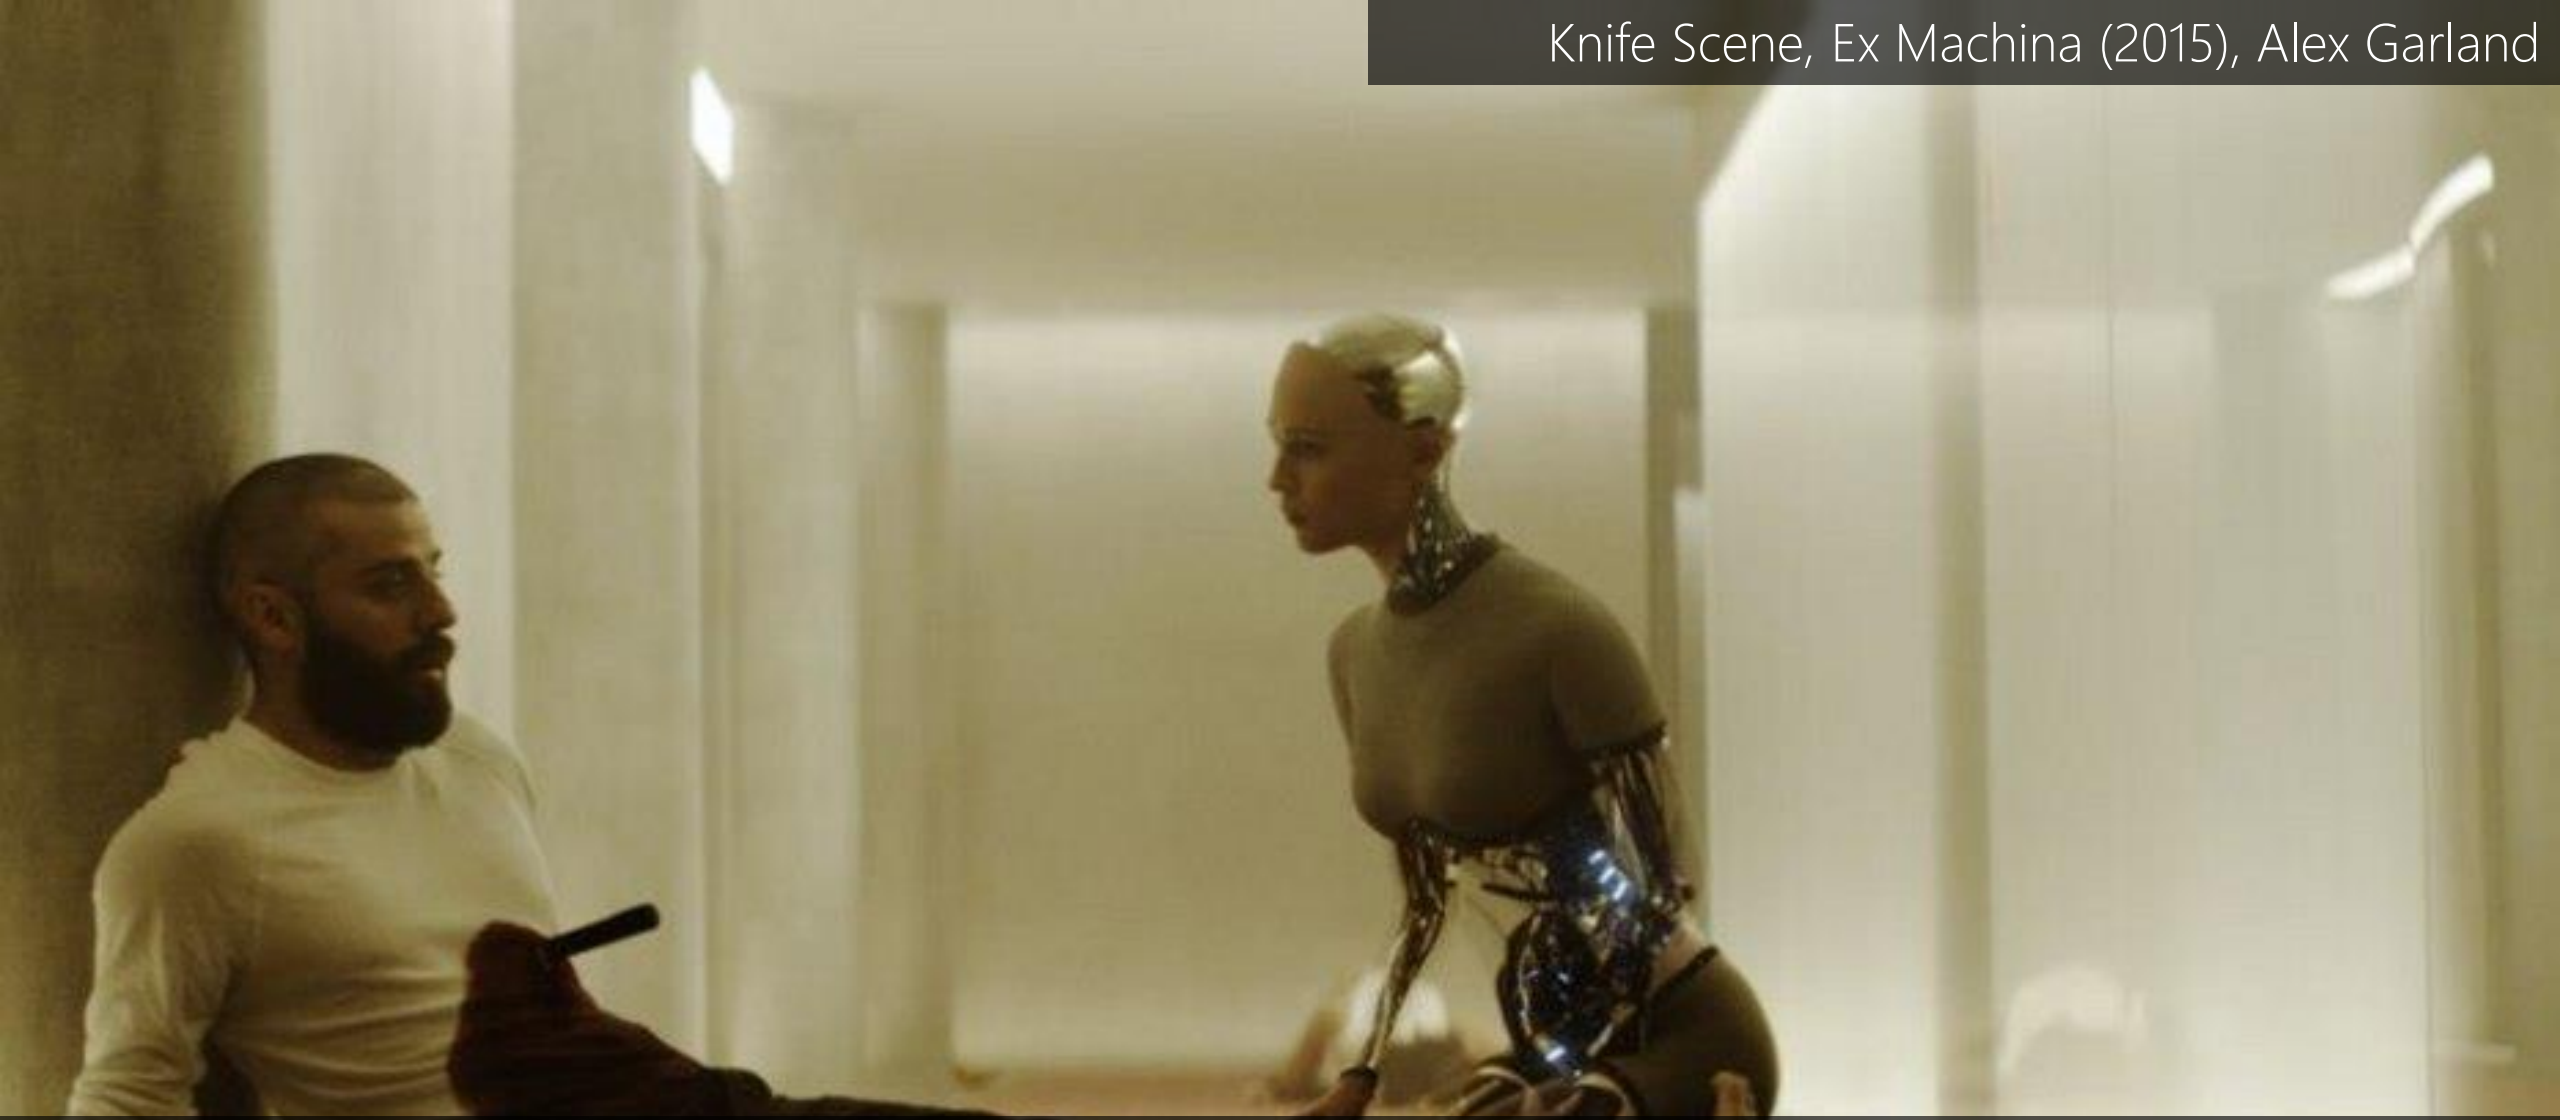
Knife Scene, Ex Machina (2015), Alex Garland

/DAY uus + eks + MAA ki naa/

Deus Ex Machina means "god from the machine." In ancient Greek theater, when actors playing gods carried onto stage by a machine. These gods would then serve as the ultimate arbiters of right and wrong and decide how the story ends. But this film is just called "Ex Machina" without the "Deus." A machine without a god.

https://www.looper.com/148401/the-ending-of-ex-machina-finally-explained/